

Binding and Storage



COMP 524: Programming Language Concepts
Björn B. Brandenburg

The University of North Carolina at Chapel Hill

What's the most striking difference?

```
/*
 * A simple interface to describe finite automata. {@link GraphvizFormatter} uses
 * this to visualize arbitrary finite automata.
 */
public interface Automaton {
    /**
     * @return Something that can be iterated to visit all states.
     */
    public Iterable<State> getStates();

    /**
     * @return the start state (which must be, by definition, unique)
     */
    public State getStartState();

    /**
     * @return Something that can be iterated to visit all final states.
     */
    public Iterable<State> getFinalStates();

    /**
     * @return Return the regular expression that is equivalent to this
     *         automaton.
     */
    public String getRegex();
}
```

Java Interface definition

```
/usr/bin/java: section
(_TEXT,__text) section
00001c74    pushl  $0x00
00001c76    movl   %esp,%ebp
00001c78    andl   $0xf0,%esp
00001c7b    subl   $0x10,%esp
00001c7e    movl   0x04(%ebp),%ebx
00001c81    movl   %ebx,(%esp)
00001c84    leal   0x08(%ebp),%ecx
00001c87    movl   %ecx,0x04(%esp)
00001c8b    addl   $0x01,%ebx
00001c8e    shll   $0x02,%ebx
00001c91    addl   $0x04,%ebx
00001c93    movl   %ebx,0x08(%esp)
00001c97    movl   (%ebx),%eax
00001c99    addl   $0x04,%ebx
00001c9c    testl  %eax,%eax
00001c9e    jne    0x00001c97
00001ca0    movl   %ebx,0x0c(%esp)
00001ca4    calll  0x00001cb2
00001ca9    movl   %eax,(%esp)
00001cac    calll  0x00002a42
00001cb1    hlt
00001cb2    pushl  %ebp
00001cb3    movl   %esp,%ebp
00001cb5    pushl  %edi
00001cb6    pushl  %esi
00001cb7    pushl  %ebx
```

x86 Assembly

What's the most striking difference?

```
/*
 * A simple interface to describe finite automata. {@link GraphvizFormatter} uses
 * this to visualize arbitrary finite automata.
 */
public interface Automaton {
    /**
     * @return Something that can be iterated to visit all states.
     */
    public Iterable<State> getStates();

    /**
     * @return the start state (which must be, by definition, unique)
     */
    public State getStartState();

    /**
     * @return Something that can be iterated to visit all final states.
     */
    public Iterable<State> getFinalStates();

    /**
     * @return Return the regular expression that is equivalent to this
     *         automaton.
     */
    public String getRegex();
}
```

Java Interface definition

Names!
High-level languages have rich facilities for naming “**things**.”

```
/usr/bin/java: section
(_TEXT, __text) section
00001c74 pushl $0x00
00001c76 movl %esp,%ebp
00001c78 andl $0xf0,%esp
00001c7b subl $0x10,%esp
00001c7e movl 0x04(%ebp),%ebx
00001c81 movl %ebx,(%esp)
00001c84 leal 0x08(%ebp),%ecx
00001c87 movl %ecx,0x04(%esp)
00001c8b addl $0x01,%ebx
00001c8e shll $0x02,%ebx
00001c91 addl $0x04,%ebx
00001c93 movl %ebx,0x08(%esp)
00001c97 movl (%ebx),%eax
00001c99 addl $0x04,%ebx
00001c9c testl %eax,%eax
00001c9e jne 0x00001c97
00001ca0 mo 0x00001cb2,%esp
00001ca4
00001ca9 movl %eax,(%esp)
00001caa calll 0x00002a42
00001ca1 hlt %ebp
00001cb2 pushl %esp,%ebp
00001cb3 movl %edi
00001cb5 pushl %esi
00001cb6 pushl %ebx
```

x86 Assembly

Names

Required for abstraction.

- Assembly only has **values** & **addresses** & registers.
 - Machine dependence!
- **Names enable abstraction.**
 - Can refer to something **without knowing the details** (e.g., exact address, exact memory layout).
 - Let the compiler worry about the details.
- Can refer to things that **do not yet exist!**
 - E.g., during development, we can (and often do) write code for (Java) interfaces that have not yet been implemented.

Abstraction

Control Abstraction vs. Data Abstraction

```
/**  
 * A simple interface to describe finite automata. {@link GraphvizFormatter} uses  
 * this to visualize arbitrary finite automata.  
 */  
public interface Automaton {  
    /**  
     * @return Something that can be iterated to visit all states.  
     */  
    public Iterable<State> getStates();  
  
    /**  
     * @return the start state (which must be, by definition, unique)  
     */  
    public State getStartState();  
  
    /**  
     * @return Something that can be iterated to visit all final states.  
     */  
    public Iterable<State> getFinalStates();  
  
    /**  
     * @return Return the regular expression that is equivalent to this  
     *         automaton.  
     */  
    public String getRegex();  
}
```

Abstraction

Control Abstraction

Can **hide arbitrary complex code** behind a simple name.
For example, addition can be simple (int) or “difficult” (vector).

```
public interface Automaton {  
    /**  
     * @return Something that can be iterated to visit all states.  
     */  
    public Iterable<State> getStates();  
  
    /**  
     * @return the start state (which must be, by definition, unique)  
     */  
    public State getStartState();  
  
    /**  
     * @return Something that can be iterated to visit all final states.  
     */  
    public Iterable<State> getFinalStates();  
  
    /**  
     * @return Return the regular expression that is equivalent to this  
     *         automaton.  
     */  
    public String getRegex();  
}
```

Abstraction

Control Abstraction vs. Data Abstraction

```
/**  
 * A simple interface to describe finite automata. {@link GraphvizFormatter} uses  
 * this to visualize arbitrary finite automata.  
 */  
public interface Automaton {  
    /**  
     * @return Something that can be iterated  
     */  
    public Iterable<State> getStates();  
  
    /**  
     * @return the start state (which must be,  
     */  
    public State getStartState();  
  
    /**  
     * @return Something that can be iterated  
     */  
    public Iterable<State> getFinalStates();  
  
    /**  
     * @return Return the regular expression that  
     *         defines this automaton.  
     */  
    public String getRegex();  
}
```

Data Abstraction

Abstract Data Types (ADTs)

Reason about **concepts** instead of **impl. details.**

Programmer doesn't know memory layout, address, whether other interfaces are implemented, what invariants need to be ensured, etc.

Binding

Associating a **name** with some **entity**.
(or “object,” but not the Java notion of an object)

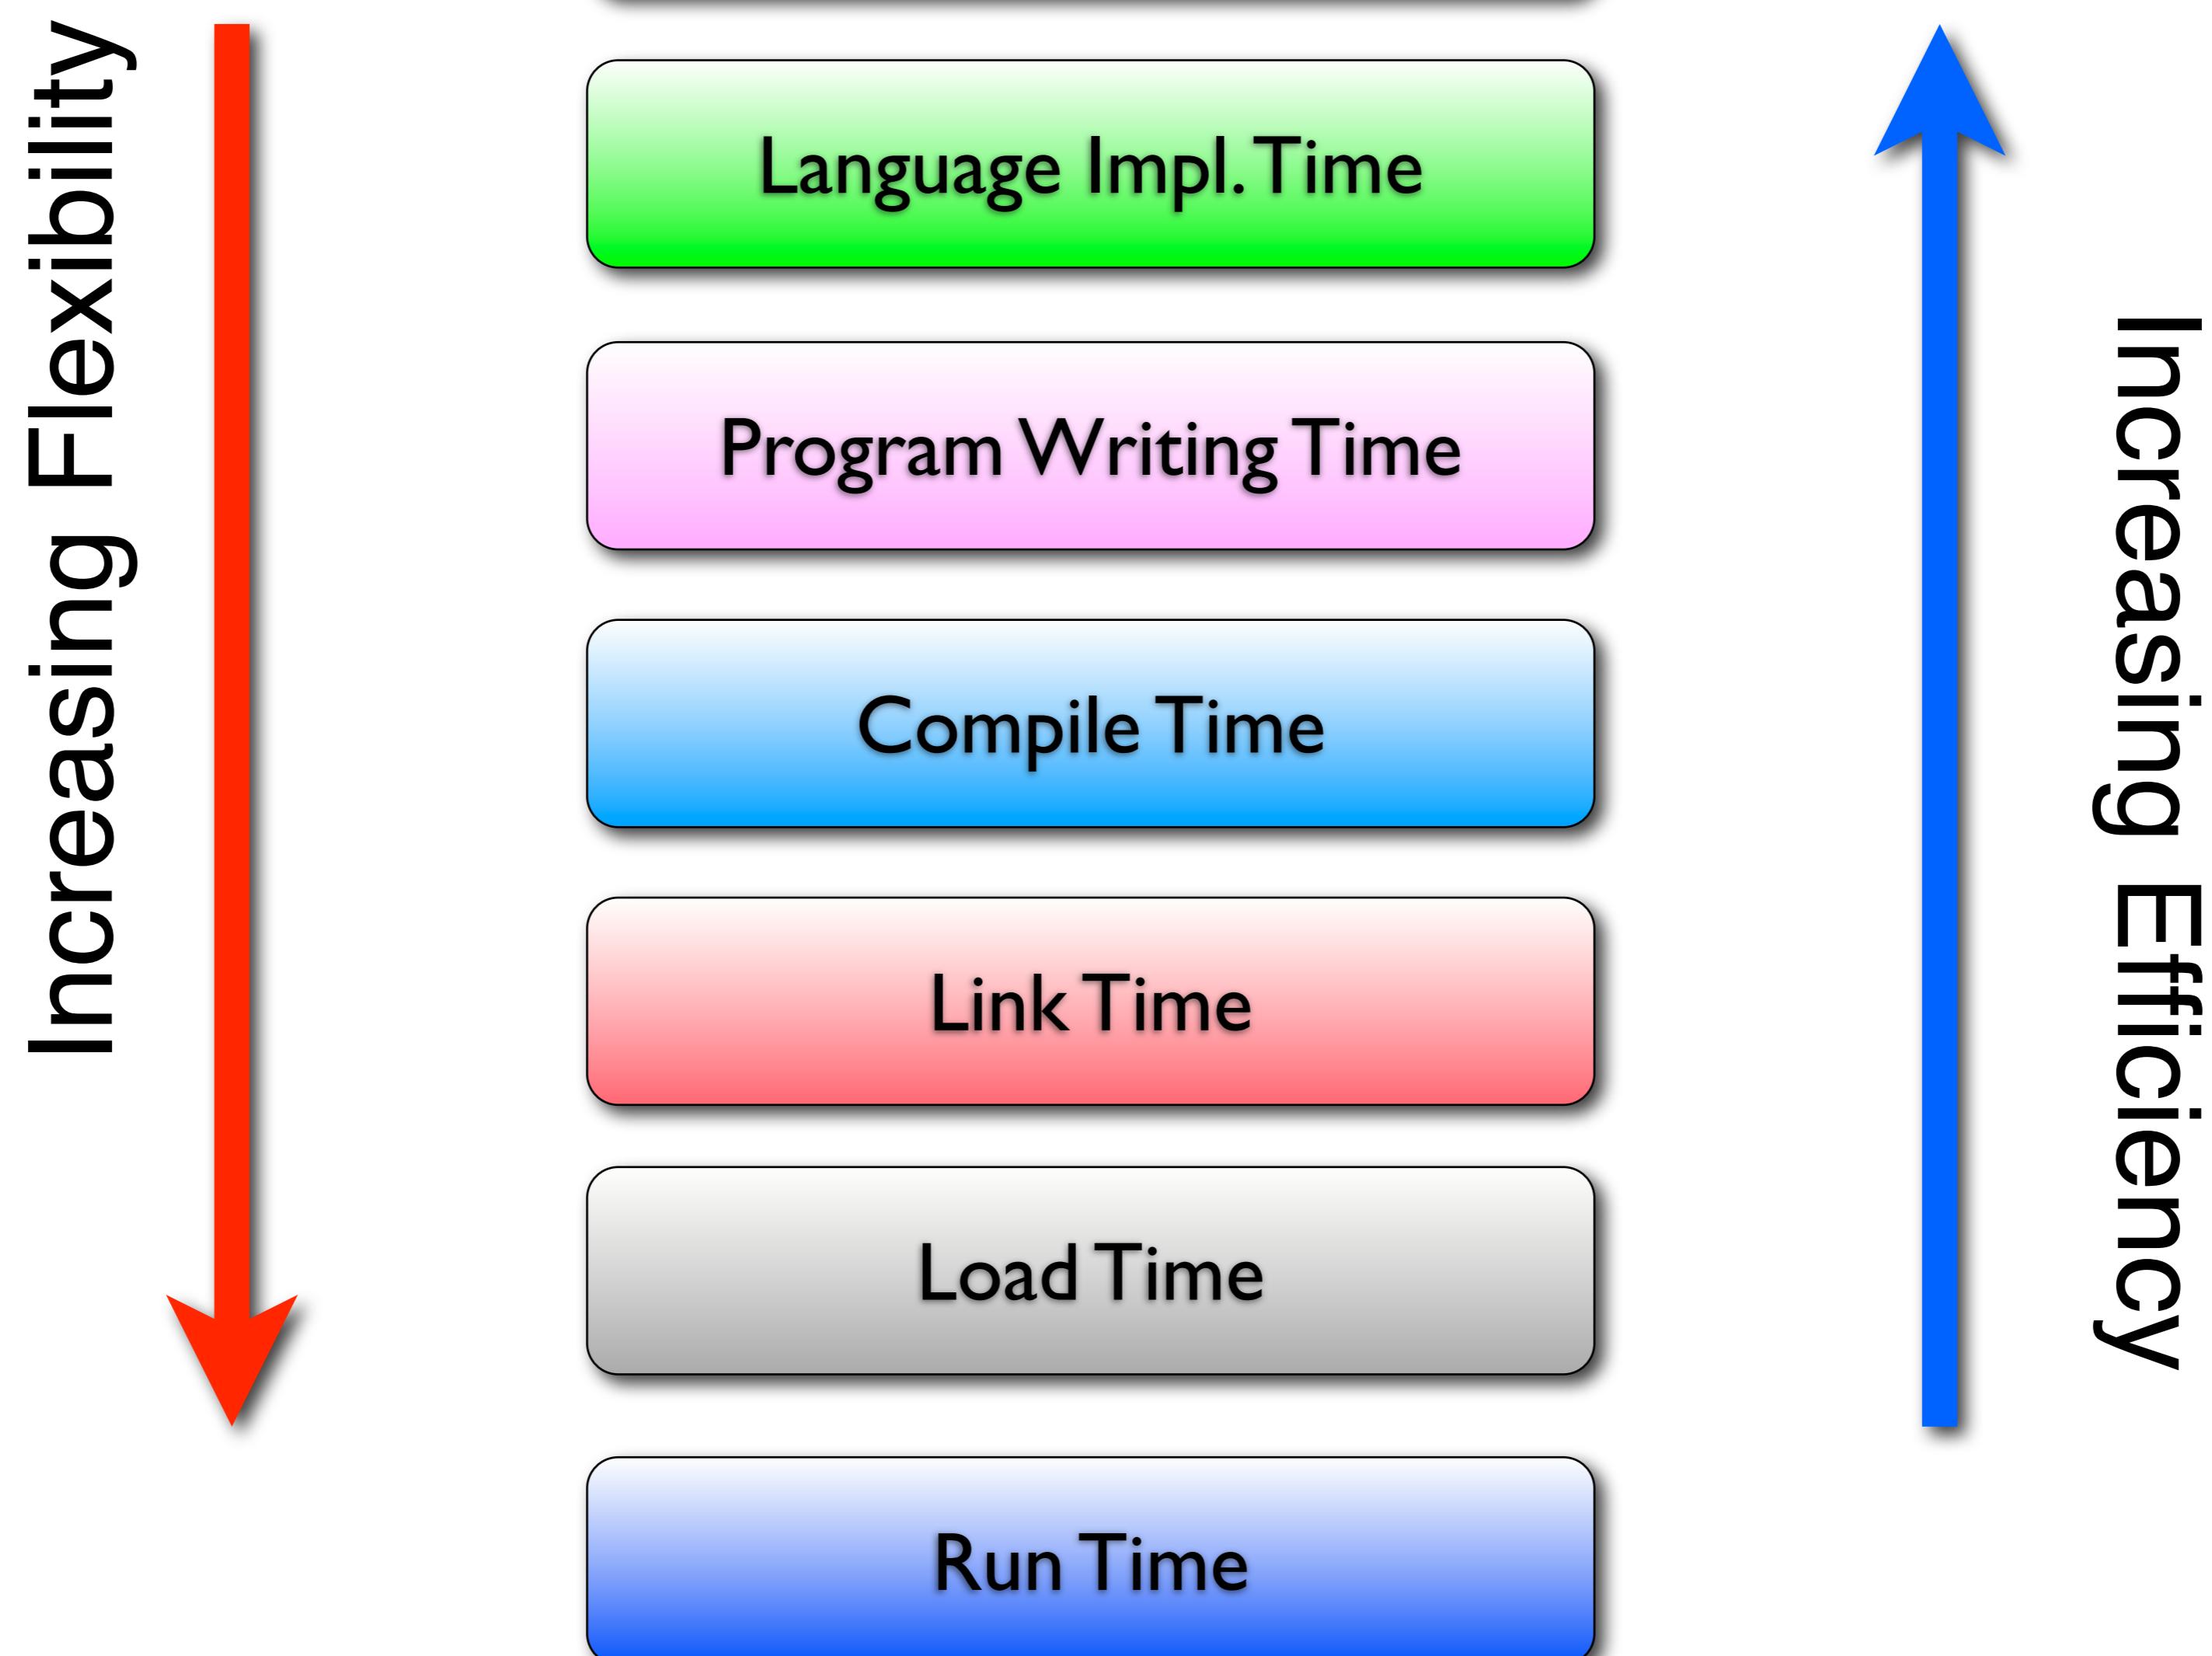
Binding vs. Abstraction.

- Introducing a **name creates an abstraction**.
- **Binding** a name to an entity **resolves an abstraction**.

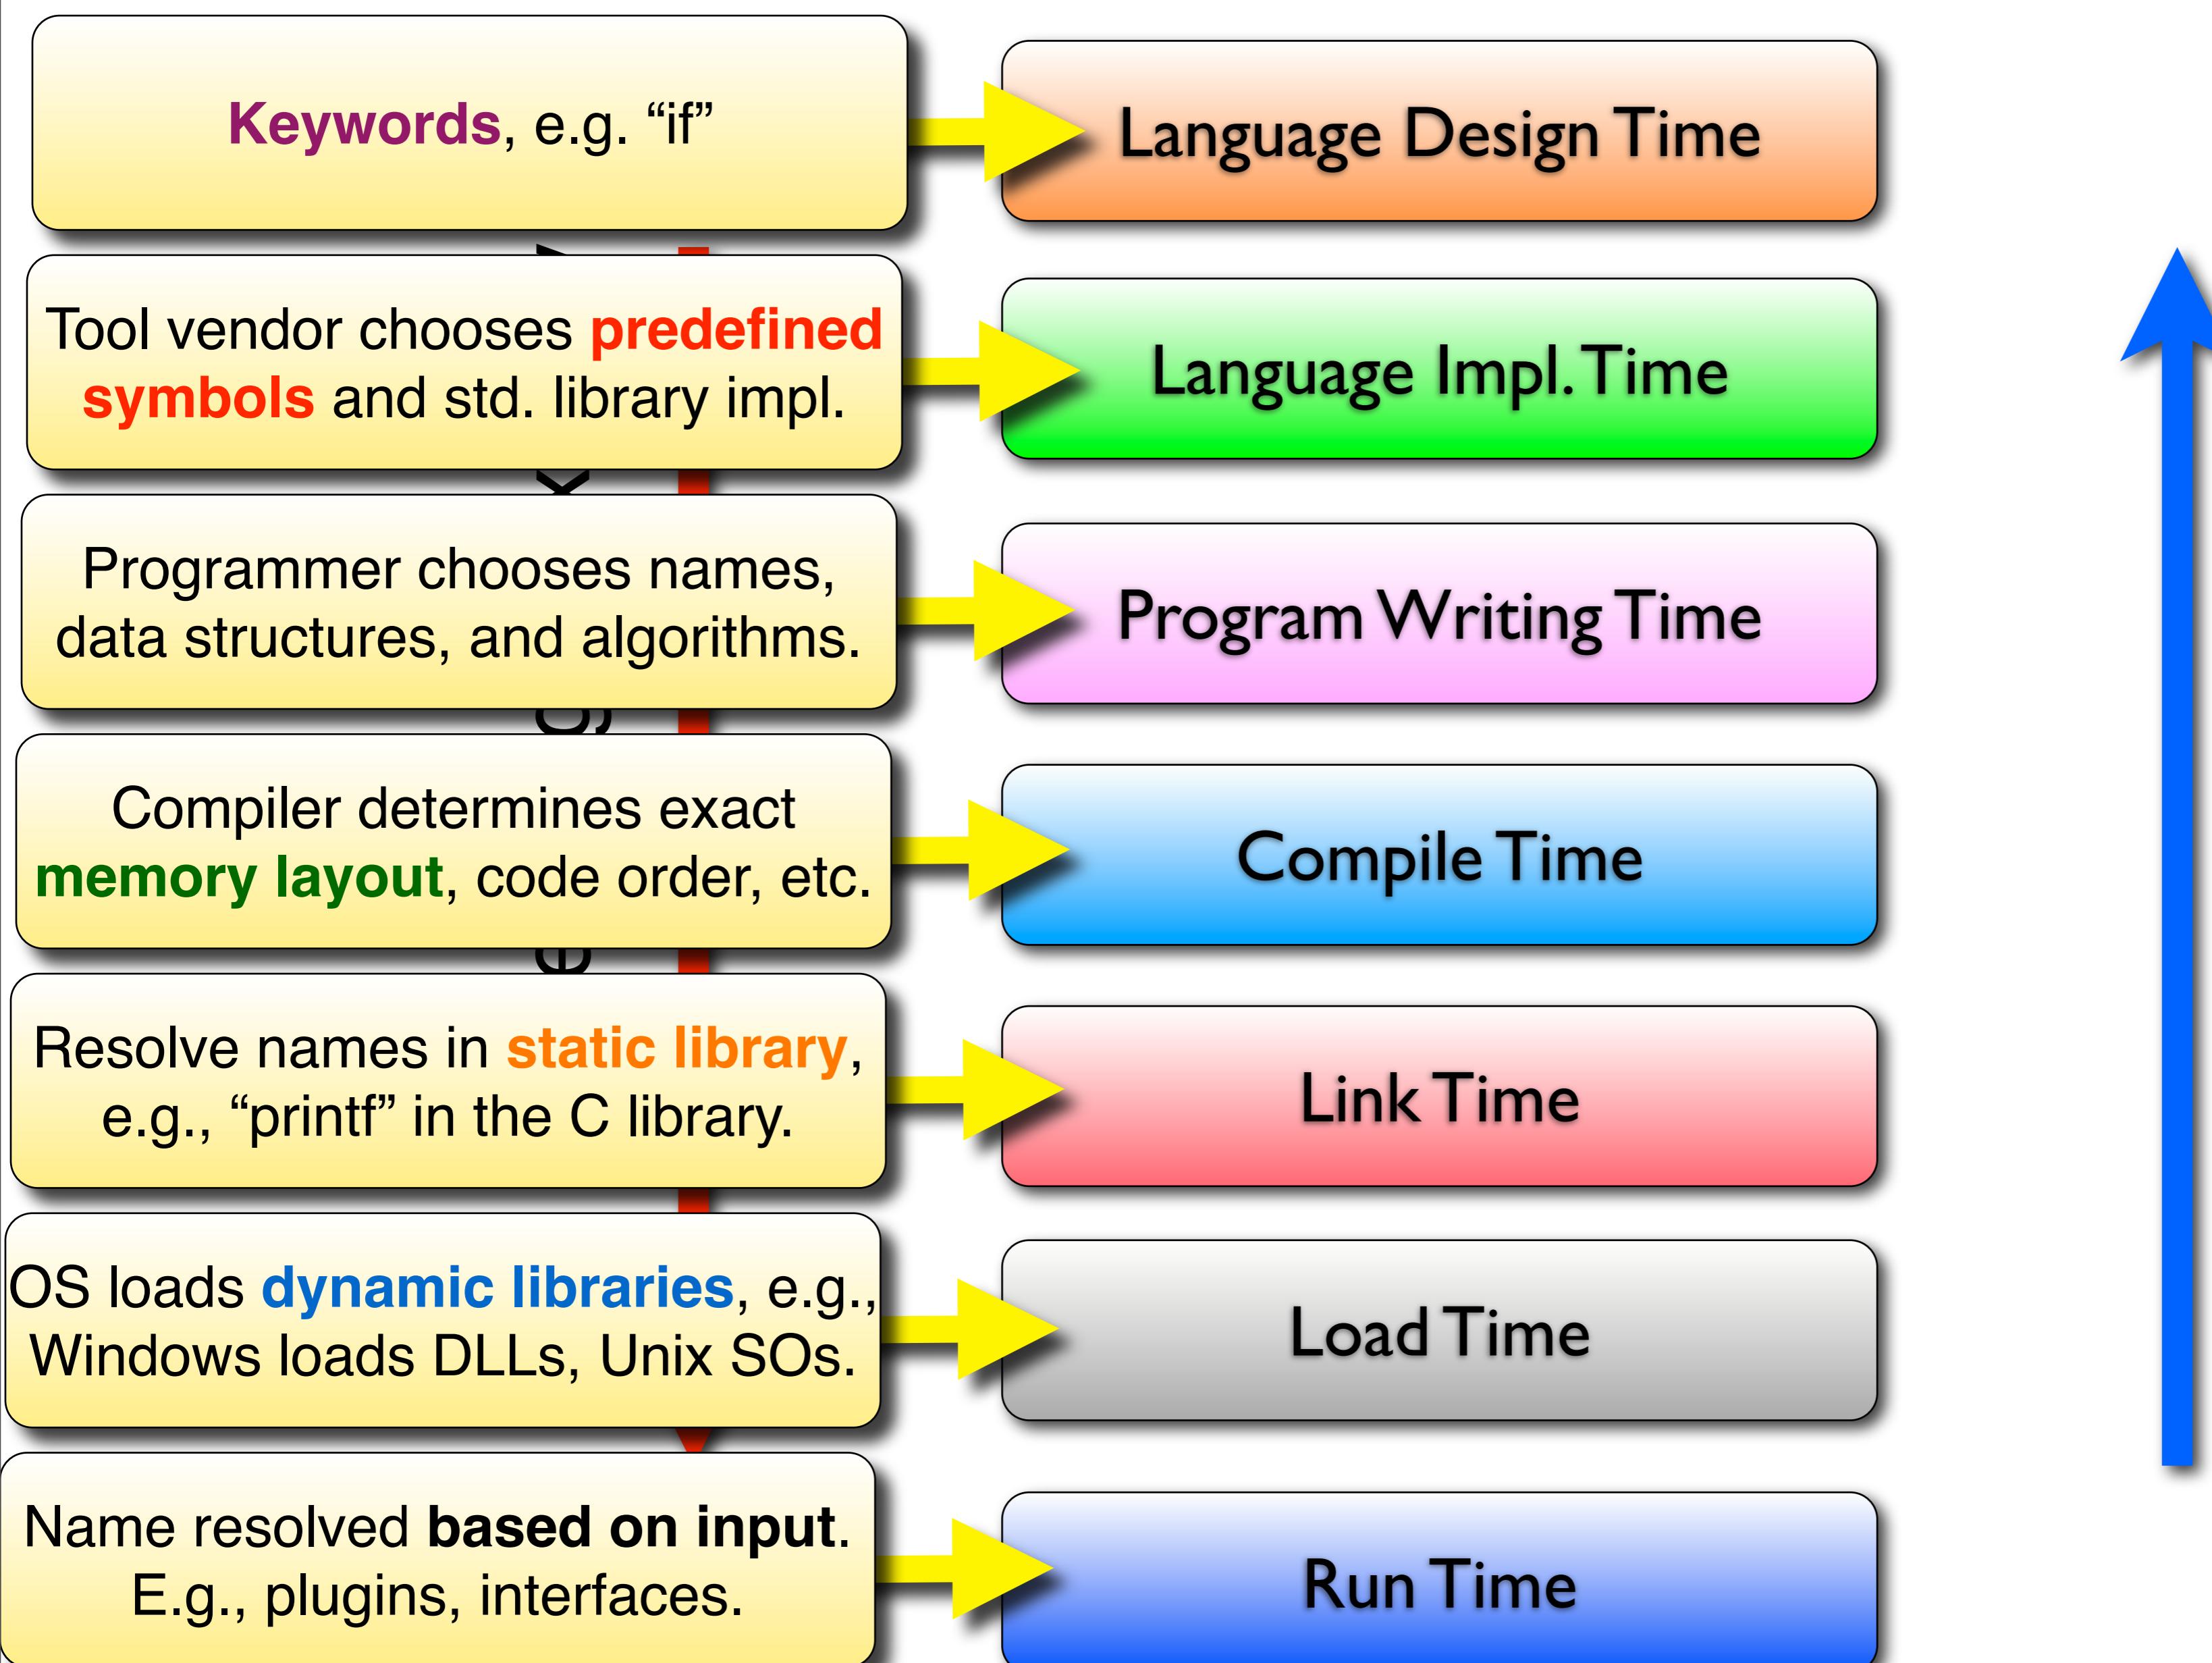
Binding time:

- When is a name resolved?

Binding Time



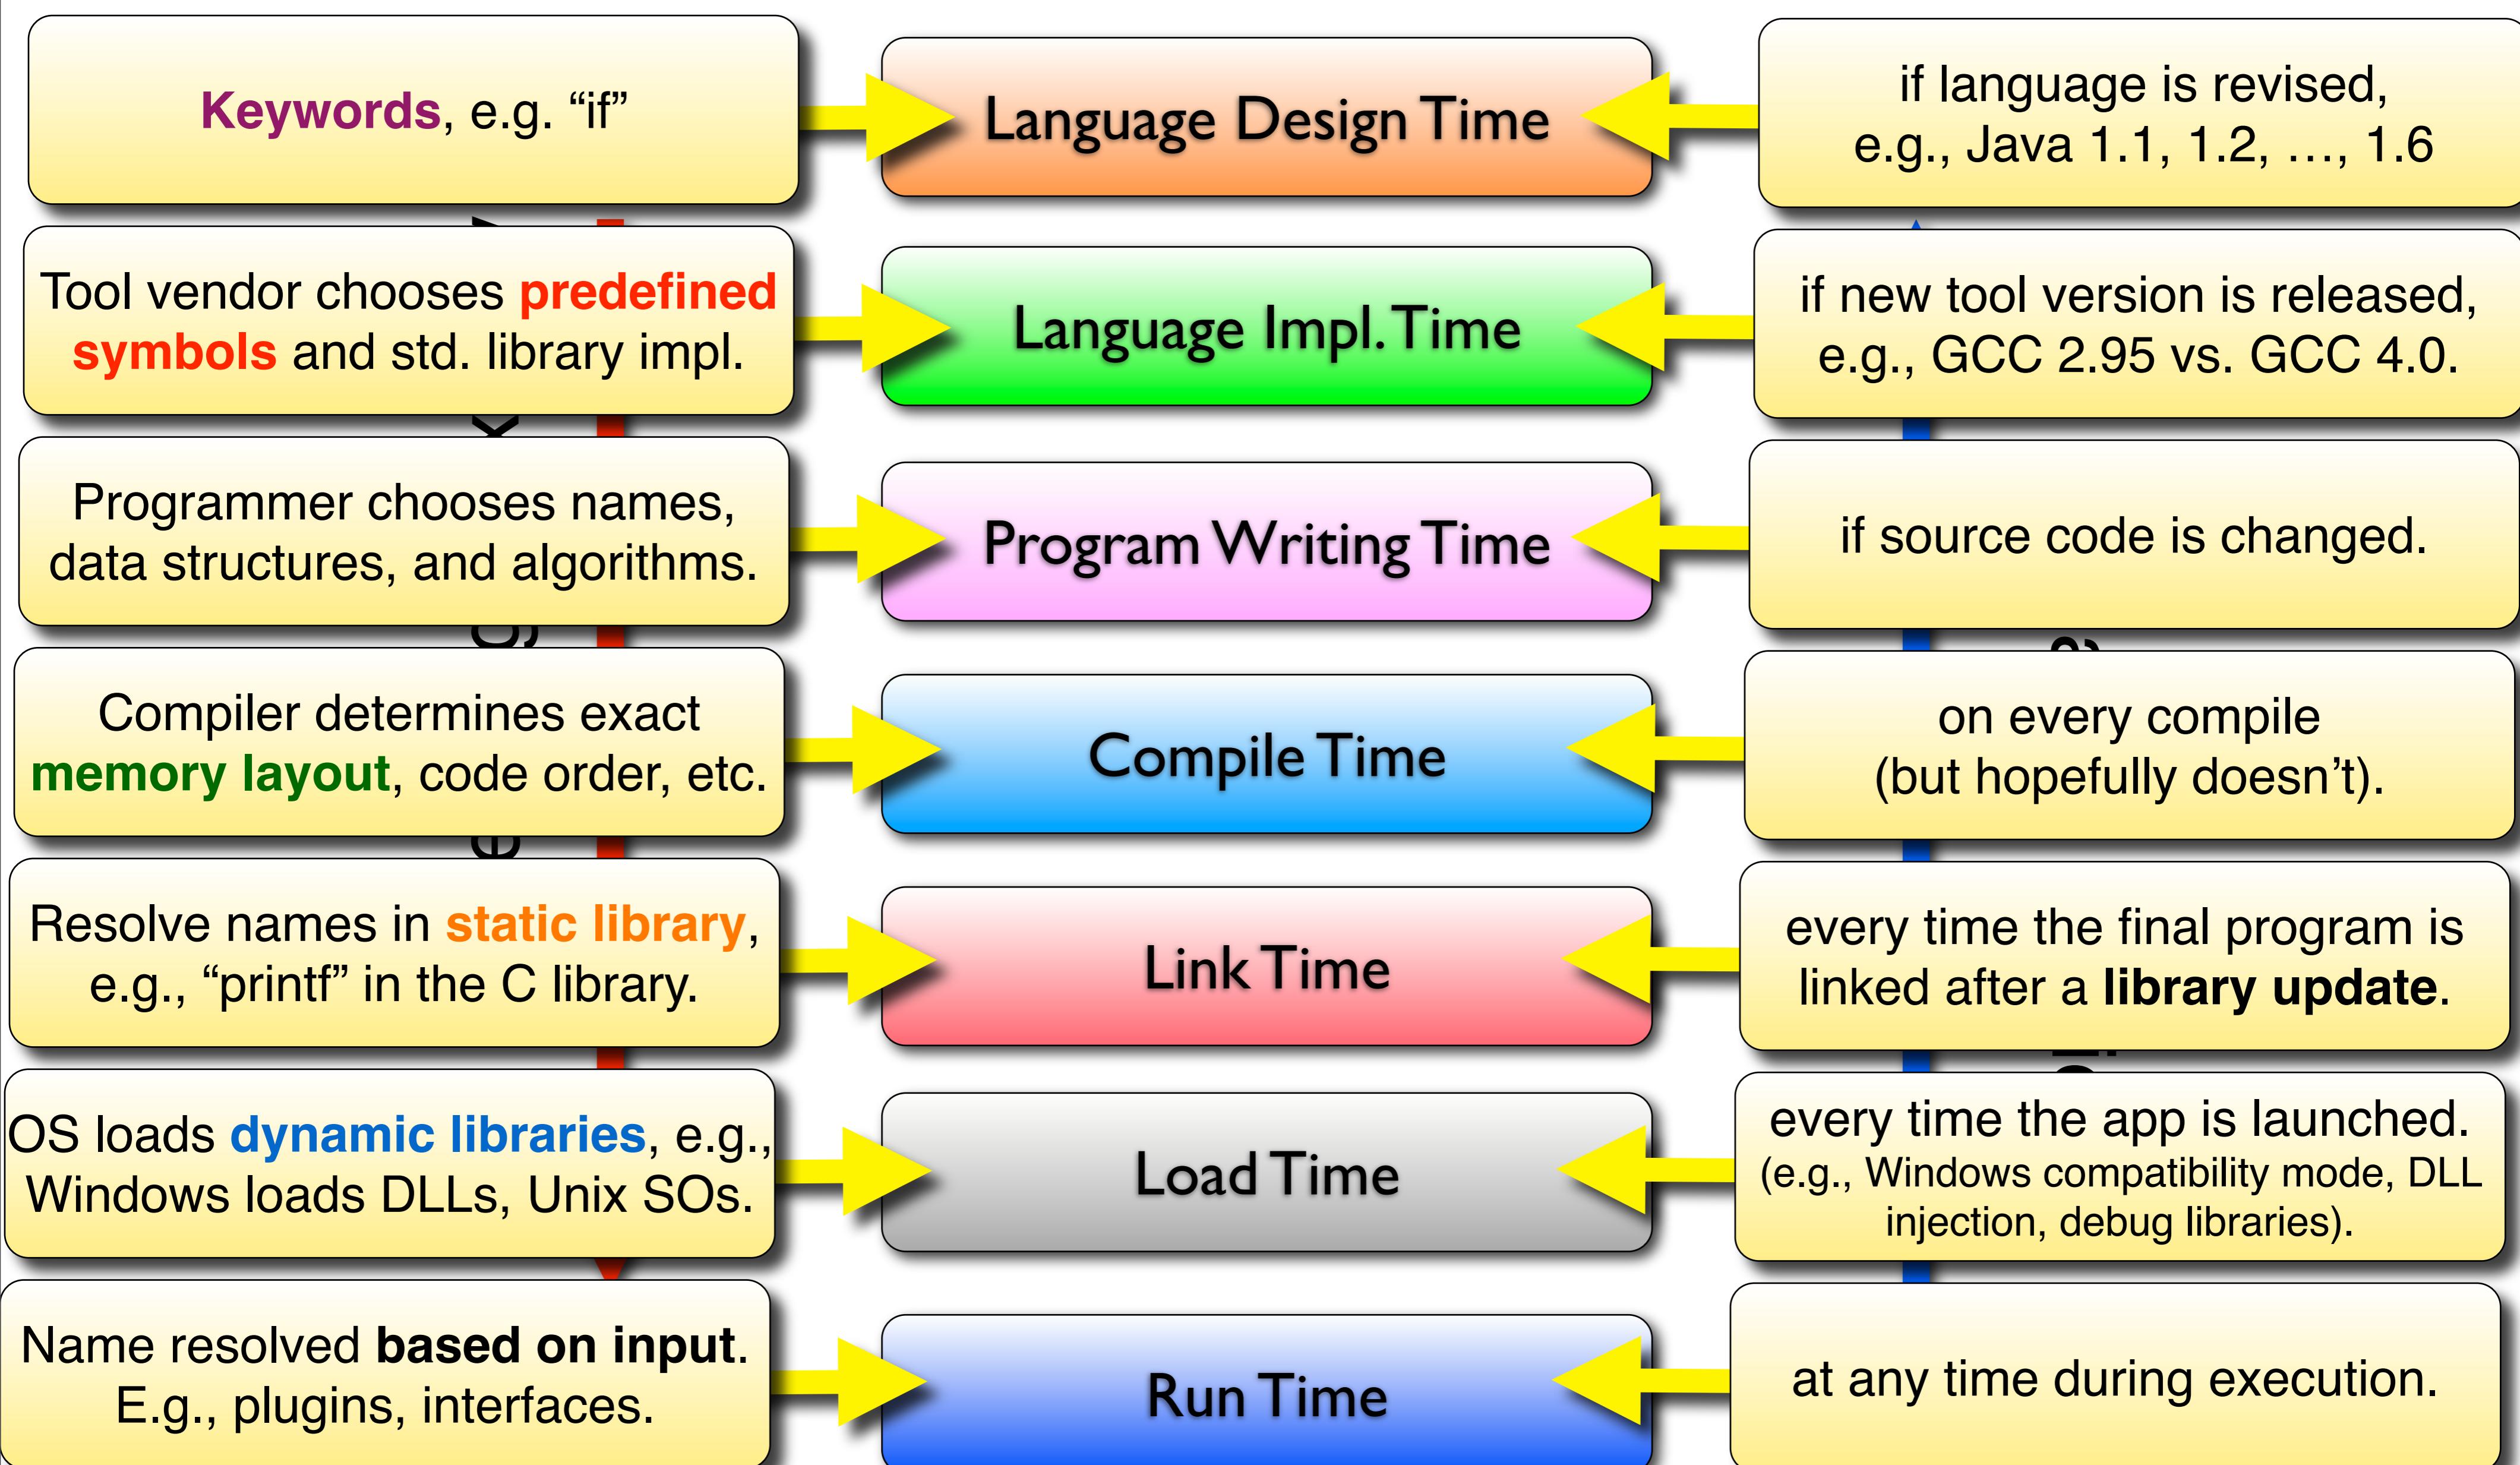
Binding Time



Increasing Efficiency

Binding Time

Binding may change...



Binding Time

Called **static** or **early** binding.

Increasing Flexibility



Language Design Time

Language Impl. Time

Program Writing Time

Compile Time

Link Time

Load Time

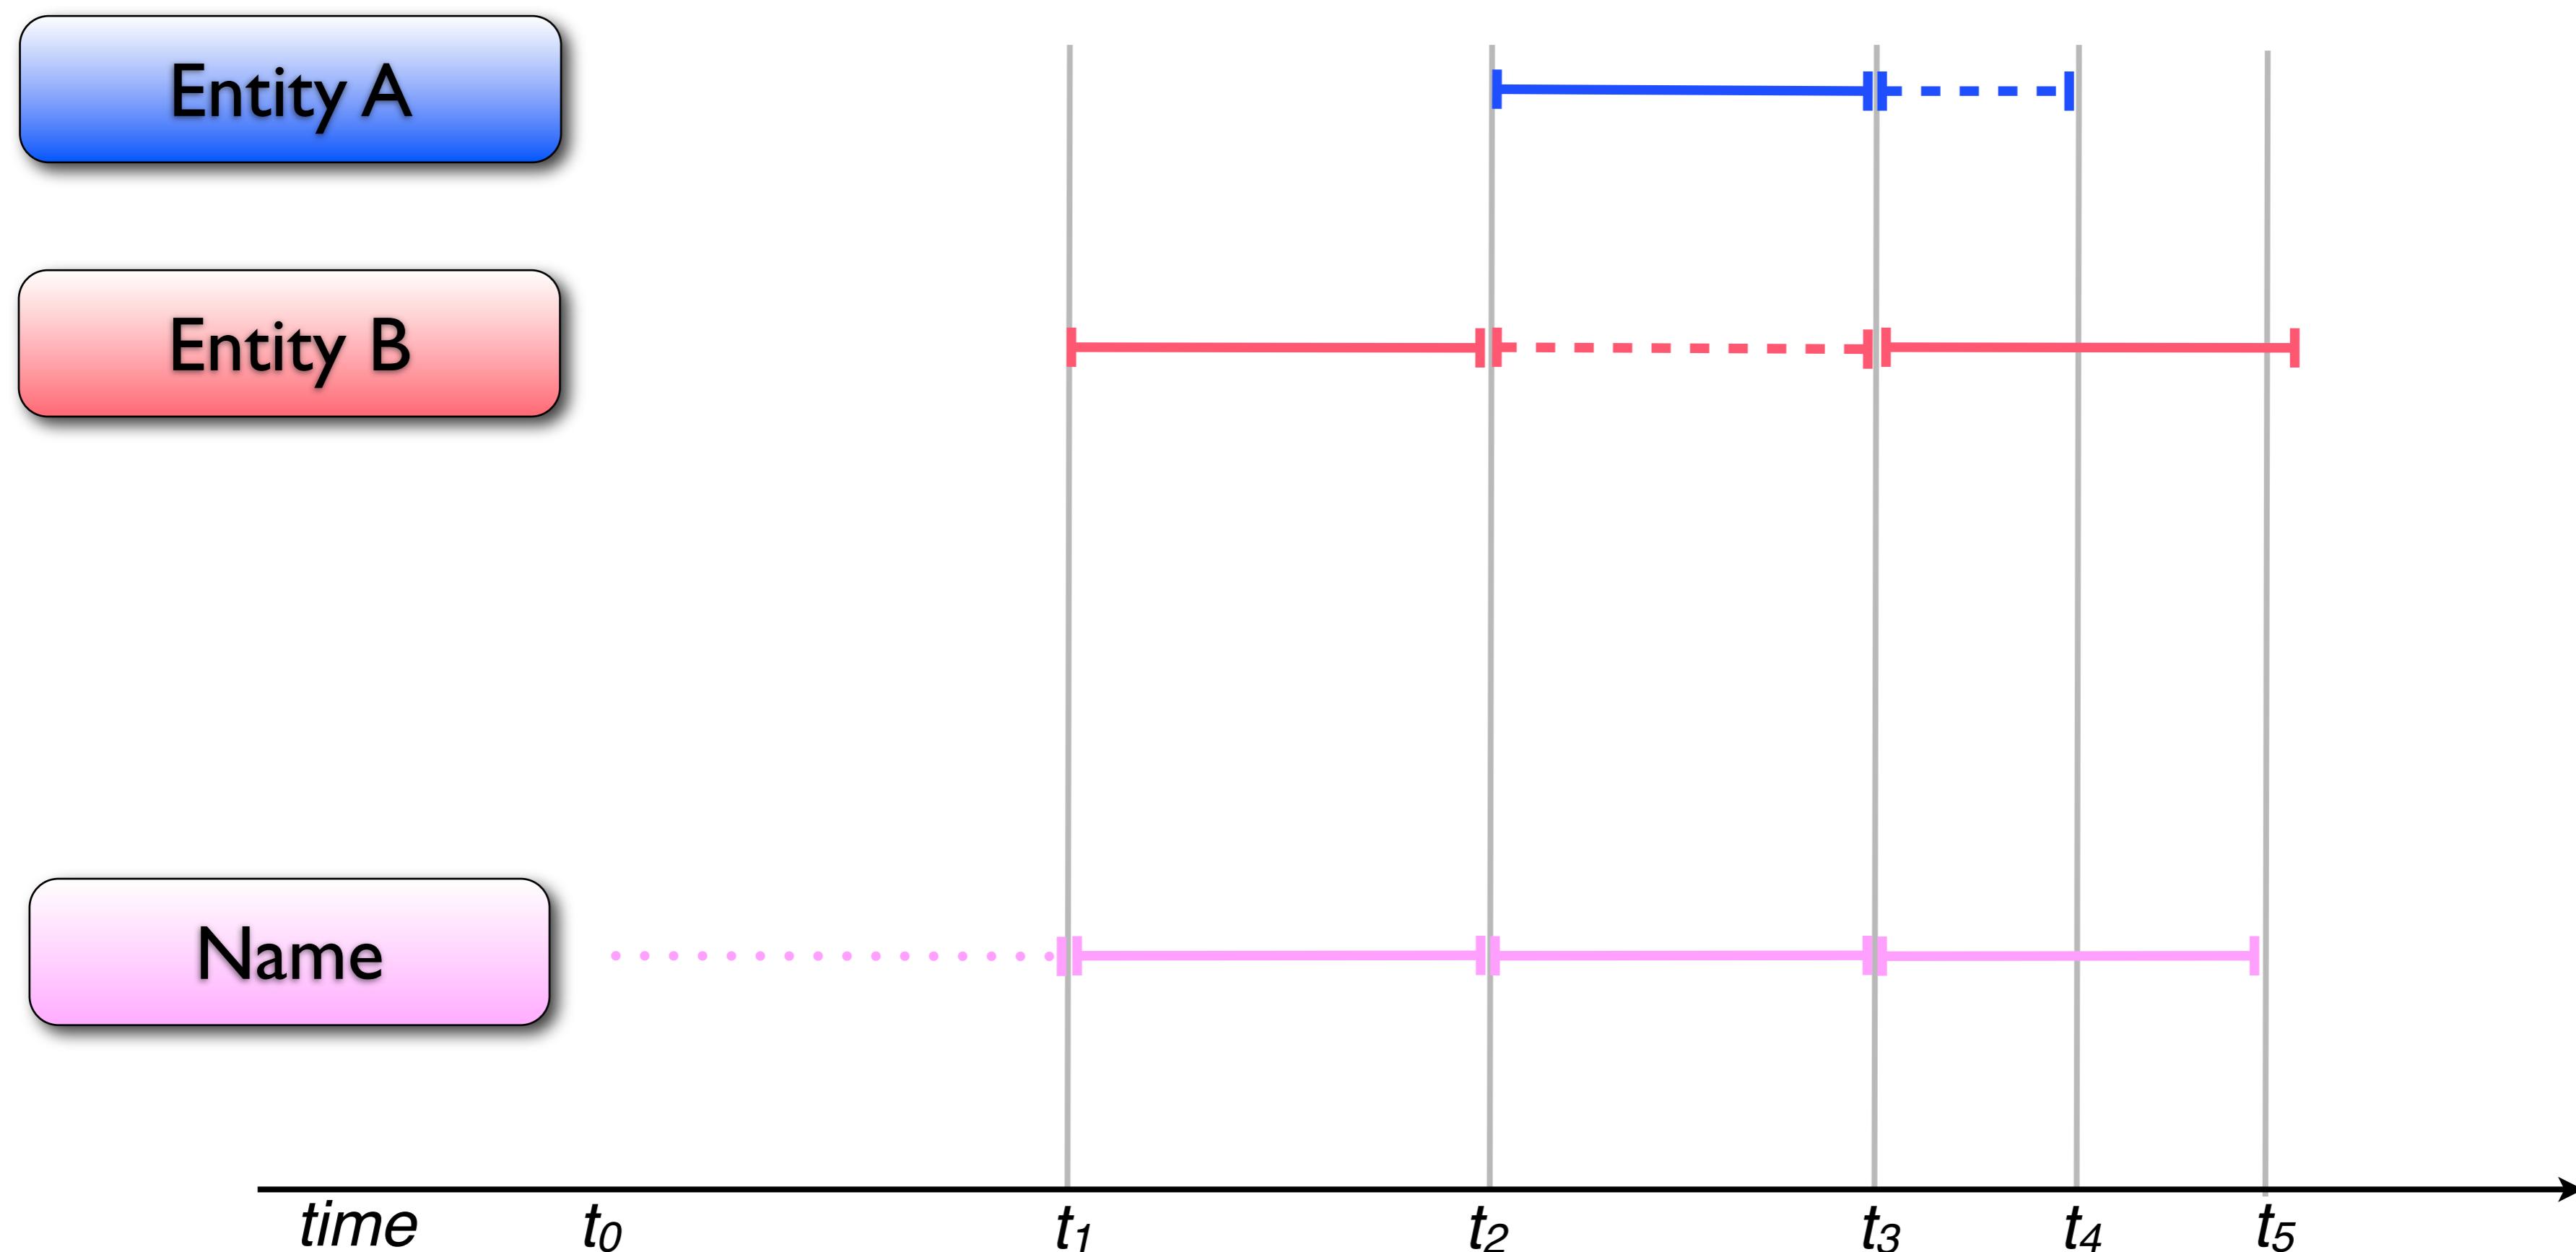
Run Time



Increasing Efficiency

Called **dynamic** or **late** binding.

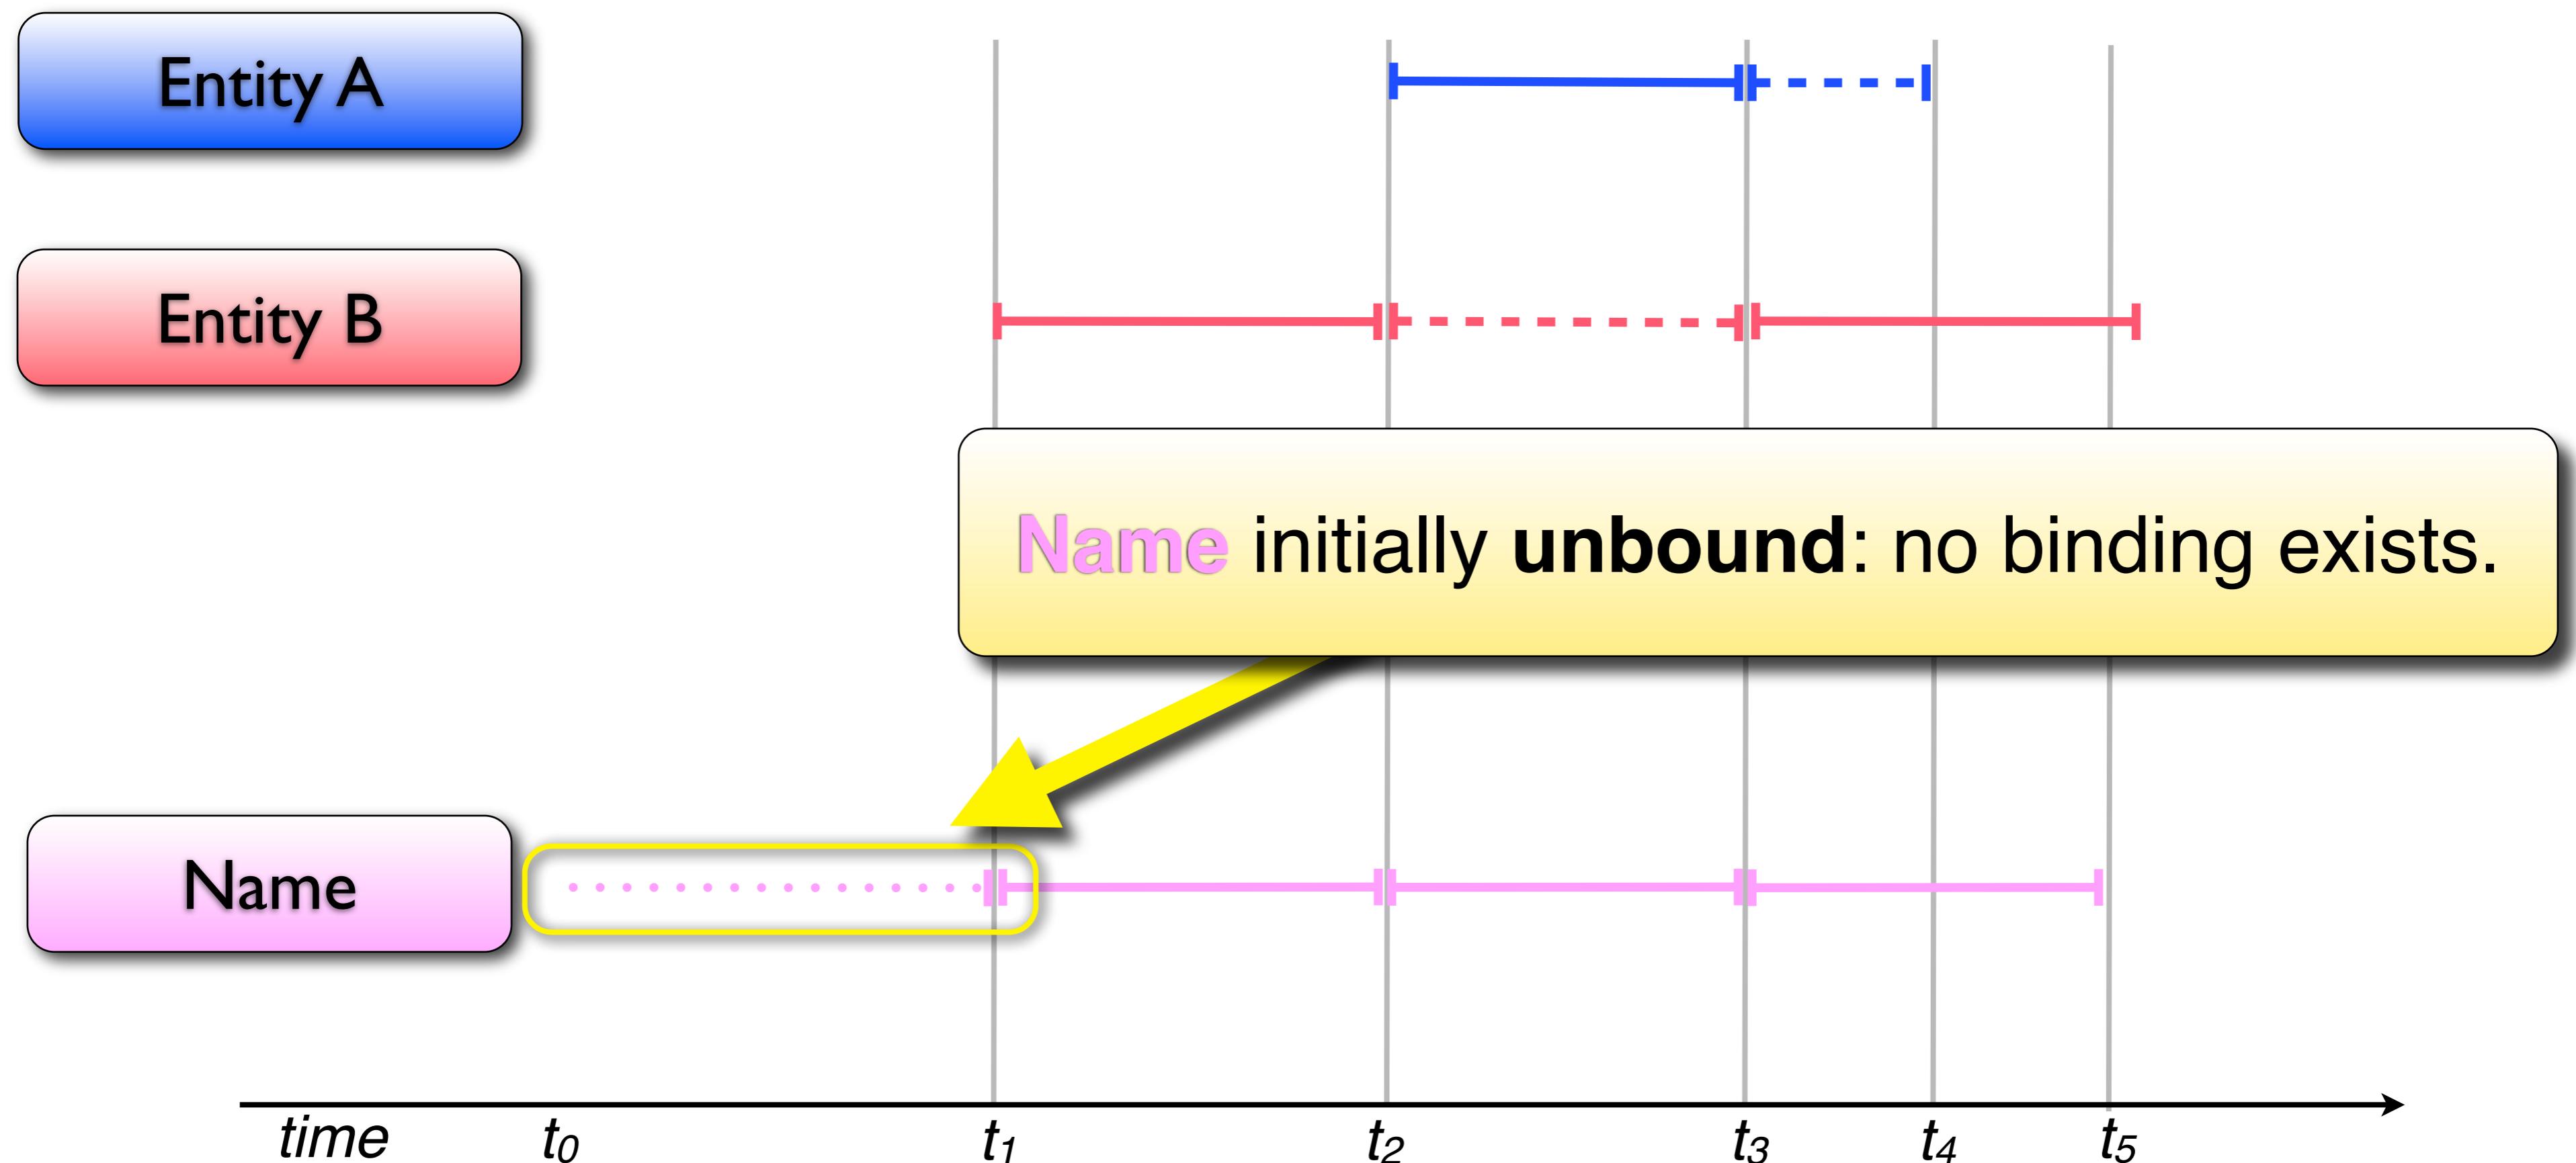
Object Lifetime vs. Binding Lifetime



Lifetime

- Entity: “alive” if memory is allocated (and initialized).
- Binding: “alive” if the name refers to some entity.

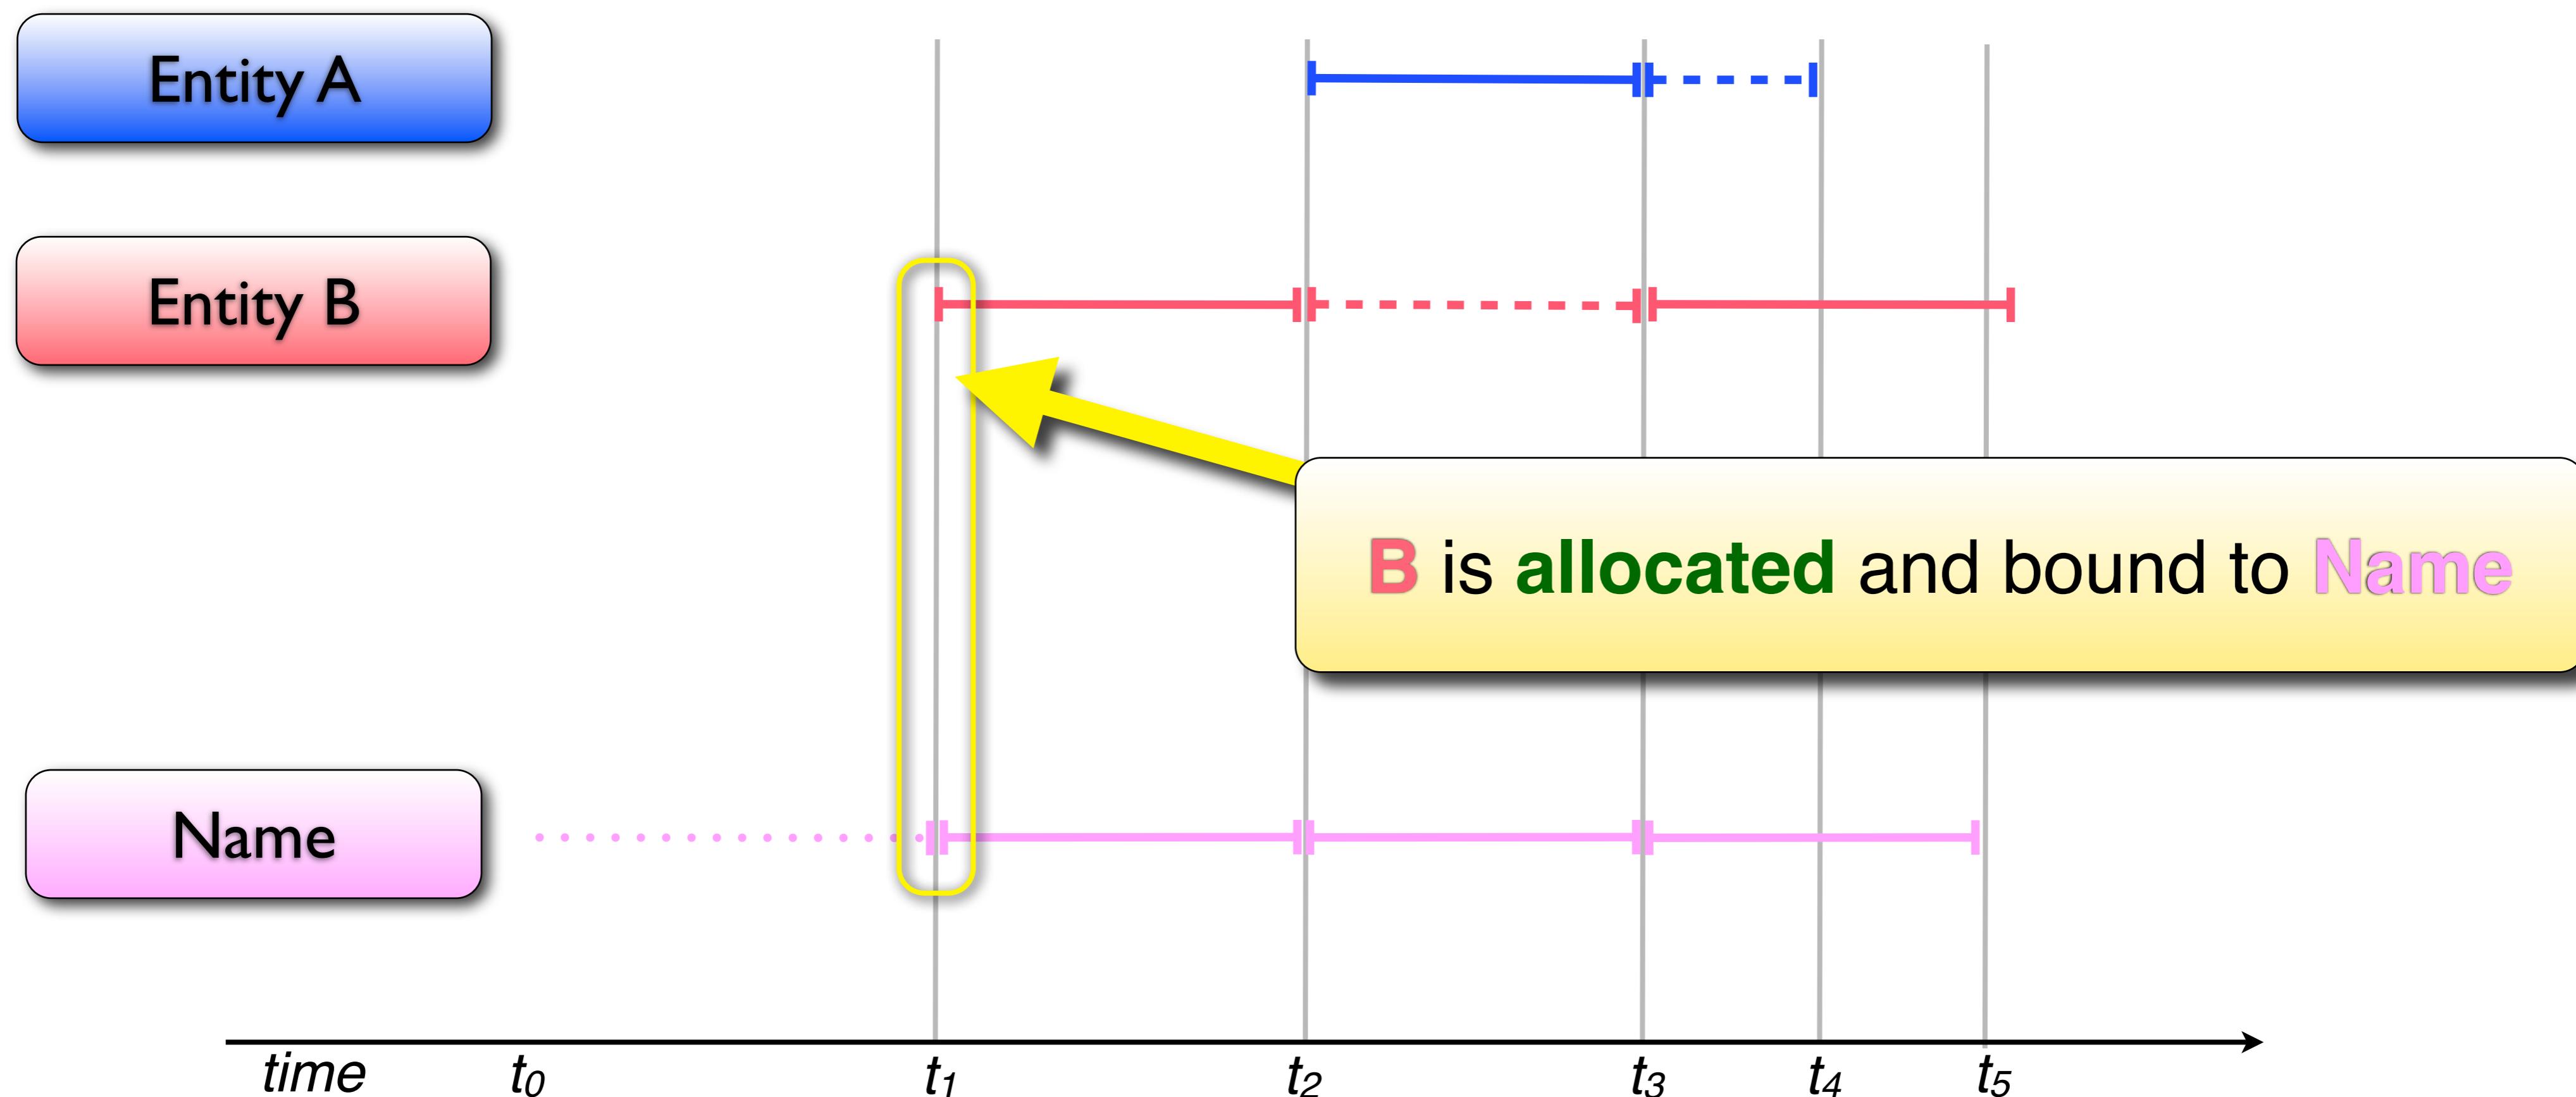
Object Lifetime vs. Binding Lifetime



Lifetime

- Entity: “alive” if memory is allocated (and initialized).
- Binding: “alive” if the name refers to some entity.

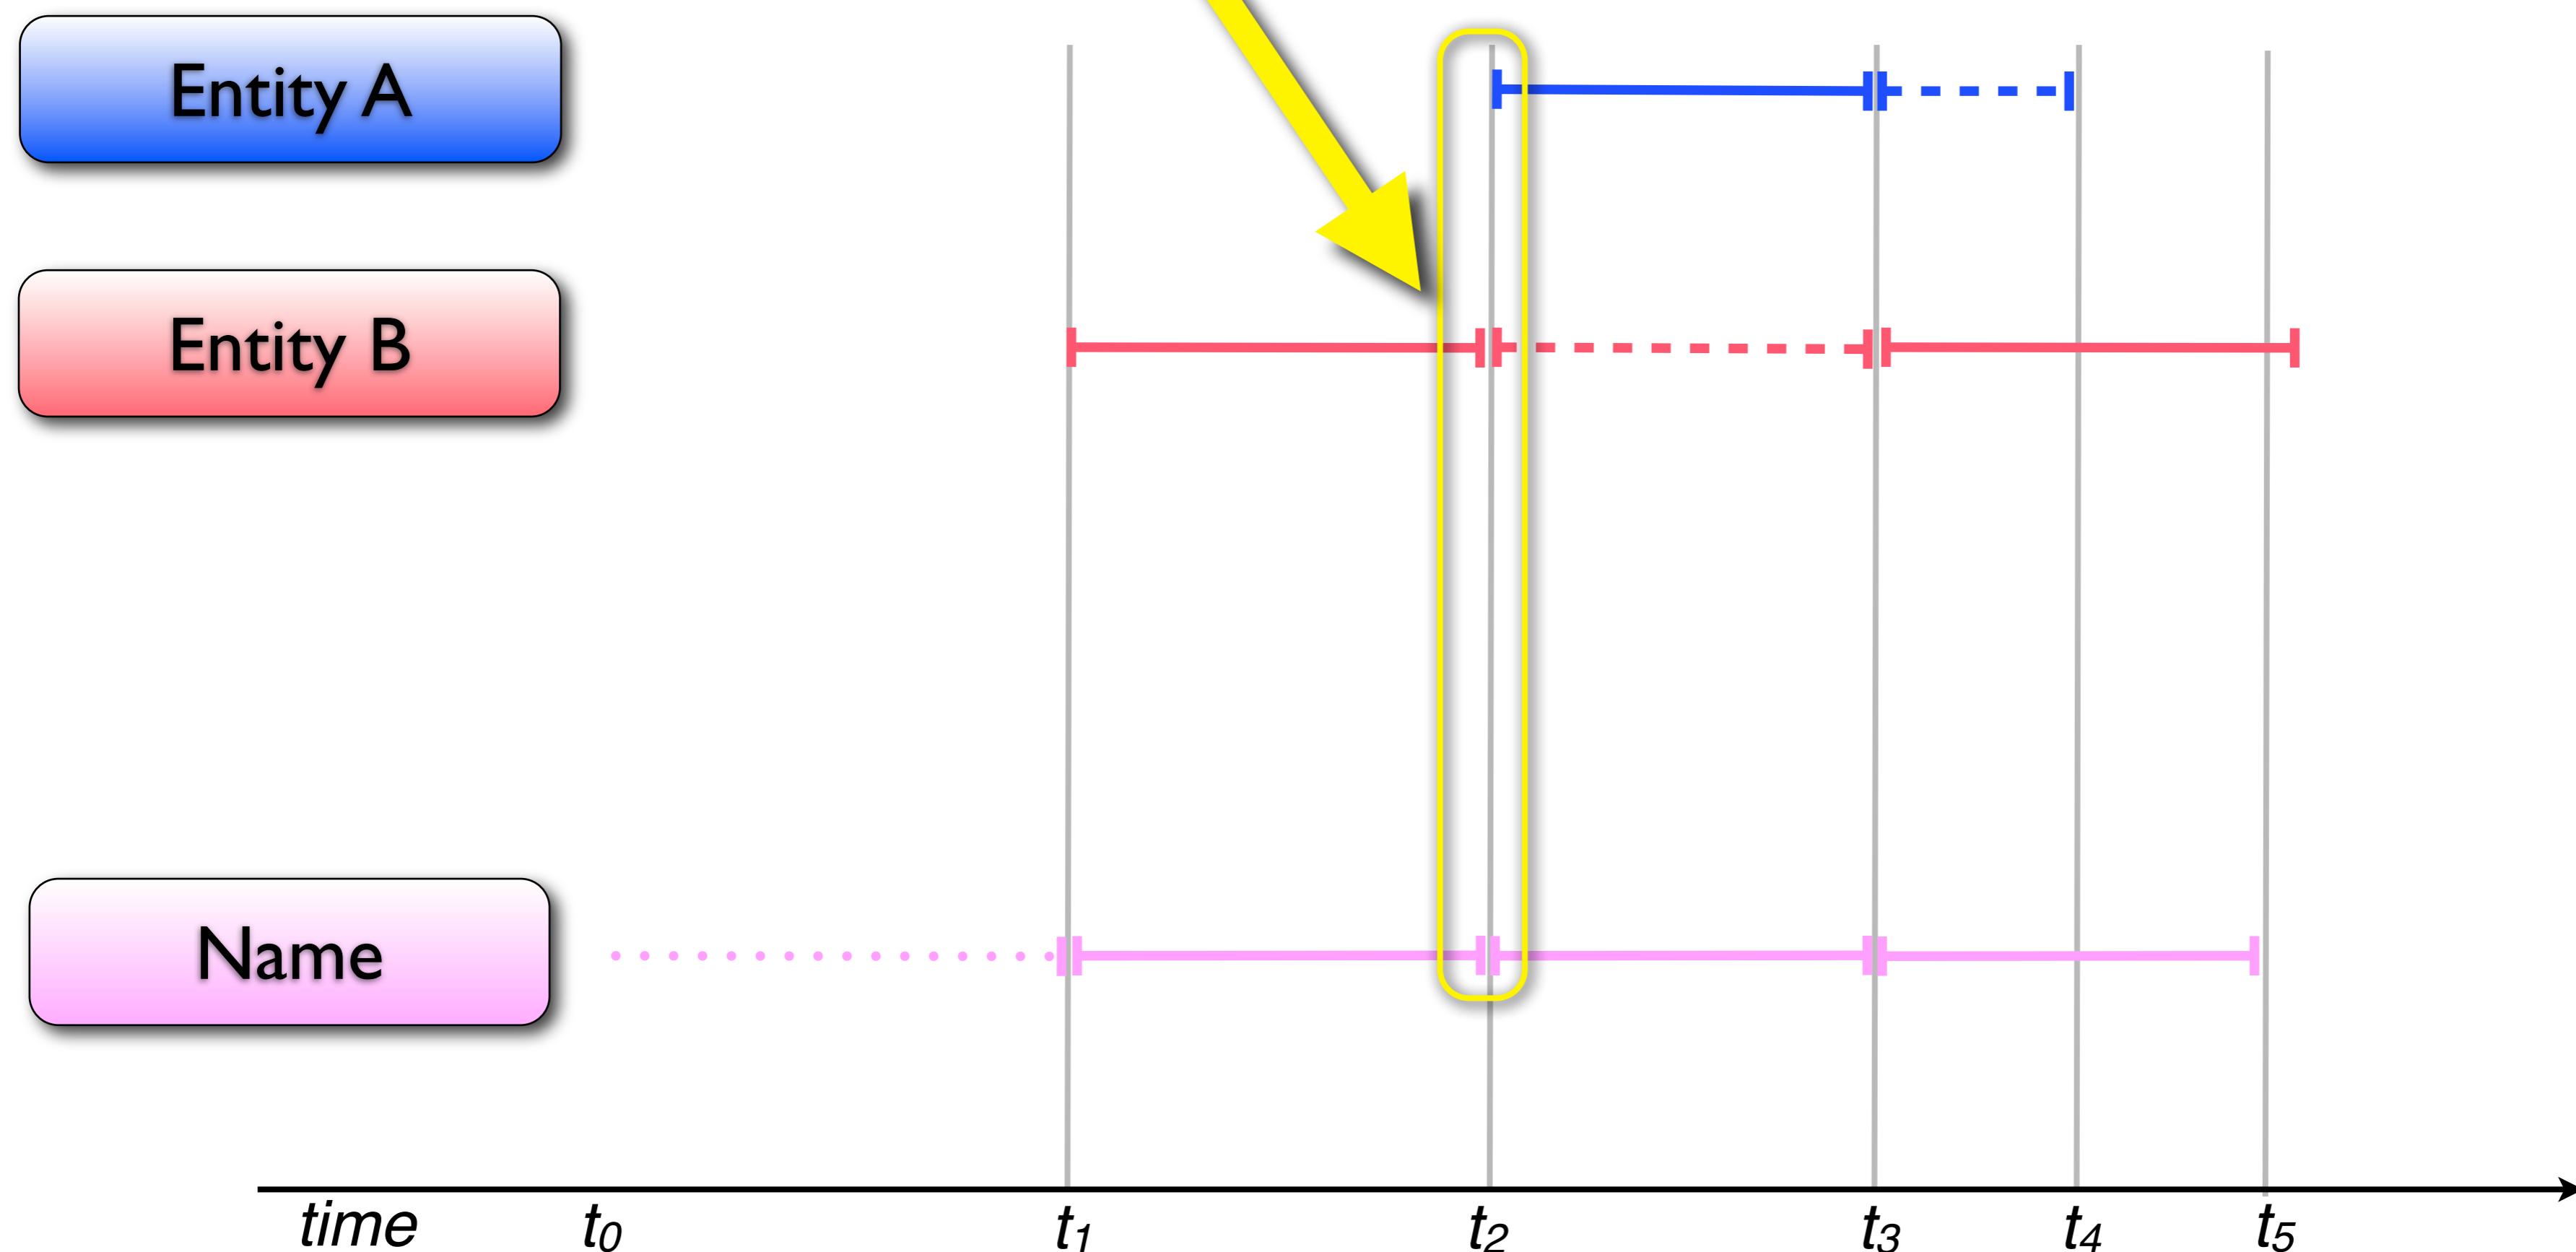
Object Lifetime vs. Binding Lifetime



Lifetime

- Entity: “alive” if memory is allocated (and initialized).
- Binding: “alive” if the name refers to some entity.

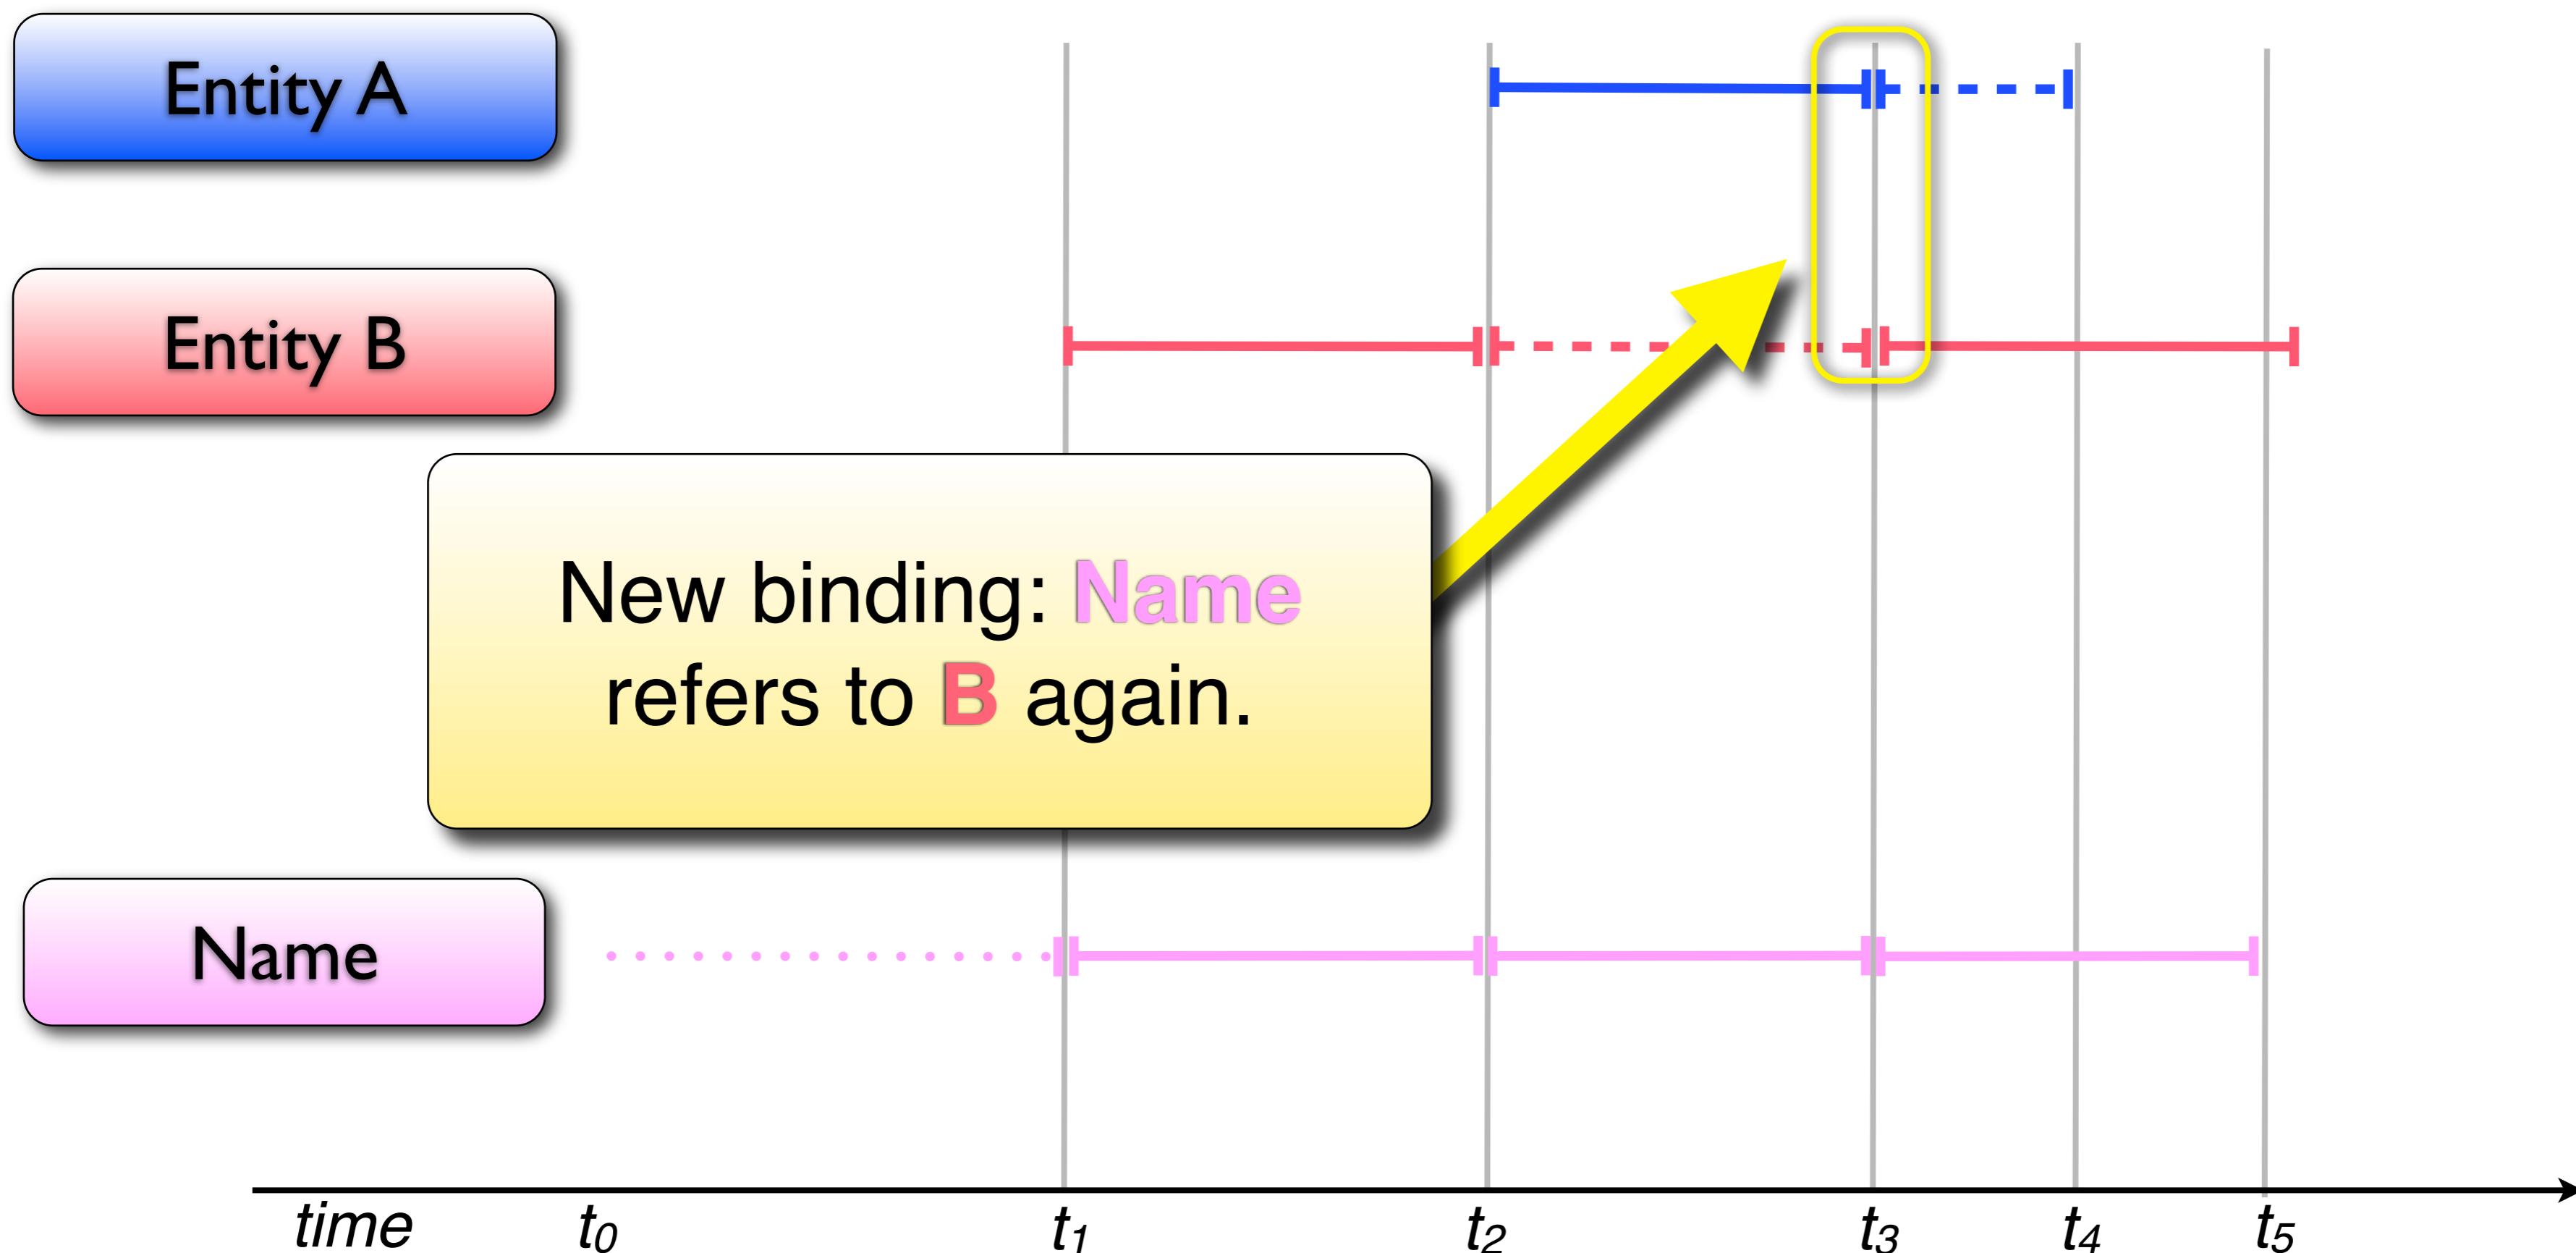
Object Binding Lifetime



Lifetime

- Entity: “alive” if memory is allocated (and initialized).
- Binding: “alive” if the name refers to some entity.

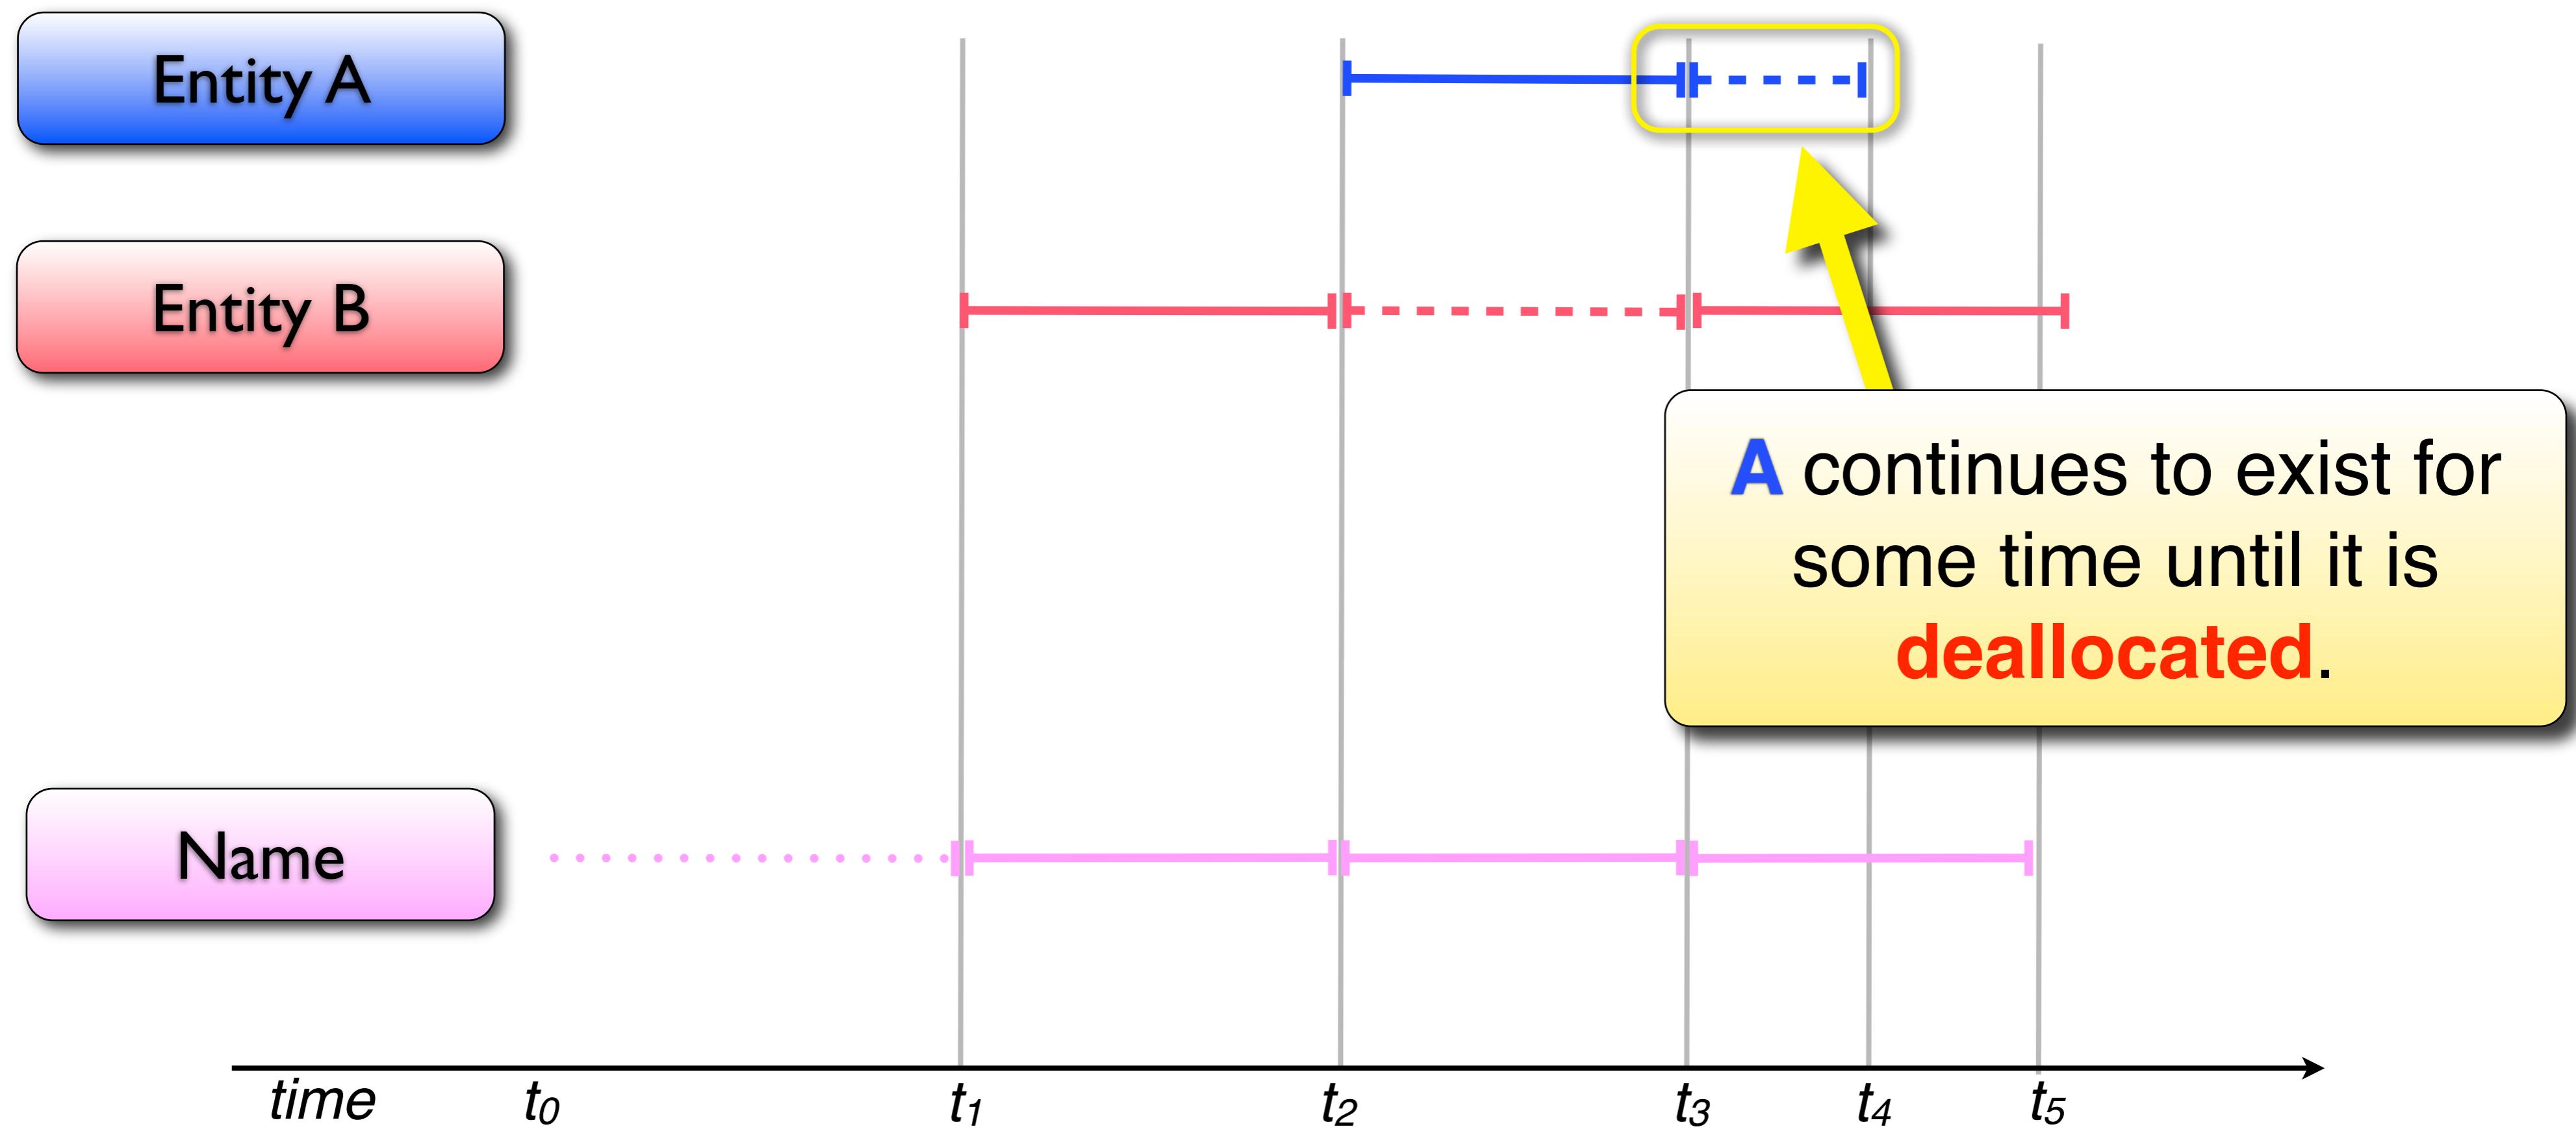
Object Lifetime vs. Binding Lifetime



Lifetime

- Entity: “alive” if memory is allocated (and initialized).
- Binding: “alive” if the name refers to some entity.

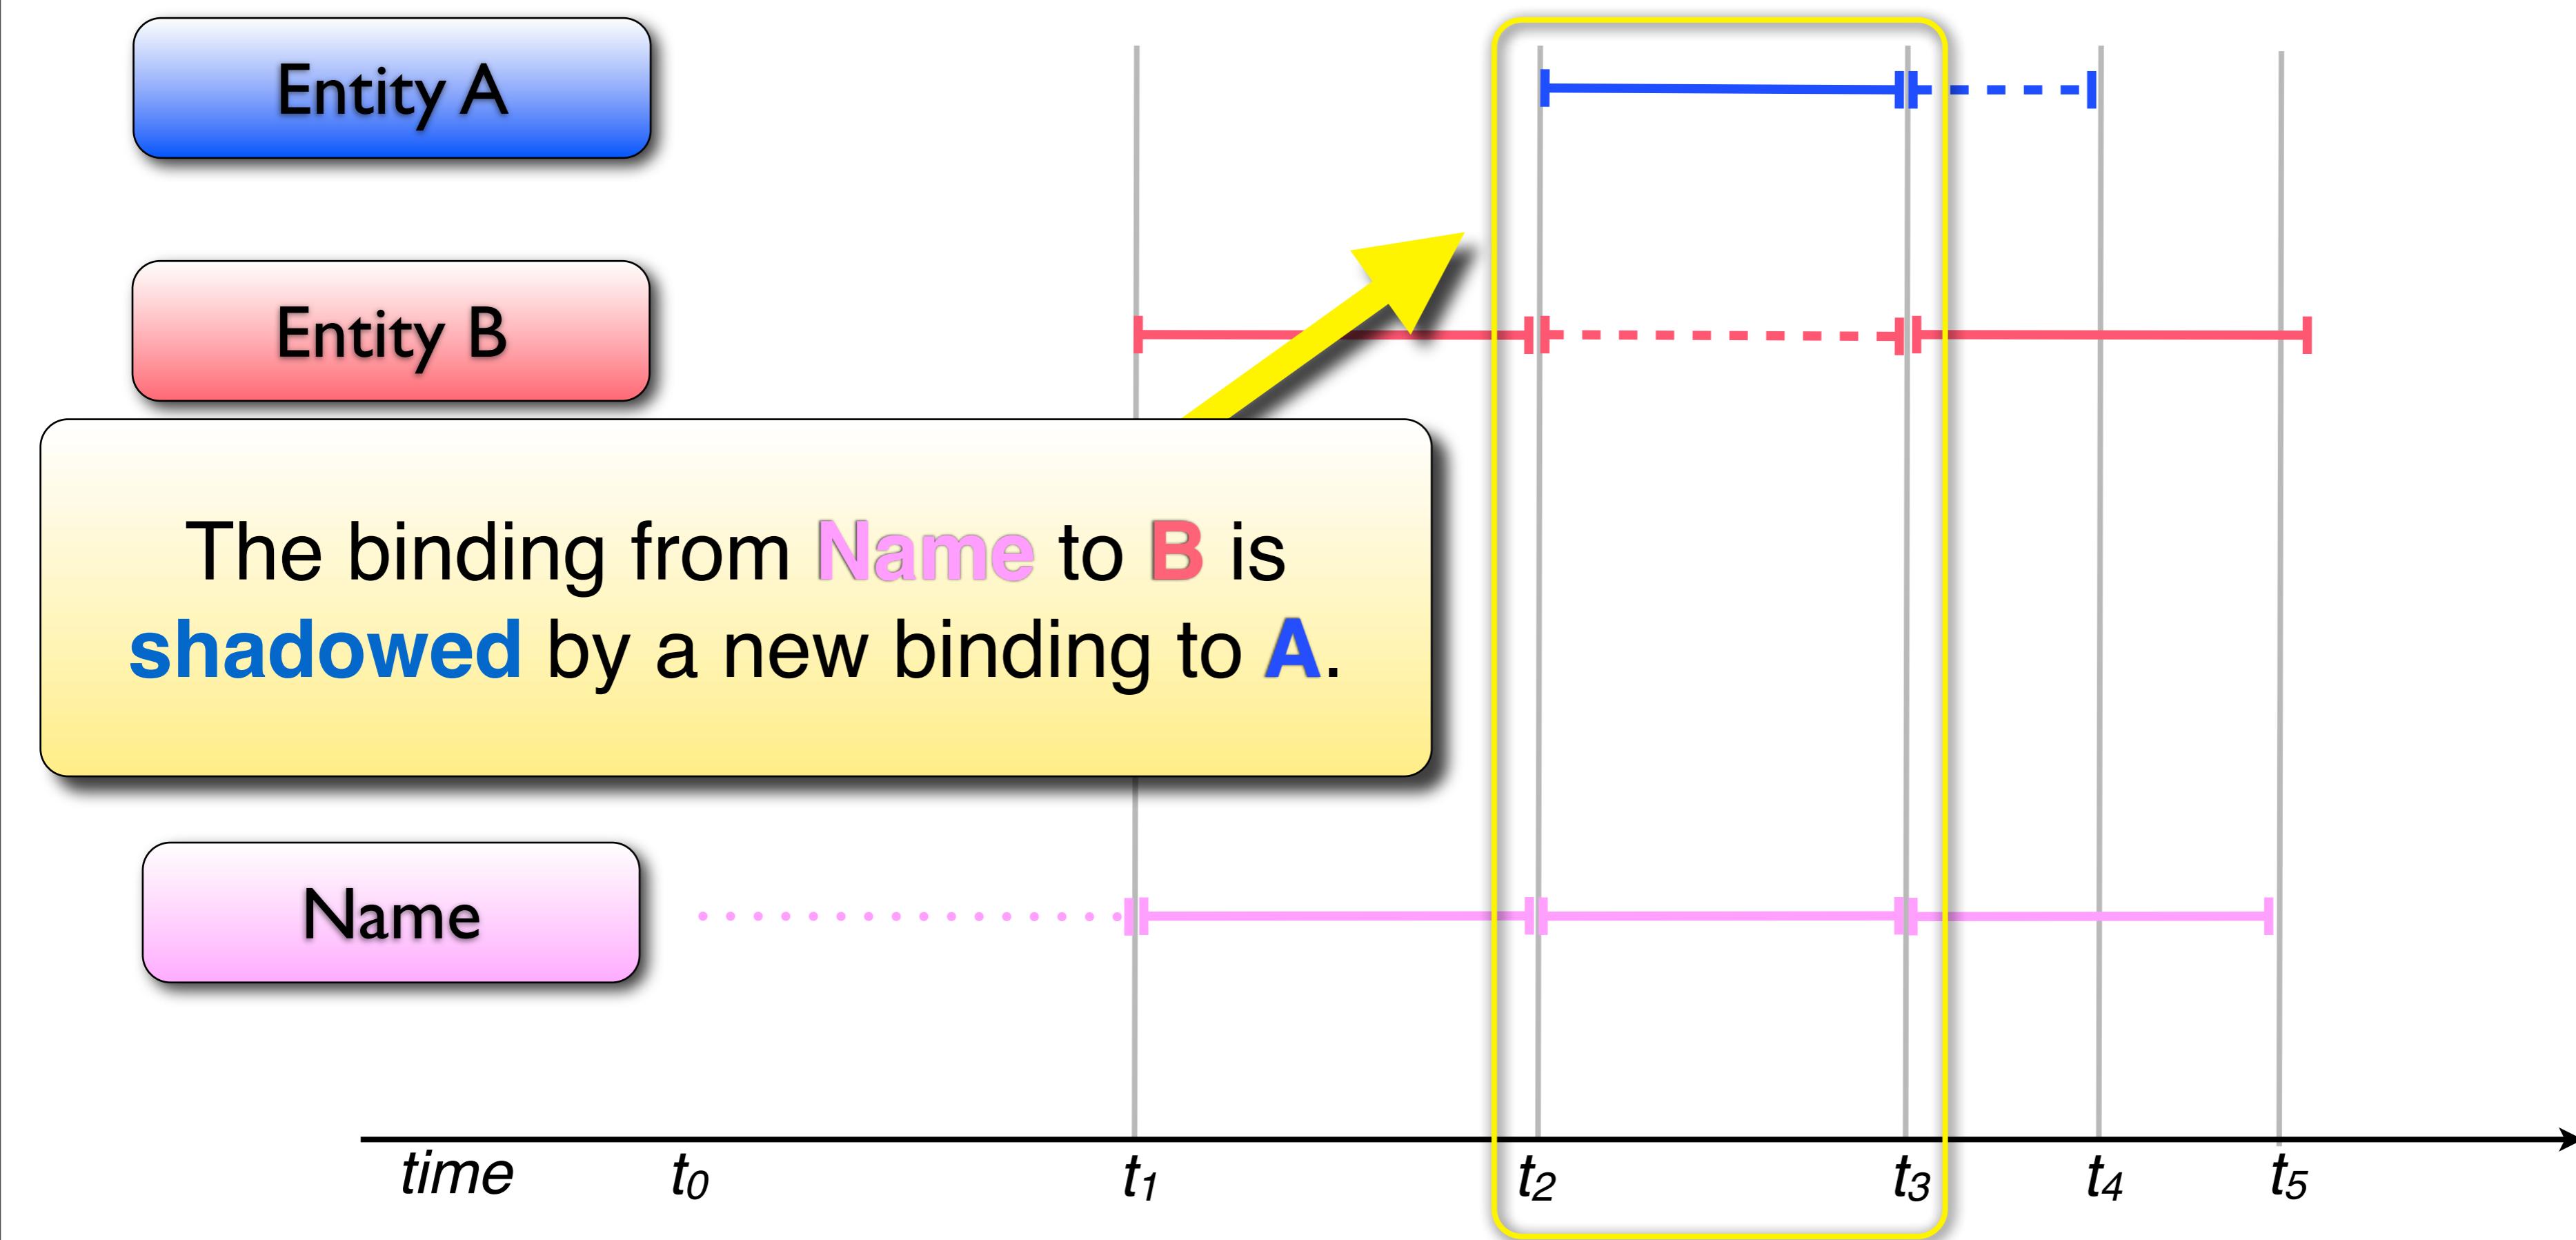
Object Lifetime vs. Binding Lifetime



Lifetime

- Entity: “alive” if memory is allocated (and initialized).
- Binding: “alive” if the name refers to some entity.

Object Lifetime vs. Binding Lifetime



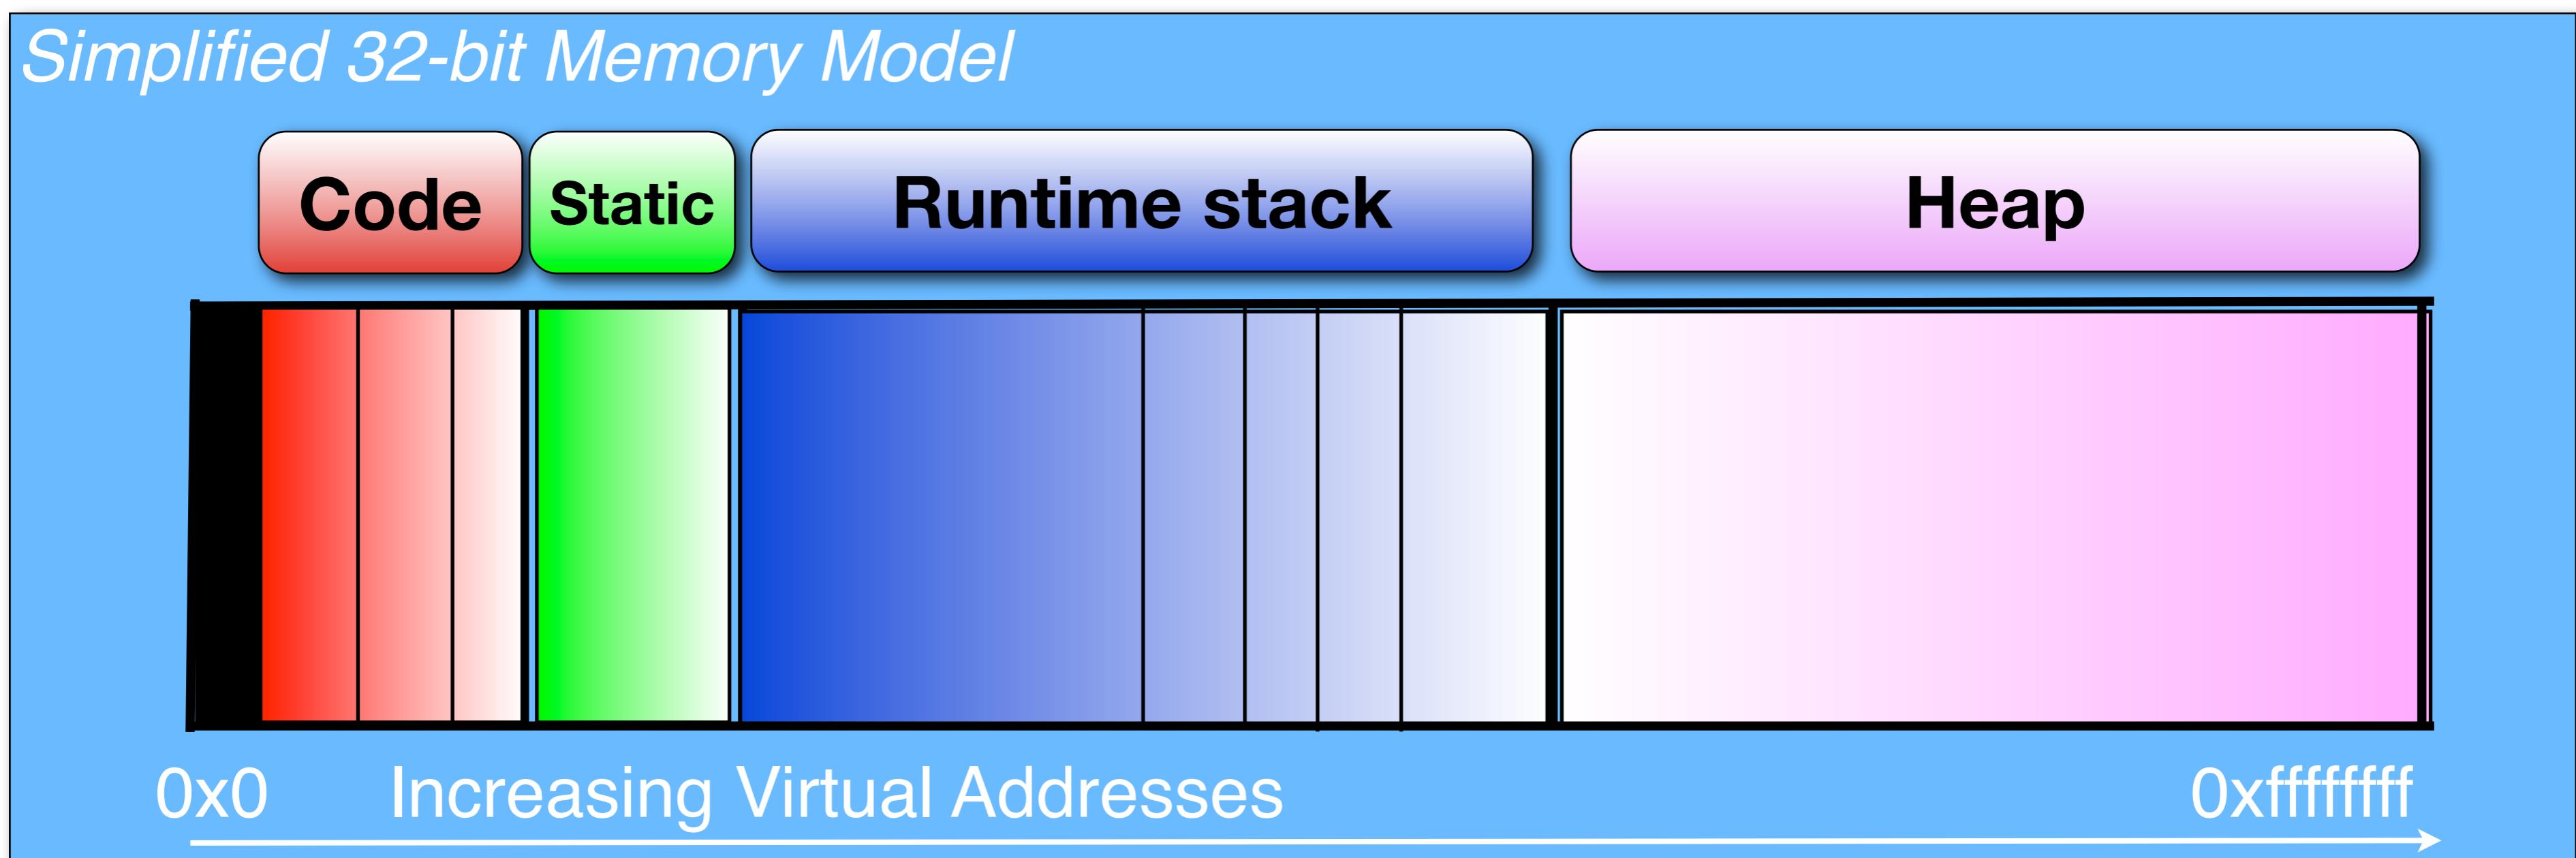
Lifetime

- Entity: “alive” if memory is allocated (and initialized).
- Binding: “alive” if the name refers to some entity.

Object Lifetimes

Defined by three principal storage allocation mechanisms:

- **Static** objects, which have a **fixed address** and are not deallocated before program termination.
- **Stack** objects, which are allocated and deallocated in a **Last-In First-Out** (LIFO) order.
- **Heap** objects, which are allocated and deallocated at **arbitrary times**.



Static Allocation

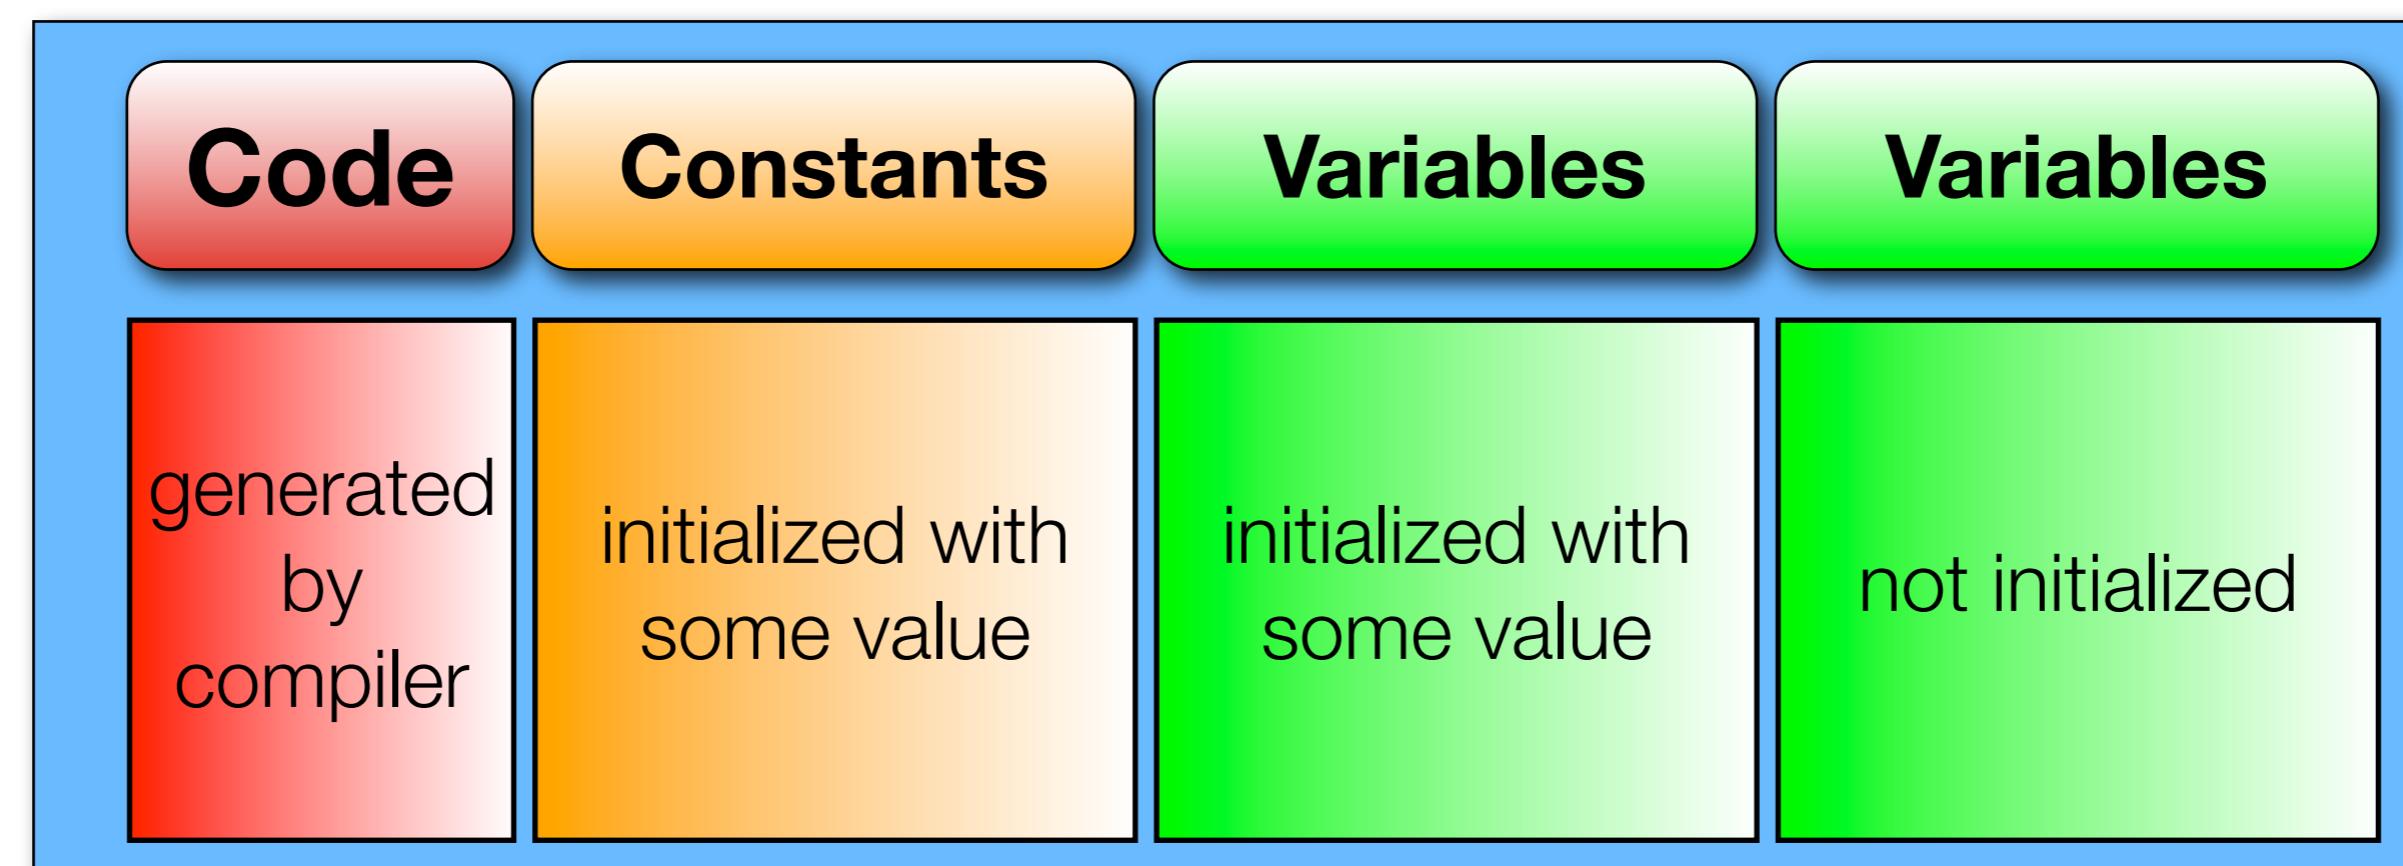
Some memory is required throughout program execution.

- Multi-byte **constants**.
 - Strings (“hello world”).
 - Lookup tables.
- **Global variables**.
- The program **code**.

Caution: this is not the same as Java’s **static**.

Must be allocated before program execution starts.

- Requirements specified in program file.
- **Allocated by OS** as part of program loading.
- The **size of static allocation is constant** between runs.



Static Allocation

Some memory is required throughout program execution.

- Multi-byte **constants**
 - String
 - Long
- **Globally visible variables**
- The stack

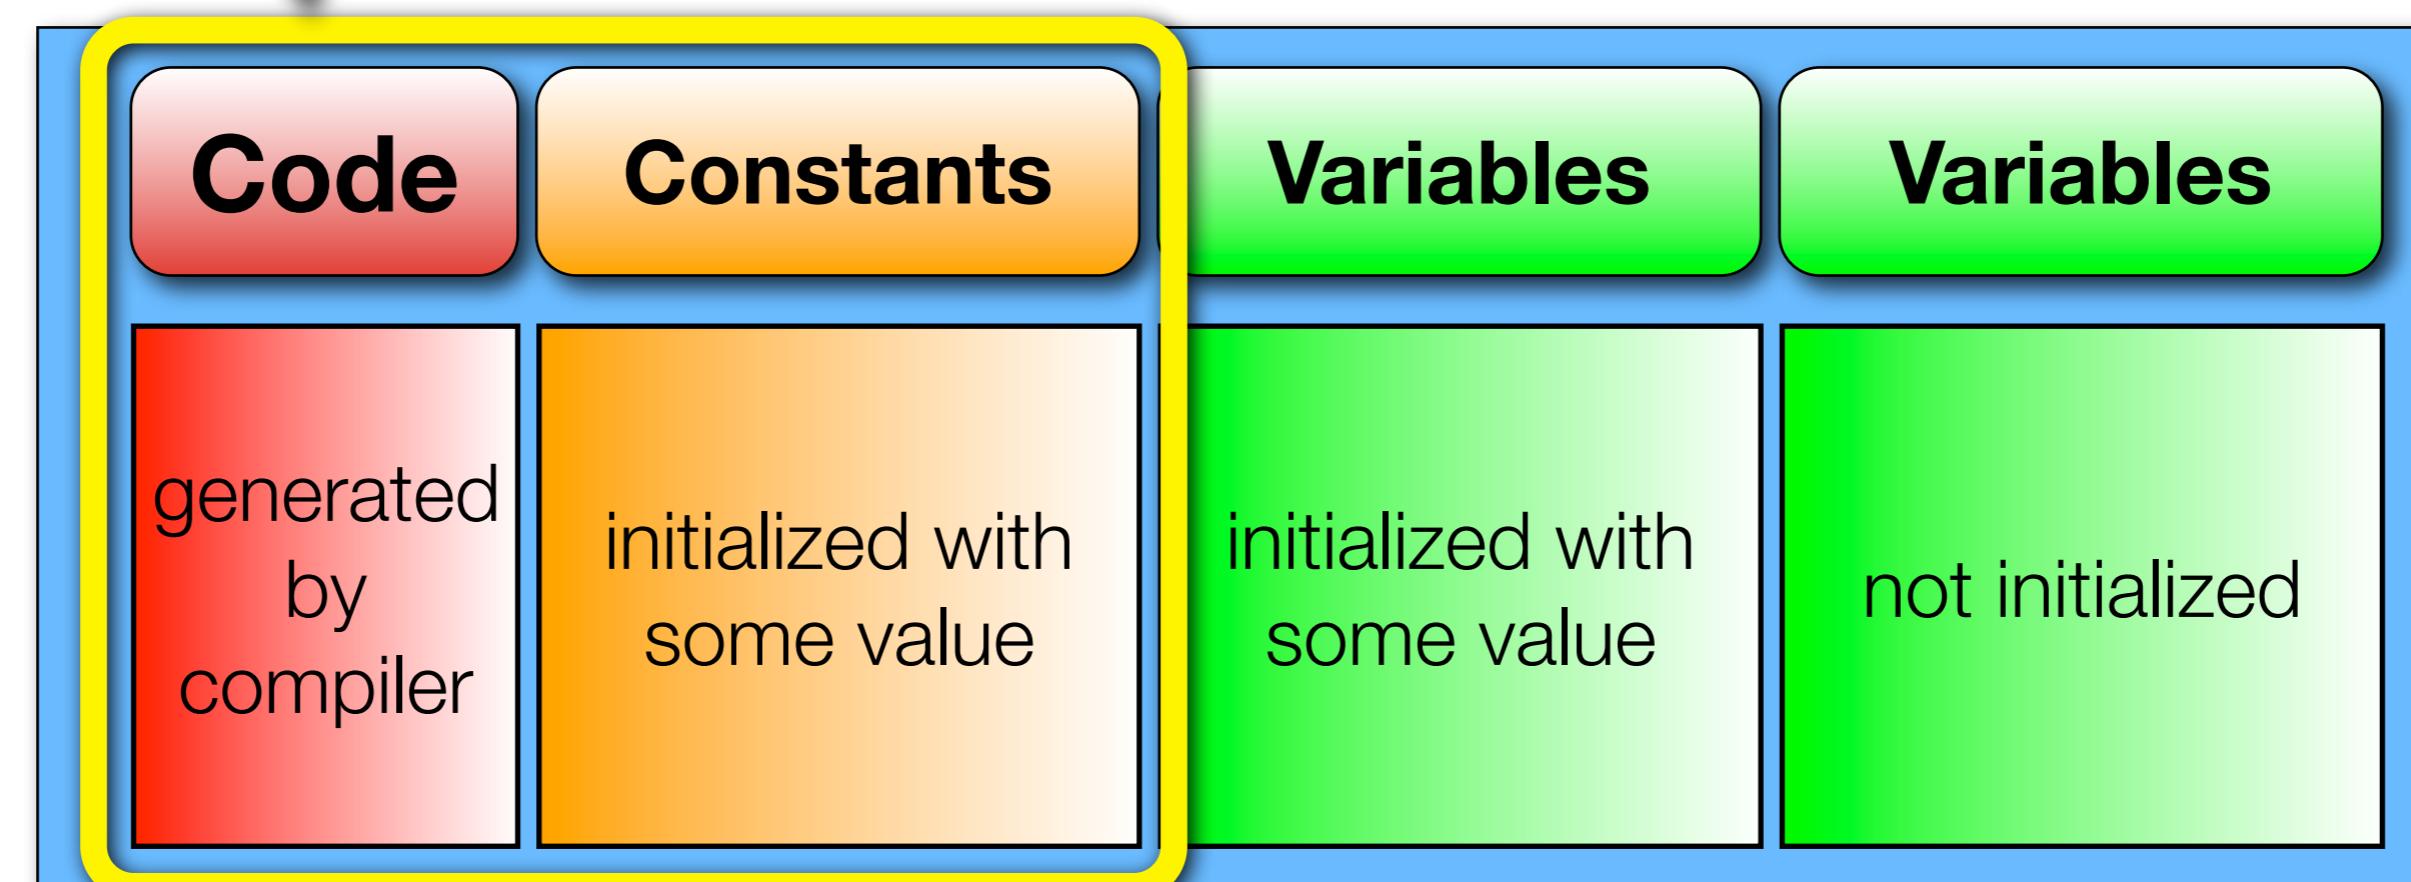
Read-only data.

Attempt to update illegal
in many operating systems.

is not the same as Java's **static**.

Must be allocated before program execution starts.

- Requirements specified in program file.
- **Allocated by OS** as part of program loading.
- The **size of static allocation is constant** between runs.



Static Allocation

Some memory is required throughout program execution.

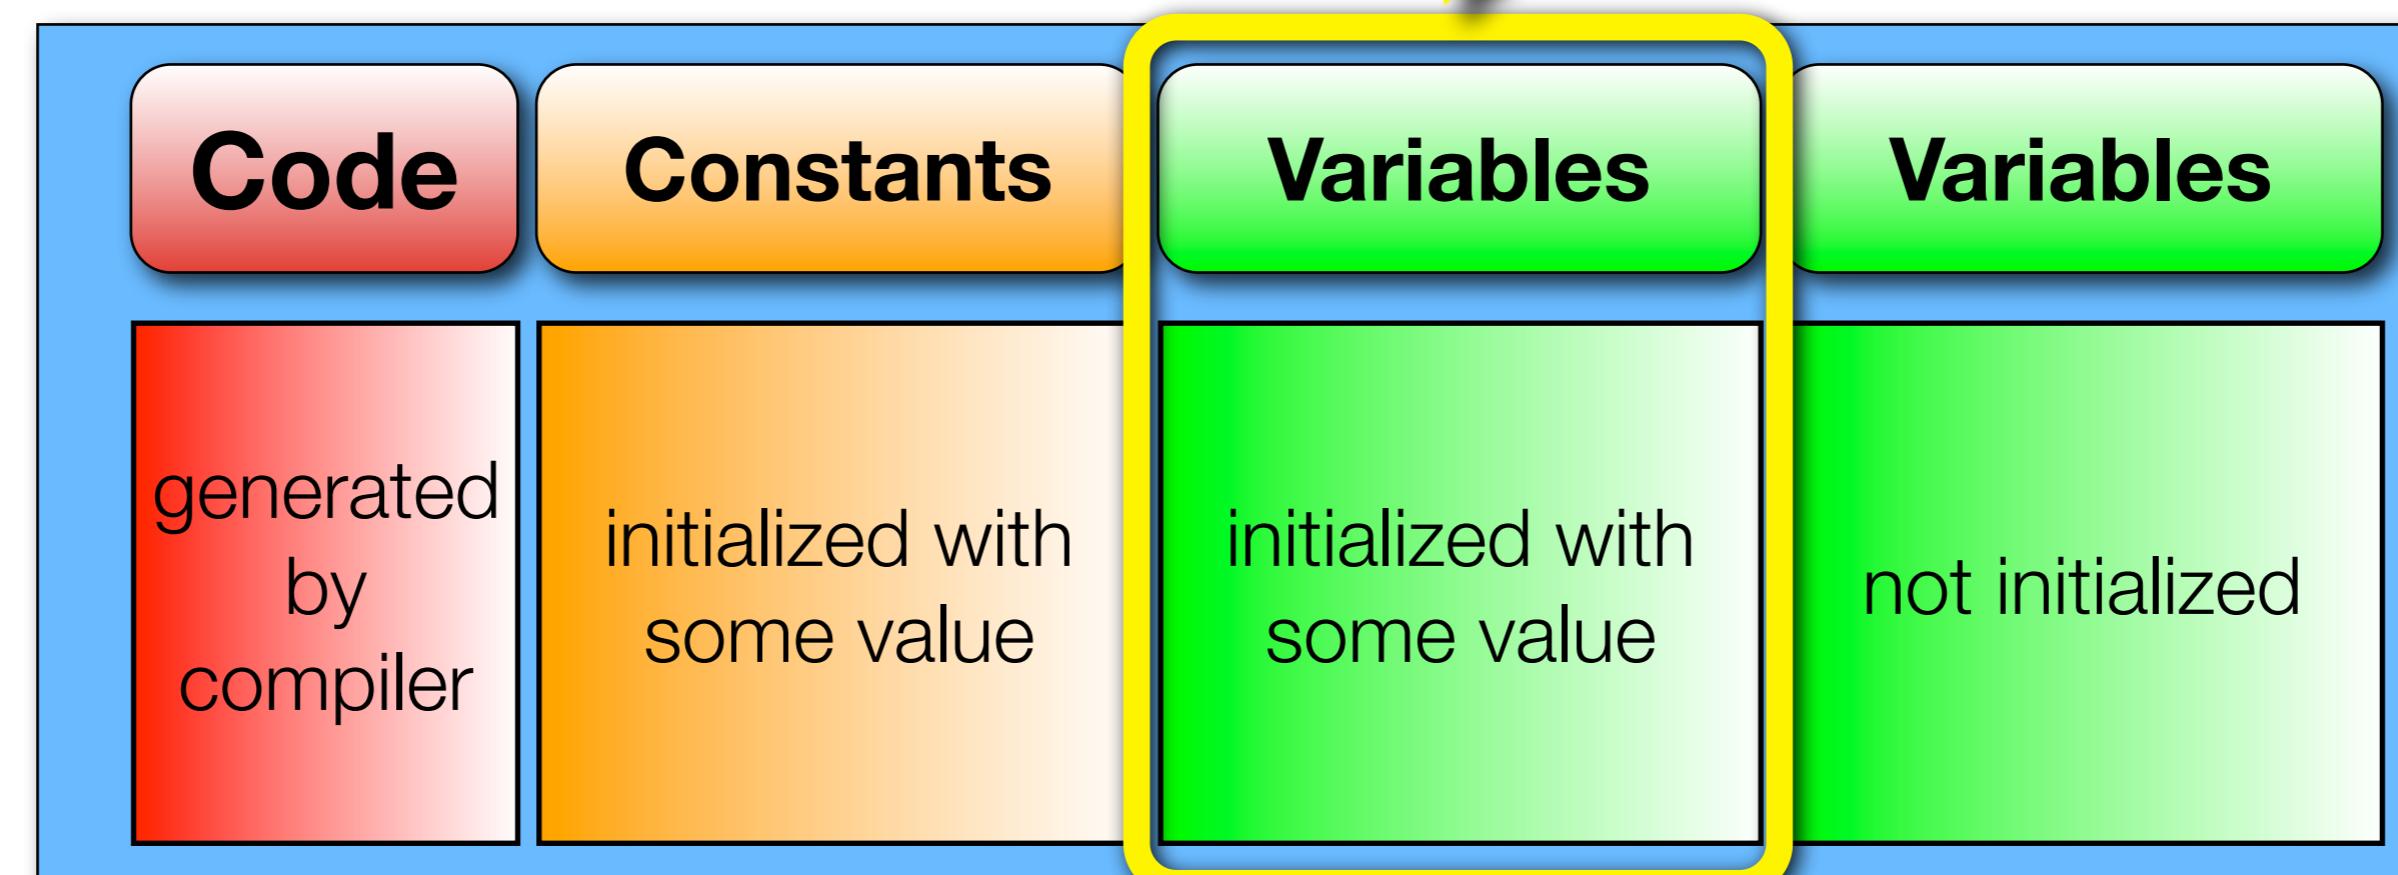
- Multi-byte **constants**.
 - Strings (“hello world”).
 - Lookup tables.
- **Global variables**.
- The program **code**.

Caution

Global Initialized Variables.
E.g., `int meaning = 42;`

Must be allocated before program execution.

- Requirements specified in program file.
- **Allocated by OS** as part of program loading.
- The **size of static allocation is constant** between runs.



Static Allocation

Some memory is required throughout program execution.

- Multi-byte **constants**.
 - Strings ("hello world")
 - Lookup tables.
- **Global variables**.
- The program **code**

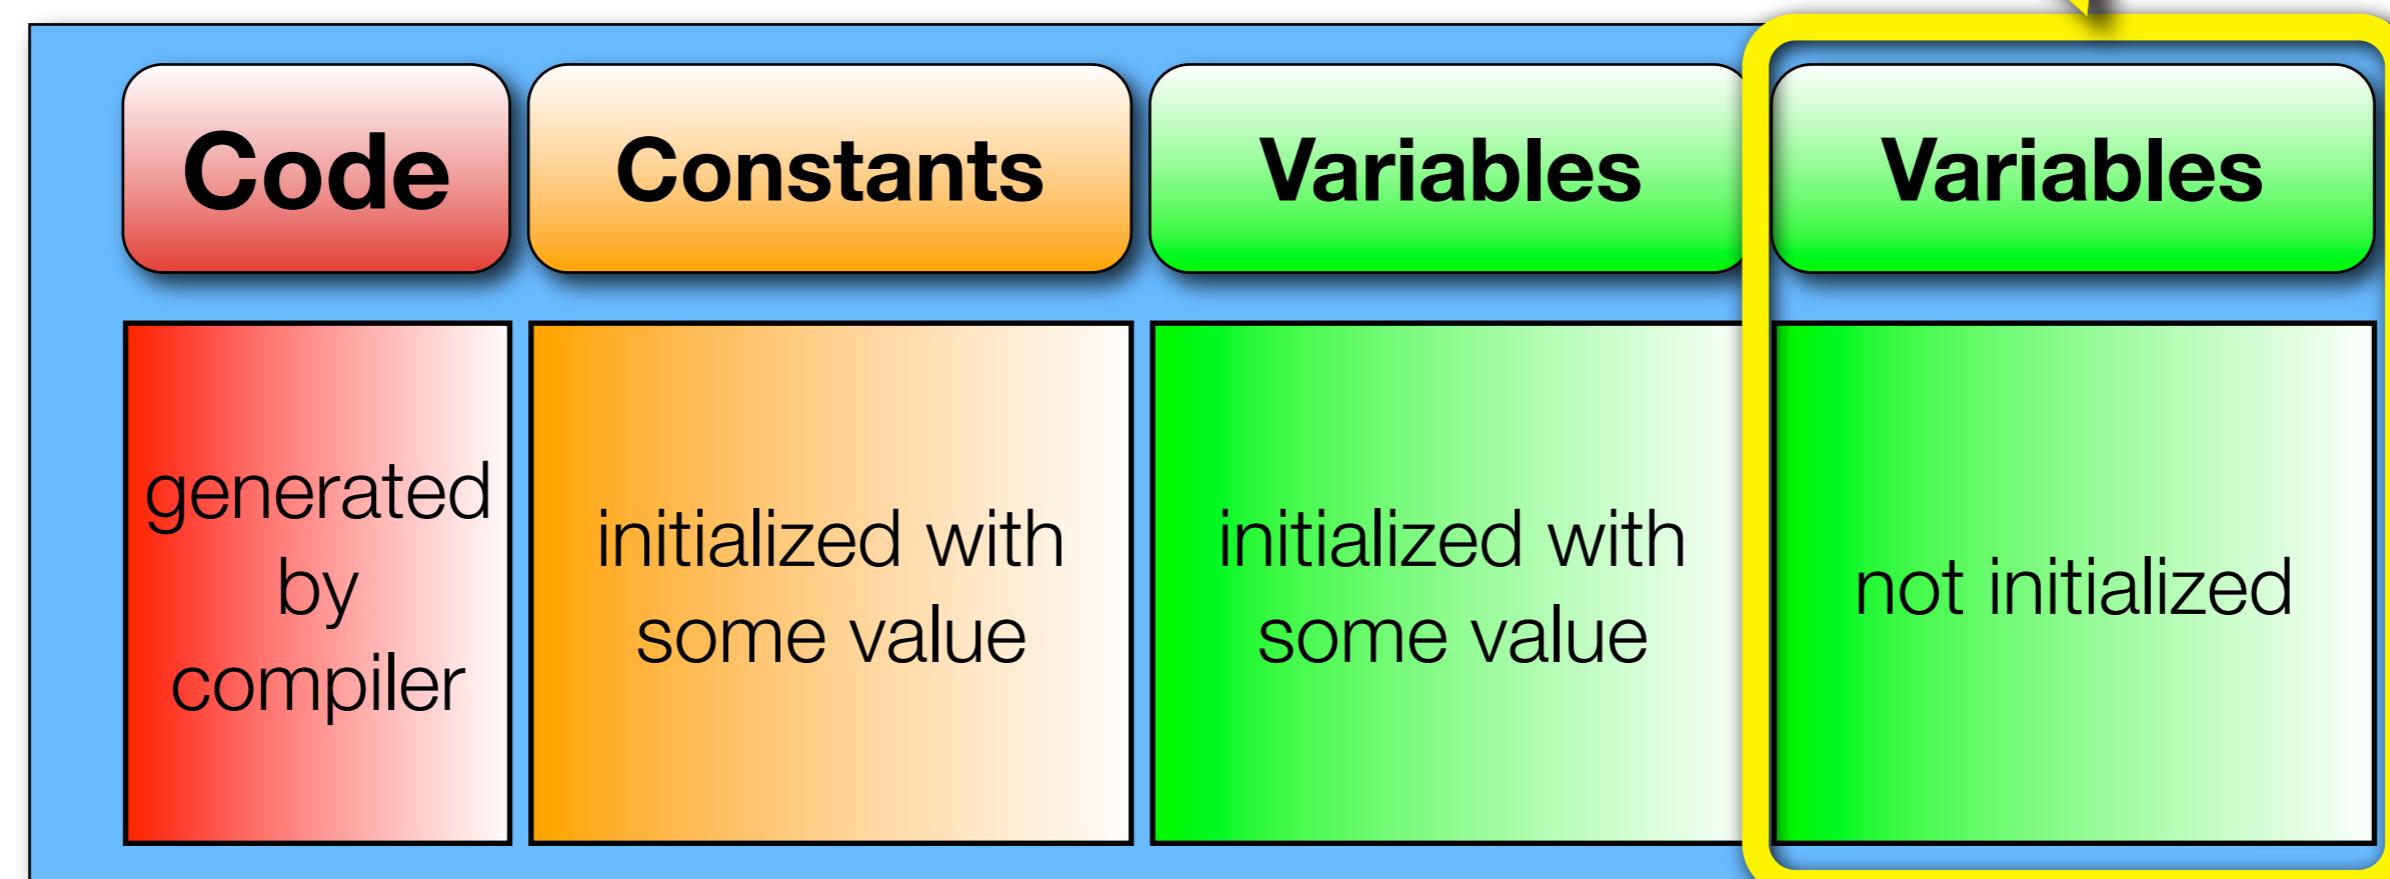
Global Uninitialized Variables.

E.g., `int last_error;`

(For historic reasons called the **.bss** segment.)

Must be allocated before execution.

- Requirements specified in program file.
- **Allocated by OS** as part of program loading.
- The **size of static allocation is constant** between runs.



Static Allocation

Some memory is required throughout program execution.

- Multi-byte **constants**

Compile-time constants.

Value must be known at compile time.

in the program code.

Must be allocated before program execution.

- Requirements specified in program file.

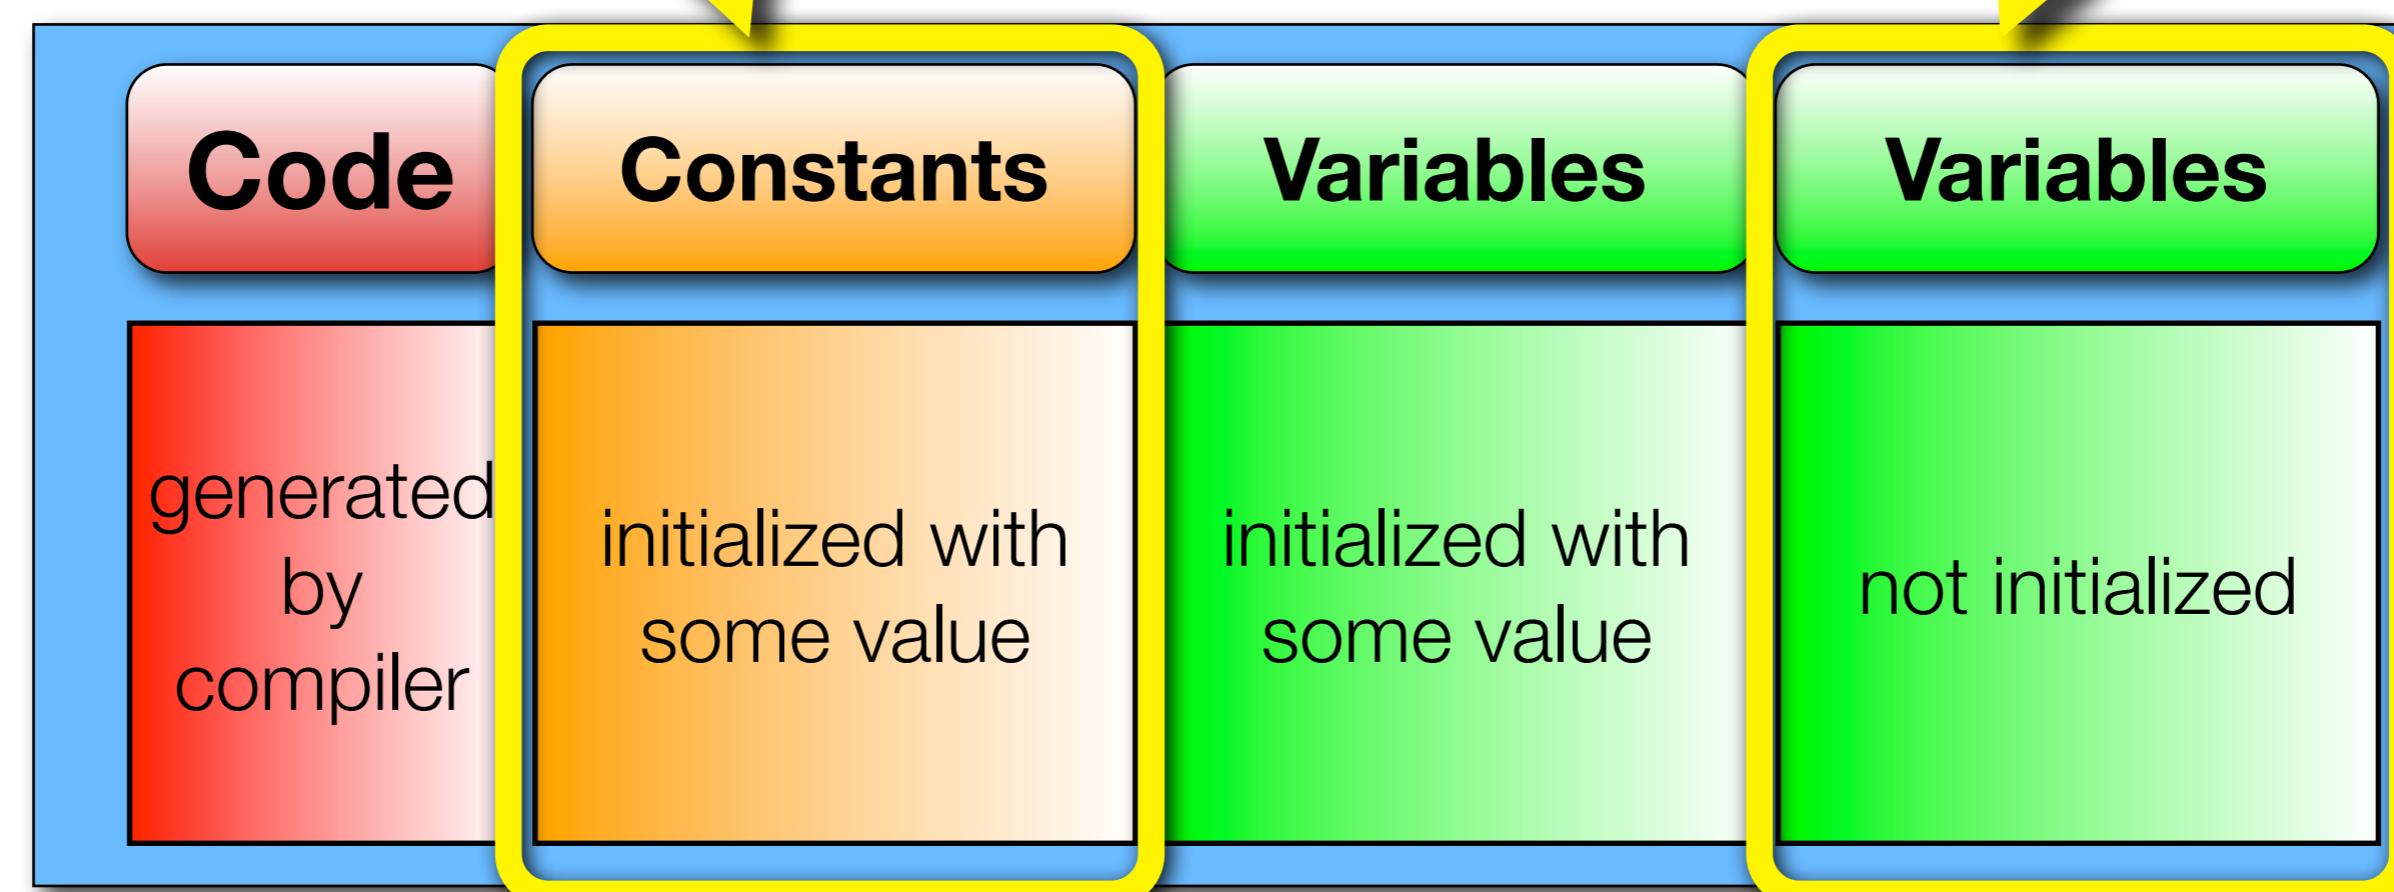
- **Allocated by OS** as part of program loading.

- The **size of static allocation is constant** between runs.

This is not the same as Java's **static**.

Elaboration-time constants.

Value computed at runtime; compiler disallows subsequent updates.



Advantages & Disadvantages

Advantages.

- No allocation/deallocation **runtime overhead**.
- Static addresses.
- Compiler can **optimize** accesses.

Limitations.

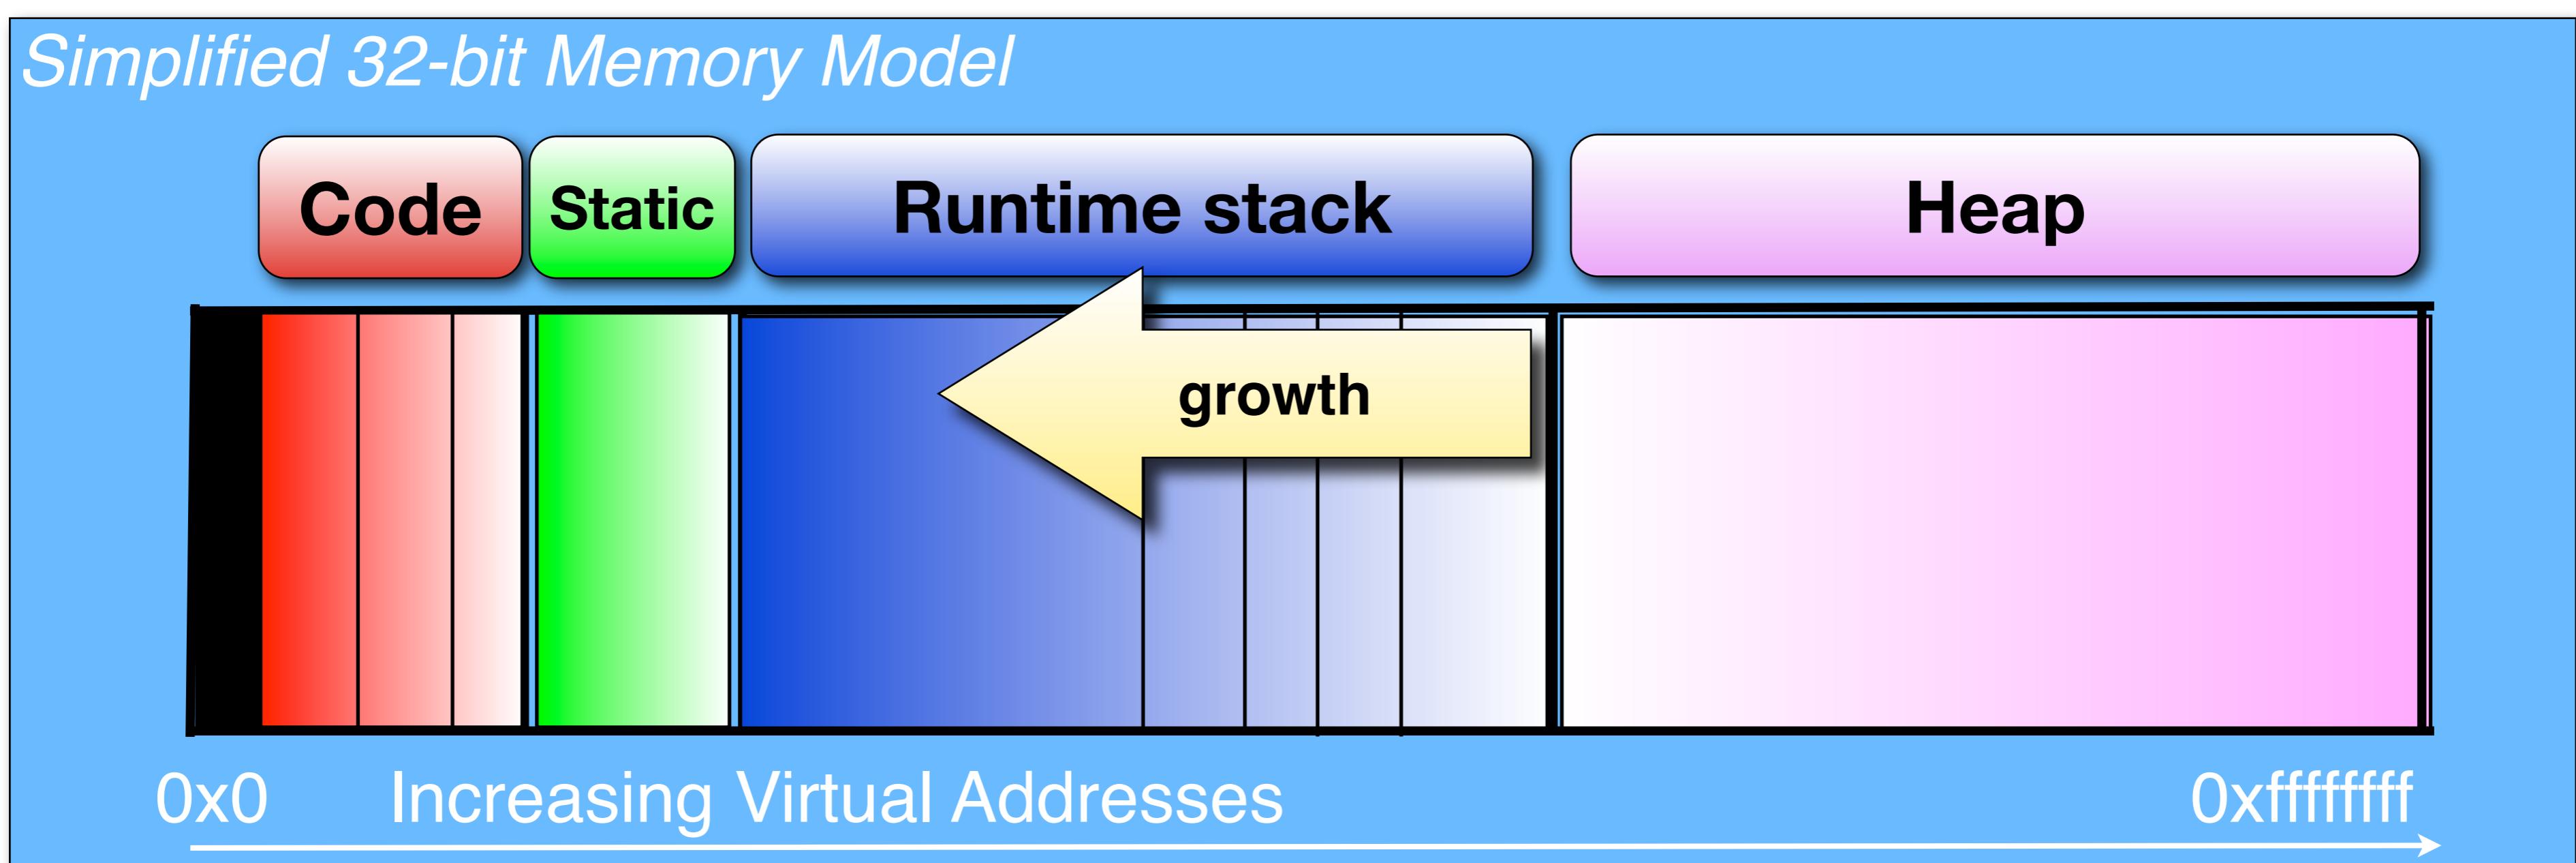
- Allocation **size fixed**; cannot depend on input.
- **Wasteful**; memory is always allocated.
- Global variables are **error-prone**.

Advice: avoid global variables.

Runtime Stack

Hardware-supported allocation area.

- Essential for **subprogram calls**.
- Grows **top-down** in many architectures.
- **Size limit**: stack overflow if available space is exhausted.
- Max. size of stack can be **adjusted at runtime**.
- OS is often involved in stack management.



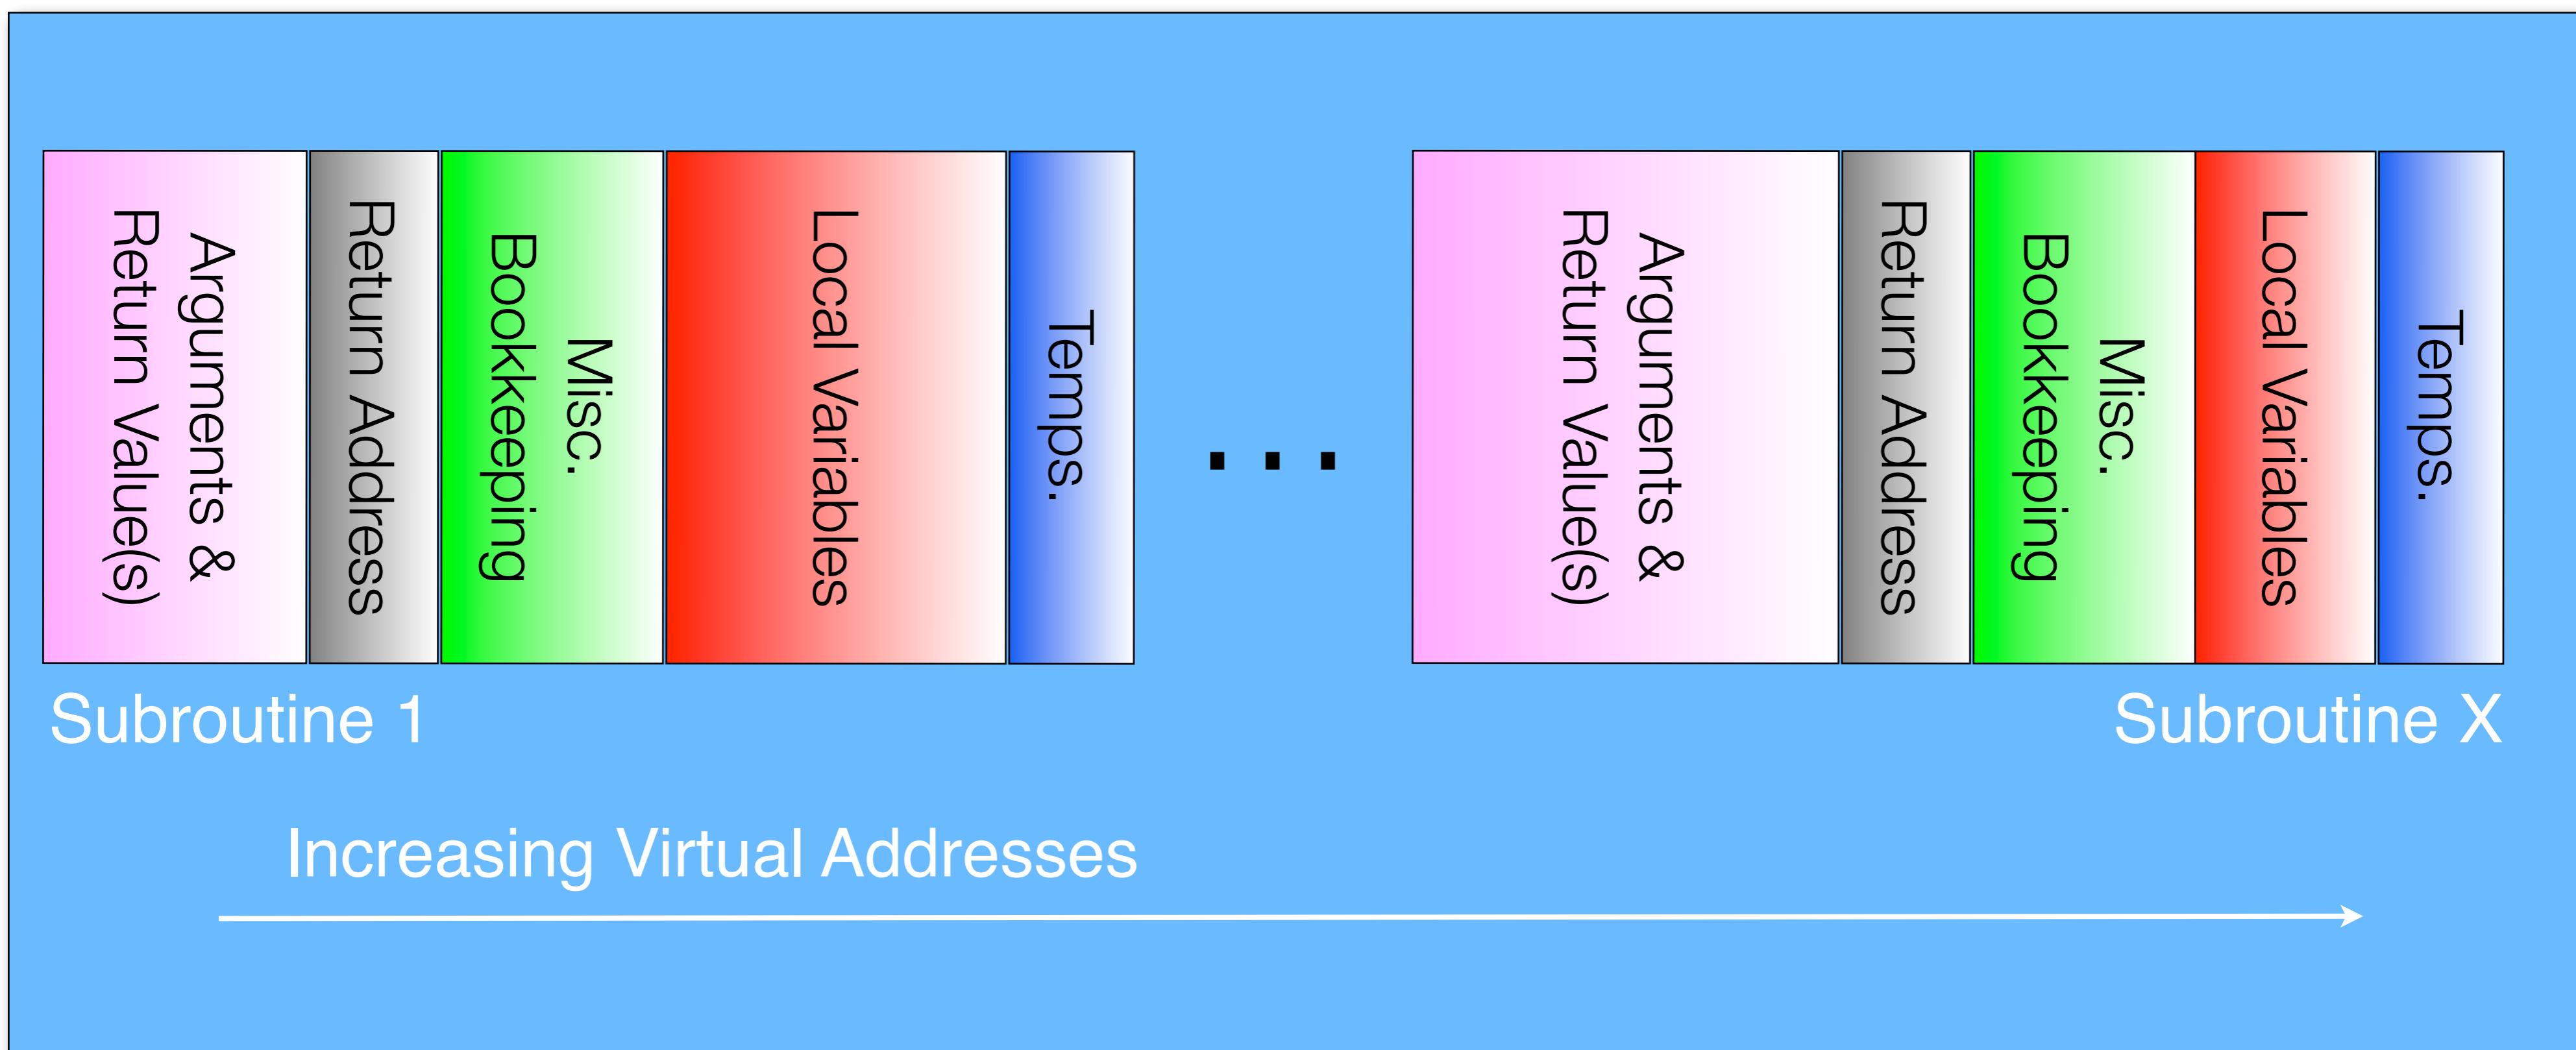
Subroutines & Static Allocation

Calling a function/method/subroutine requires memory.

- **Arguments**.
- **Local variables**.
- **Return address**
(where to continue execution after subroutine completes).
- Some **bookkeeping** information.
 - E.g., to support **exceptions** and call **stack traces**.

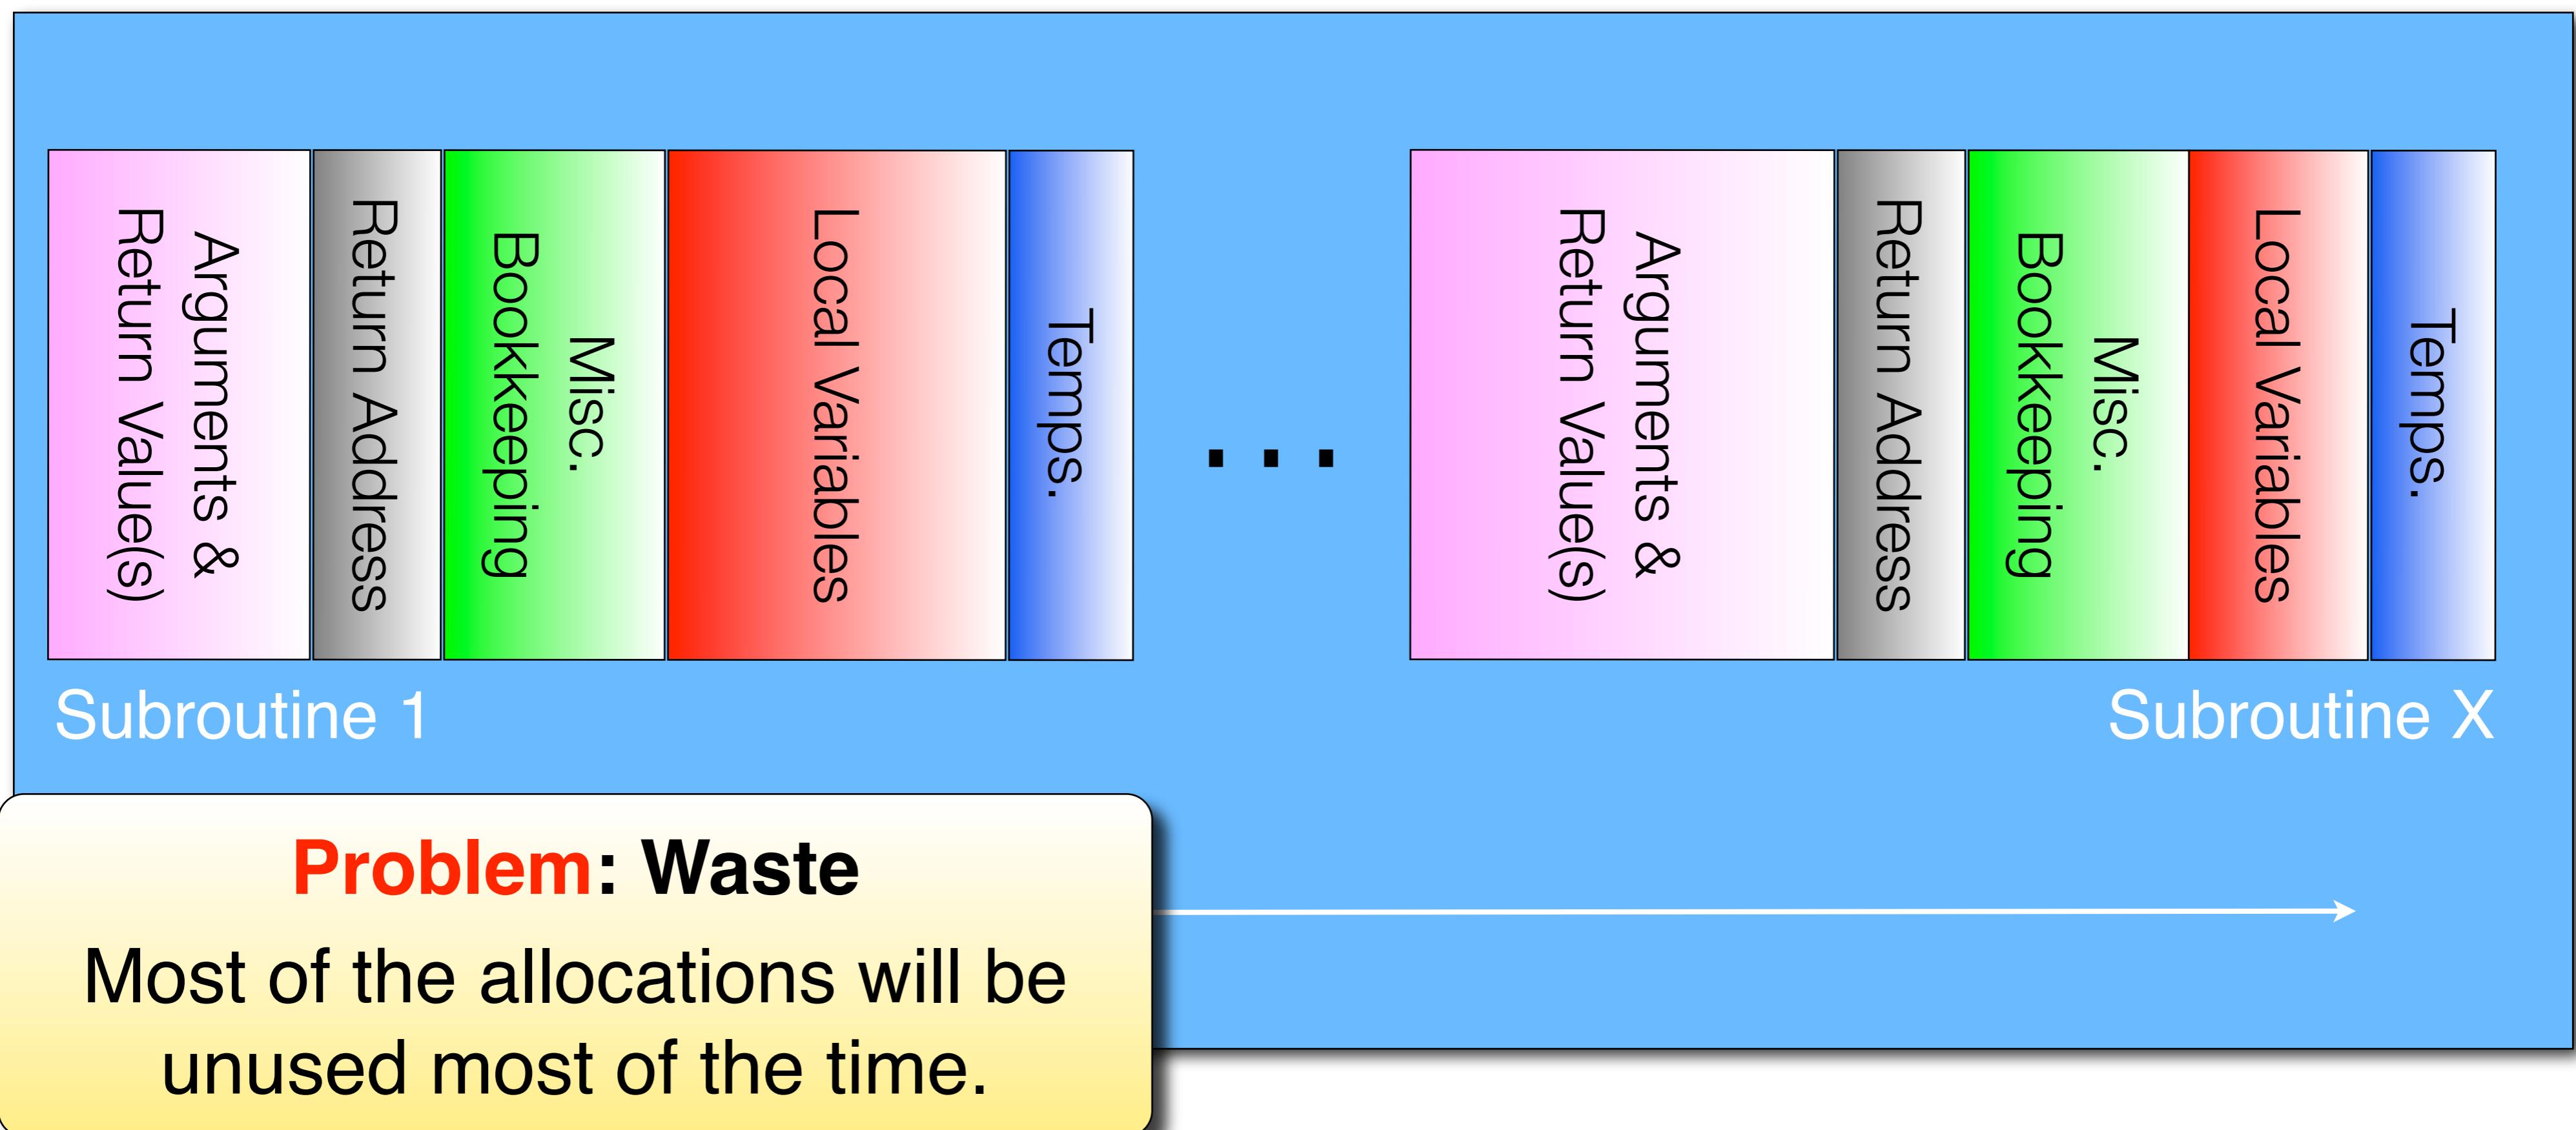
Where should this memory be allocated?

Static Allocation of Subroutine Memory



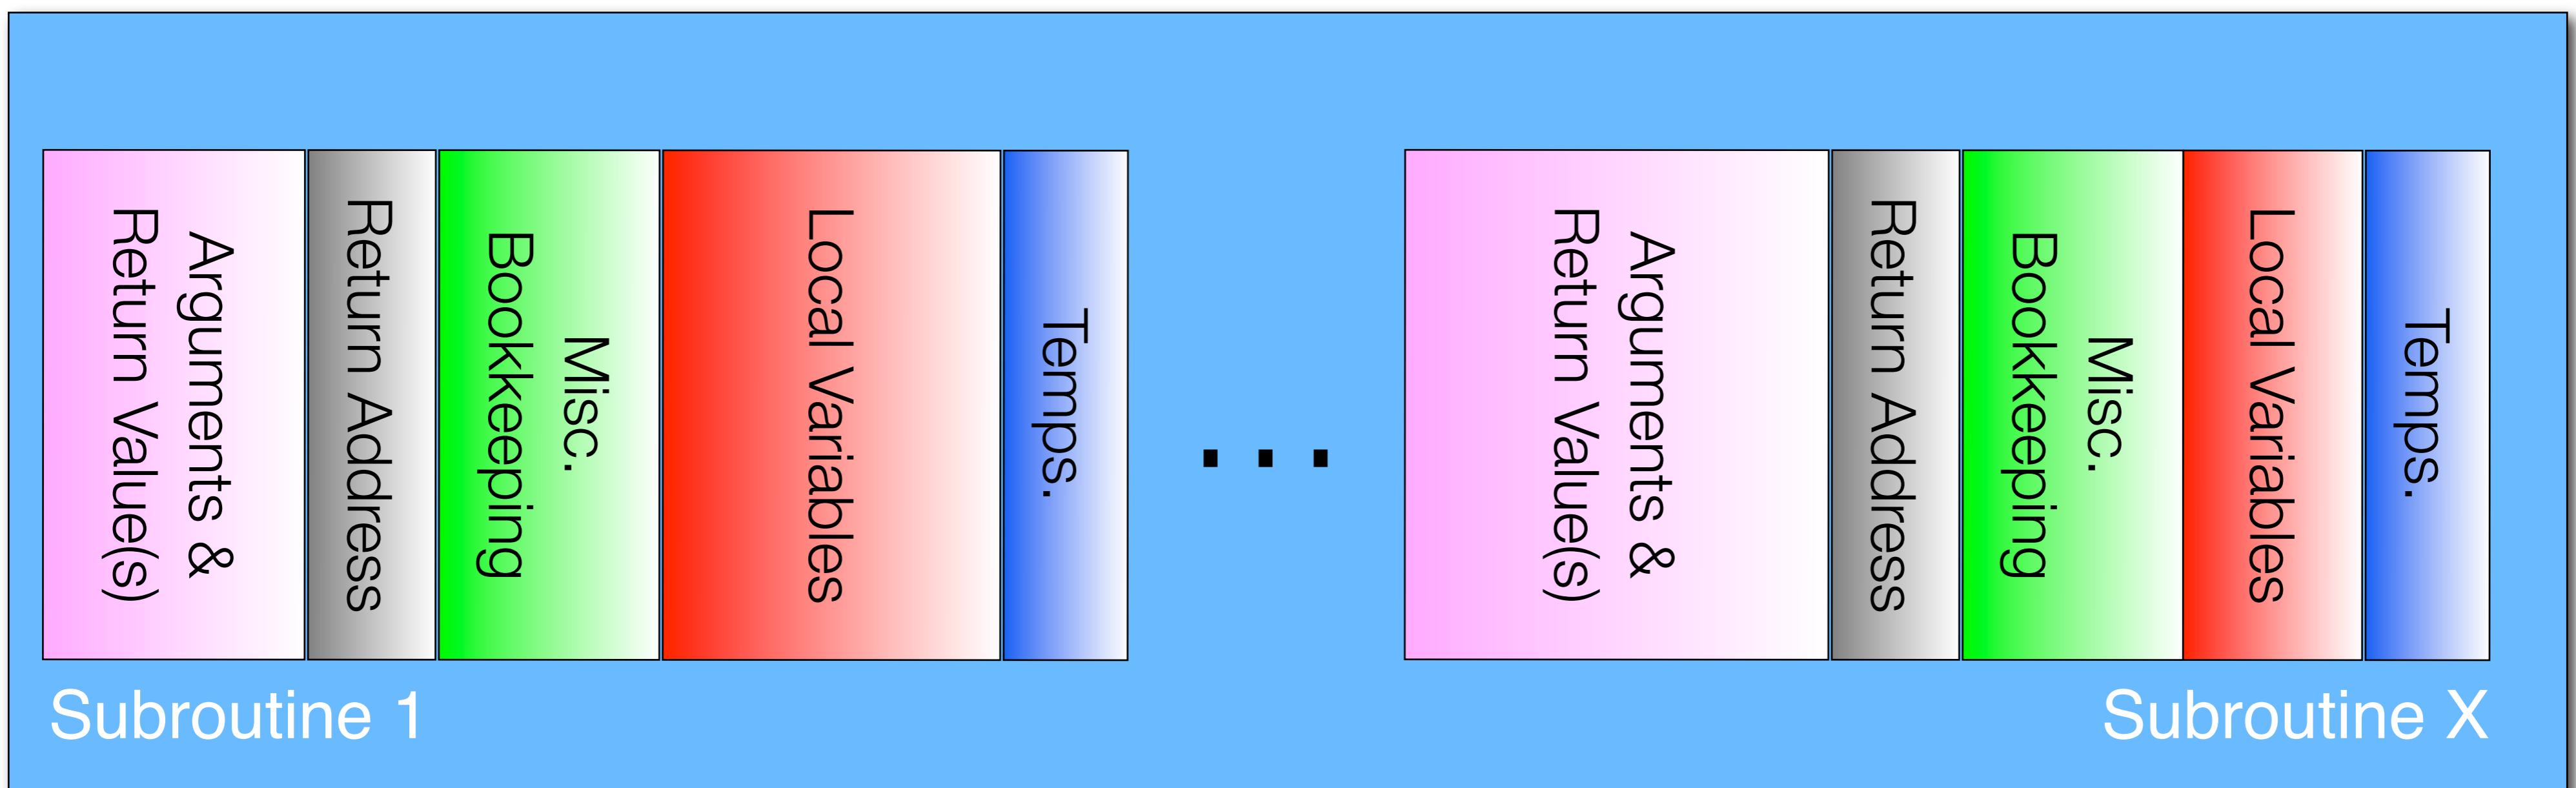
*One approach: statically allocate memory for each subroutine.
(e.g. early versions of Fortran)*

Static Allocation of Subroutine Memory



*One approach: statically allocate memory for each subroutine.
(e.g. early versions of Fortran)*

Static Allocation of Subroutine Memory



Problem: Waste

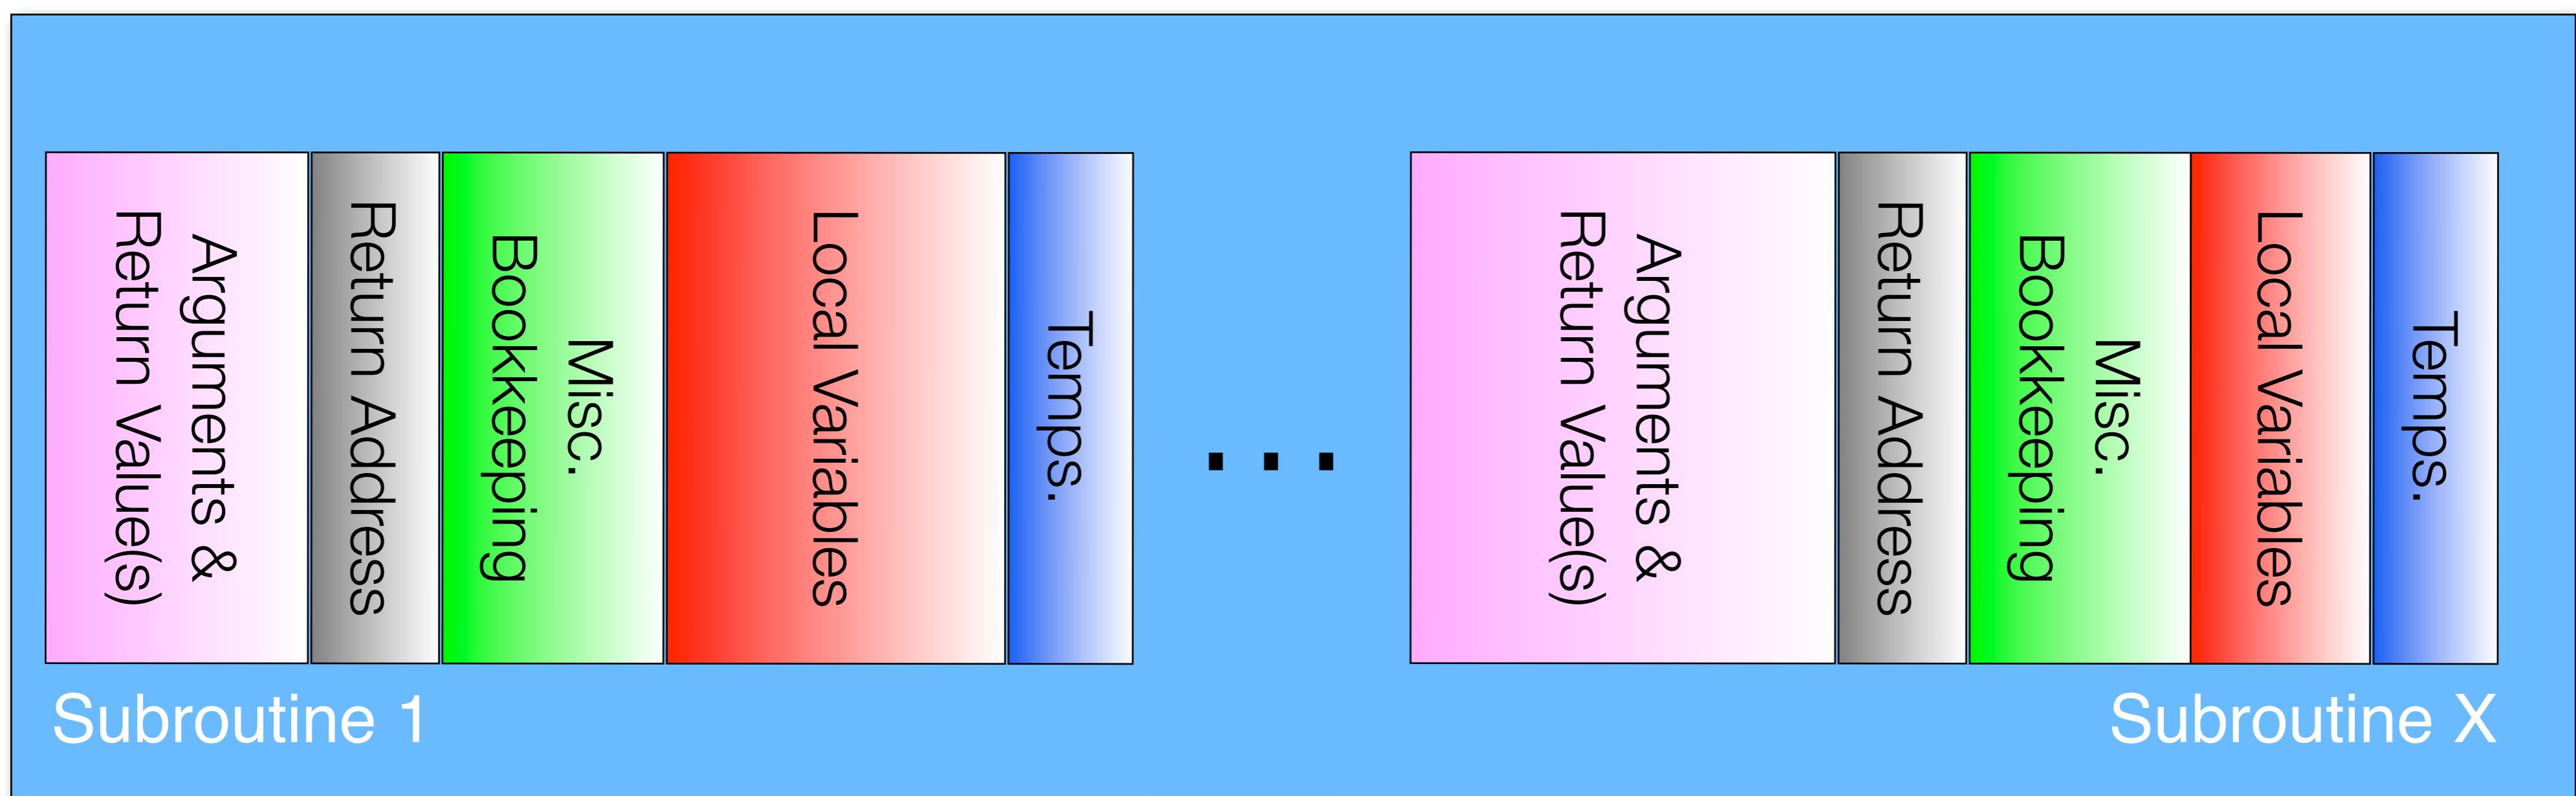
Most of the allocations will be unused most of the time.

Problem: No Recursion

Subroutines may not be called while their memory is already in use.

*One approach: statically allocate memory for each subroutine.
(e.g. early versions of Fortran)*

Static Allocation of Subroutine Memory



Problem: Waste

Most of the allocations will be unused most of the time.

Problem: No Recursion

Subroutines may not be called while their memory is already in use.

Limited recursion depth can be allowed by allocating memory for multiple subroutine *instances*. But this **increases waste...**

Runtime Stack

Idea: allocate memory for subroutines on demand.

- Reserve (large) area for subroutine calls. Allocate new **frame** for each call. Deallocate on return.
- Subroutine calls must be **fast**: need **efficient** memory management.

Observation: last-in first-out allocation pattern.

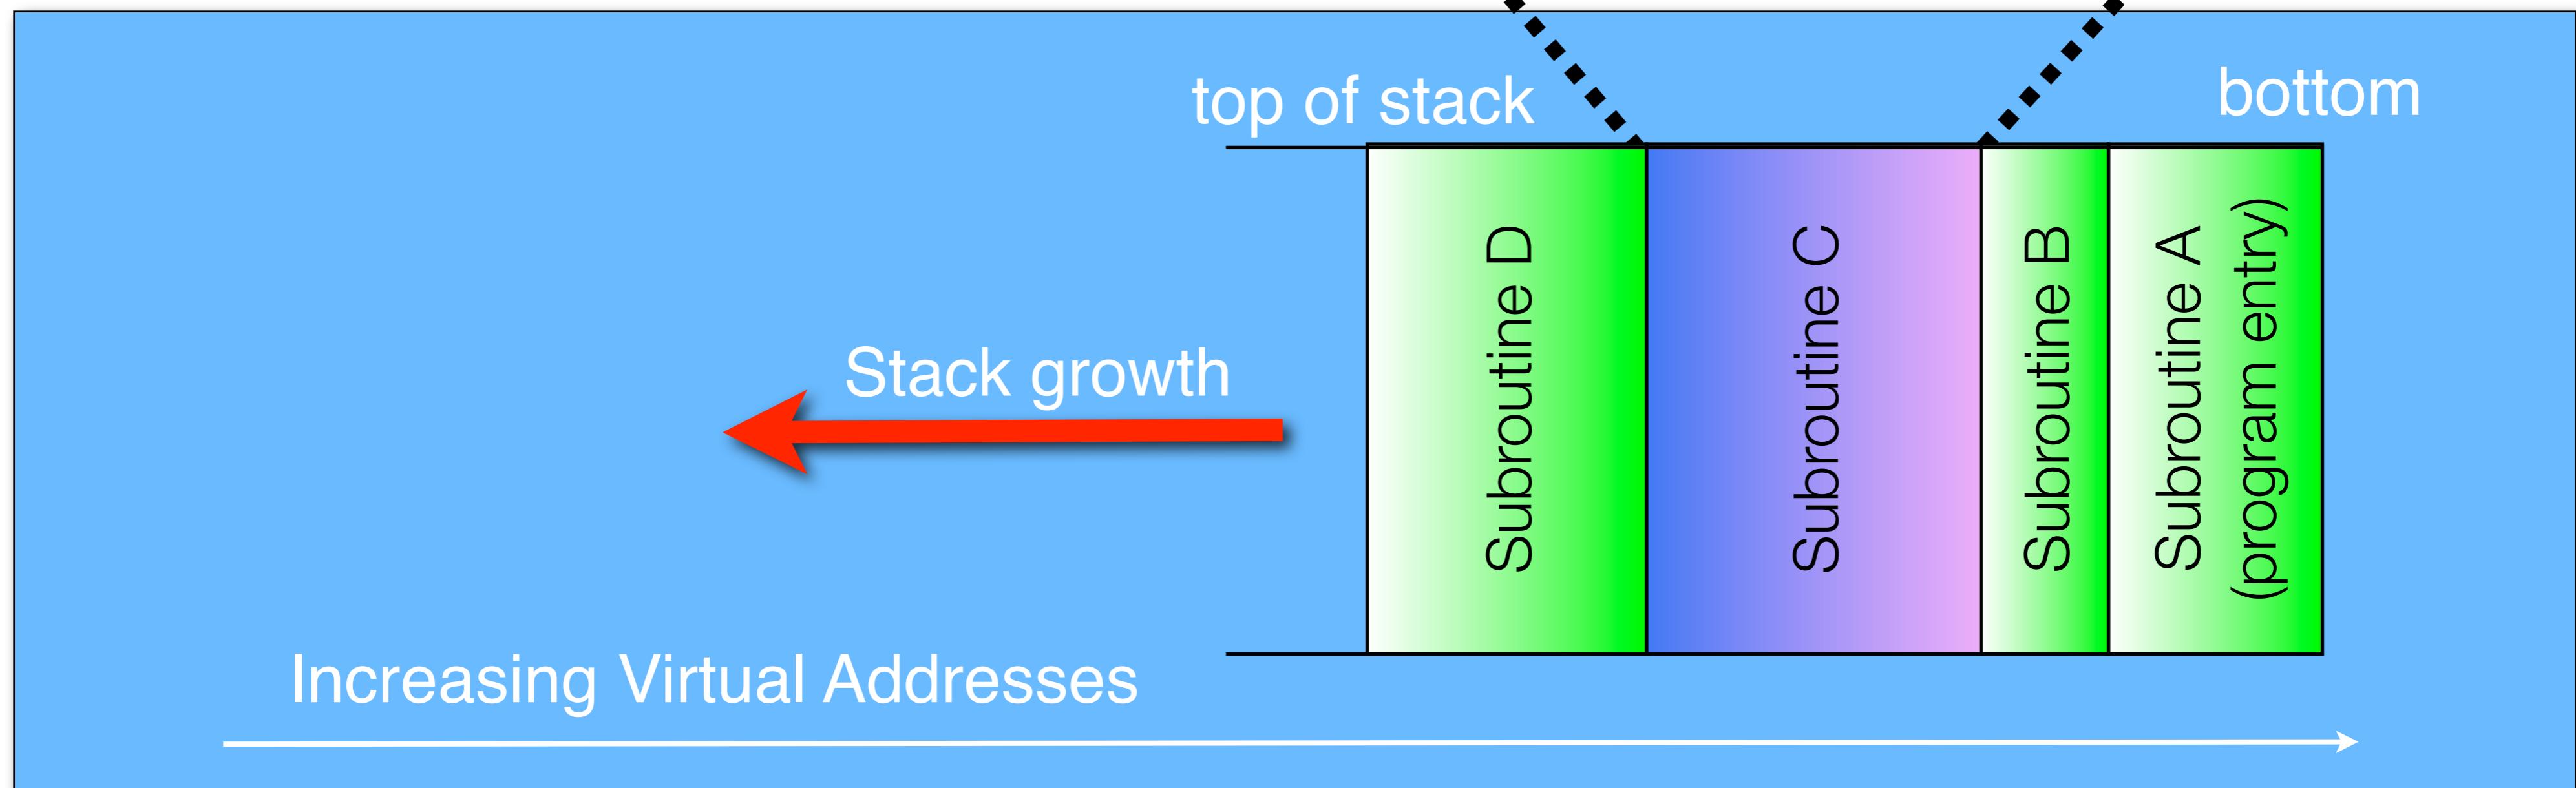
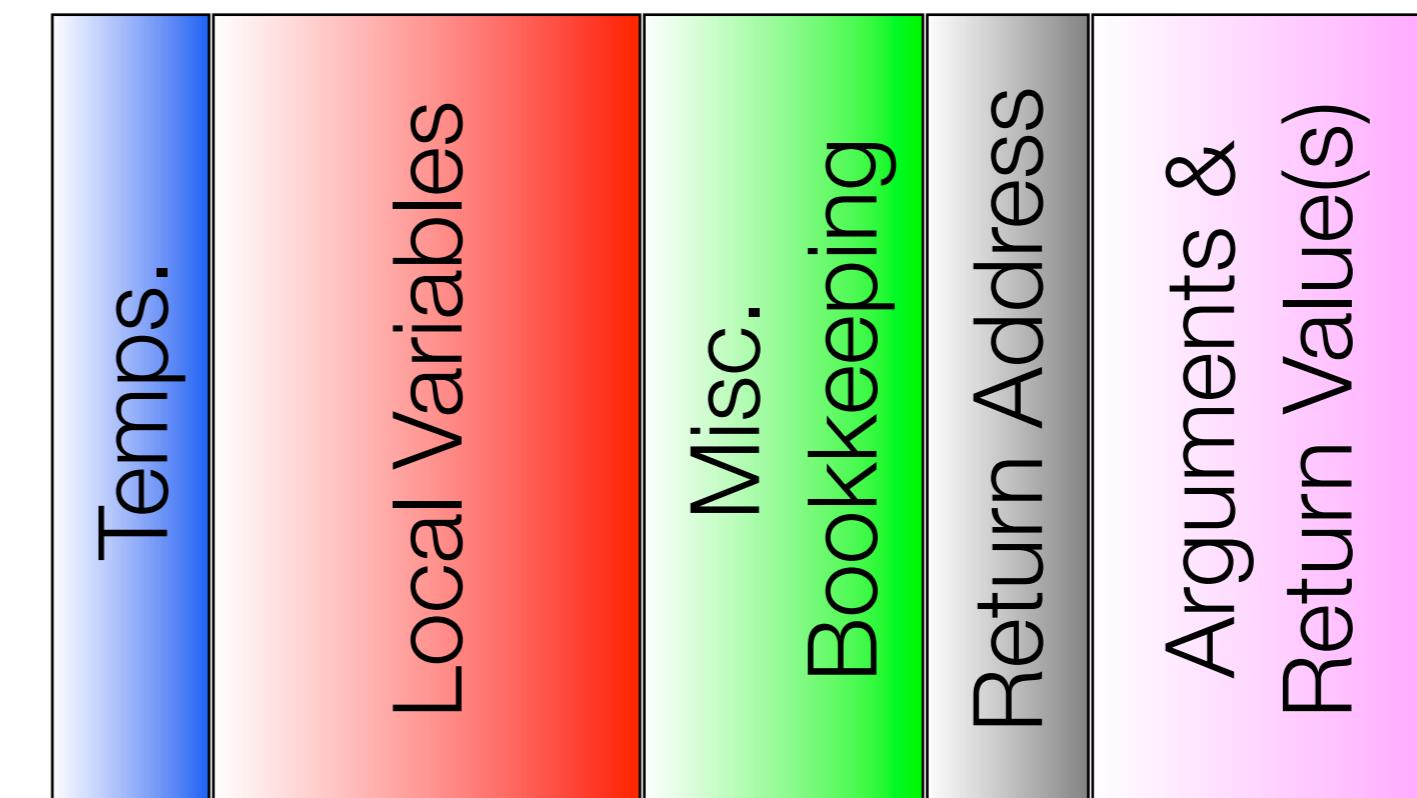
- A routine returns only after all called subroutines have returned.
- Thus, allocations can be “piled” on top of each other.

**Subroutine memory is allocated
on-demand from the runtime stack.**

Stack Frames

On a subroutine call:

- New **stack frame pushed** onto stack.
- Stack frames **differ in size** (depending on local variables).
- **Recursion** only limited by total size of stack.
- **Reduced waste**: unused subroutines only consume memory for code, **not variables**.
- Stack frame **popped** from stack on return.



Calling Sequence

Compilers generate code to manage the runtime stack.

- **Setup**, before call to subroutine.
- **Prologue**, before subroutine body executes.
- **Epilogue**, after subroutine body completes (the **return**).
- **Cleanup**, right after subroutine call.

```
private void checkForKleeneClosure(NFA fa) throws IOException {
    if (readNextChar() == '*') {
        // indeed, it's a closure.
        fa.createKleeneClosure();
    } else
        unreadLastChar();
}
```

Calling Sequence

Compilers generate code to manage the runtime stack.

- **Setup**, before call to subroutine.
- **Prologue**, before subroutine body executes.
- **Epilogue**, after subroutine body completes (the **return**).
- **Cleanup**, right after subroutine call.

```
private void checkForKleeneClosure(NFA fa) throws IOException {  
    if (readNextChar() == '*') {  
        // indeed, it's a closure.  
        fa.createKleeneClosure();  
    } else  
        unreadLastChar();  
}
```

Calling Sequence Example

```
int add(int a, int b)
{
    int c;
    c = a + b;
    return c;
}

int main(int argc, char** argv)
{
    int sum = add(10, 20);
    printf("sum = %d\n", sum);
    return 0;
}
```

C program

Calling Sequence Example

```
int add(int a, int b)
{
    int c;
    c = a + b;
    return c;
}
```

```
int main(int argc, char** argv)
{
    int sum = add(10, 20);
    printf("sum = %d\n", sum);
    return 0;
}
```

C program, x86-64 assembly

_add:		
0000000100000ea0	pushq	%rbp
0000000100000ea1	movq	%rsp,%rbp
0000000100000ea4	movl	%edi,0xec(%rbp)
0000000100000ea7	movl	%esi,0xe8(%rbp)
0000000100000eaa	movl	0xe8(%rbp),%eax
0000000100000ead	addl	0xec(%rbp),%eax
0000000100000eb0	movl	%eax,0xfc(%rbp)
0000000100000eb3	movl	0xfc(%rbp),%eax
0000000100000eb6	leave	
0000000100000eb7	ret	

_main:		
0000000100000eb8	pushq	%rbp
0000000100000eb9	movq	%rsp,%rbp
0000000100000ebc	subq	\$0x20,%rsp
0000000100000ec0	movl	%edi,0xec(%rbp)
0000000100000ec3	movq	%rsi,0xe0(%rbp)
0000000100000ec7	movl	\$0x00000014,%esi
0000000100000ecc	movl	\$0x0000000a,%edi
0000000100000ed1	callq	0x00000ea0
0000000100000ed6	movl	%eax,0xfc(%rbp)
0000000100000ed9	movl	0xfc(%rbp),%esi
0000000100000edc	leaq	0x00000041(%rip),%rdi
0000000100000ee3	movl	\$0x00000000,%eax
0000000100000ee8	callq	0x00000efa
0000000100000eed	movl	\$0x00000000,%eax
0000000100000ef2	leave	
0000000100000ef3	ret	

Calling Sequence Example

```
int add(int a, int b)
```

```
{
    int c;
    c = a + b;
    return c;
}
```

```
int main(int argc, char** argv)
```

```
{
    int sum = add(10, 20);
    printf("sum = %d\n", sum);
    return 0;
}
```

Call of `add()` from `main()`.

`_add:`

```
0000000100000ea0
0000000100000ea1
0000000100000ea4
0000000100000ea7
0000000100000eaa
0000000100000ead
0000000100000eb0
0000000100000eb3
0000000100000eb6
0000000100000eb7
```

`pushq %rbp`

`movq %rsp,%rbp`

`movl %edi,0xec(%rbp)`

`movl %esi,0xe8(%rbp)`

`movl 0xe8(%rbp),%eax`

`addl 0xec(%rbp),%eax`

`movl %eax,0xfc(%rbp)`

`movl 0xfc(%rbp),%eax`

`leave`

`ret`

`_main:`

```
0000000100000eb8
0000000100000eb9
0000000100000ebc
0000000100000ec0
0000000100000ec3
0000000100000ec7
0000000100000ecc
0000000100000ed1
0000000100000ed6
0000000100000ed9
0000000100000edc
0000000100000ee3
0000000100000ee8
0000000100000eed
0000000100000ef2
0000000100000ef3
```

`pushq %rbp`

`movq %rsp,%rbp`

`subq $0x20,%rsp`

`movl %edi,0xec(%rbp)`

`movq %rsi,0xe0(%rbp)`

`movl $0x00000014,%esi`

`movl $0x0000000a,%edi`

`callq 0x00000ea0`

`movl %eax,0xfc(%rbp)`

`movl 0xfc(%rbp),%esi`

`leaq 0x00000041(%rip),%rdi`

`movl $0x00000000,%eax`

`0x0000efa`

`movl $0x00000000,%eax`

`leave`

`ret`

Prologue

push stack & load arguments

```
int add(int a, int b)
{
    int c;
    c = a + b;
    return c;
}

int main(int argc, char** argv)
{
    int sum = add(10, 20);
    printf("sum = %d\n", sum);
    return 0;
}
```

C program, x86-64 assembly

_add:

0000000100000ea0	pushq	%rbp
0000000100000ea1	movq	%rsp,%rbp
0000000100000ea4	movl	%edi,0xec(%rbp)
0000000100000ea7	movl	%esi,0xe8(%rbp)
0000000100000eaa	movl	0xe8(%rbp),%eax
0000000100000ead	addl	0xec(%rbp),%eax
0000000100000eb0	movl	%eax,0xfc(%rbp)
0000000100000eb3	movl	0xfc(%rbp),%eax
0000000100000eb6	leave	
0000000100000eb7	ret	

_main:

0000000100000eb8	pushq	%rbp
0000000100000eb9	movq	%rsp,%rbp
0000000100000ebc	subq	\$0x20,%rsp
0000000100000ec0	movl	%edi,0xec(%rbp)
0000000100000ec3	movq	%rsi,0xe0(%rbp)
0000000100000ec7	movl	\$0x00000014,%esi
0000000100000ecc	movl	\$0x000000a,%edi
0000000100000ed1	callq	0x00000ea0
0000000100000ed6	movl	%eax,0xfc(%rbp)
0000000100000ed9	movl	0xfc(%rbp),%esi
0000000100000edc	leaq	0x00000041(%rip),%rdi
0000000100000ee3	movl	\$0x00000000,%eax
0000000100000ee8	callq	0x00000efa
0000000100000eed	movl	\$0x00000000,%eax
0000000100000ef2	leave	
0000000100000ef3	ret	

Prologue

push stack & load arguments

```
int add(int a, int b)
{
    int c;
    c = a + b;
    return c;
}
```

```
int main(int argc, char** argv)
{
    int sum = add(10, 20);
    printf("sum = %d\n", sum);
    return 0;
}
```

C program. x86-64 assembly

Epilogue

setup return value & restore stack

_add:

0000000100000ea0	pushq	%rbp
0000000100000ea1	movq	%rsp,%rbp
0000000100000ea4	movl	%edi,0xec(%rbp)
0000000100000ea7	movl	%esi,0xe8(%rbp)
0000000100000eaa	movl	0xe8(%rbp),%eax
0000000100000ead	addl	0xec(%rbp),%eax
0000000100000eb0	movl	%eax,0xfc(%rbp)
0000000100000eb3	movl	0xfc(%rbp),%eax
0000000100000eb6	leave	
0000000100000eb7	ret	

_main:

0000000100000eb8	pushq	%rbp
0000000100000eb9	movq	%rsp,%rbp
0000000100000ebc	subq	\$0x20,%rsp
0000000100000ec0	movl	%edi,0xec(%rbp)
0000000100000ec3	movq	%rsi,0xe0(%rbp)
0000000100000ec7	movl	\$0x00000014,%esi
0000000100000ecc	movl	\$0x000000a,%edi
0000000100000ed1	callq	0x00000ea0
0000000100000ed6	movl	%eax,0xfc(%rbp)
0000000100000ed9	movl	0xfc(%rbp),%esi
0000000100000edc	leaq	0x00000041(%rip),%rdi
0000000100000ee3	movl	\$0x00000000,%eax
0000000100000ee8	callq	0x00000efa
0000000100000eed	movl	\$0x00000000,%eax
0000000100000ef2	leave	
0000000100000ef3	ret	

Prologue

push stack & load arguments

```
int add(int a, int b)
{
    int c;
    c = a + b;
    return c;
```

Call Setup

prepare arguments

```
int main(int argc, char** argv)
{
    int sum = add(10, 20);
    printf("sum = %d\n", sum);
    return 0;
}
```

C program. x86-64 assembly

Epilogue

setup return value & restore stack

_add:

0000000100000ea0	pushq	%rbp
0000000100000ea1	movq	%rsp,%rbp
0000000100000ea4	movl	%edi,0xec(%rbp)
0000000100000ea7	movl	%esi,0xe8(%rbp)
0000000100000eaa	movl	0xe8(%rbp),%eax
0000000100000ead	addl	0xec(%rbp),%eax
0000000100000eb0	movl	%eax,0xfc(%rbp)
0000000100000eb3	movl	0xfc(%rbp),%eax
0000000100000eb6	leave	
0000000100000eb7	ret	

_main:

0000000100000eb8	pushq	%rbp
0000000100000eb9	movq	%rsp,%rbp
0000000100000ebc	subq	\$0x20,%rsp
0000000100000ec0	movl	%edi,0xec(%rbp)
0000000100000ec3	movq	%rsi,0xe0(%rbp)
0000000100000ec7	movl	\$0x00000014,%esi
0000000100000ecc	movl	\$0x0000000a,%edi
0000000100000ed1	callq	0x00000ea0
0000000100000ed6	movl	%eax,0xfc(%rbp)
0000000100000ed9	movl	0xfc(%rbp),%esi
0000000100000edc	leaq	0x00000041(%rip),%rdi
0000000100000ee3	movl	\$0x00000000,%eax
0000000100000ee8	callq	0x00000efa
0000000100000eed	movl	\$0x00000000,%eax
0000000100000ef2	leave	
0000000100000ef3	ret	

Prologue

push stack & load arguments

```
int add(int a, int b)
{
    int c;
    c = a + b;
    return c;
```

Call Setup

prepare arguments

```
int main(int argc, char** argv)
{
    int sum = add(10, 20);
    printf("sum = %d\n", sum);
    return 0;
}
```

Cleanup

save return value

Epilogue

setup return value & restore stack

_add:

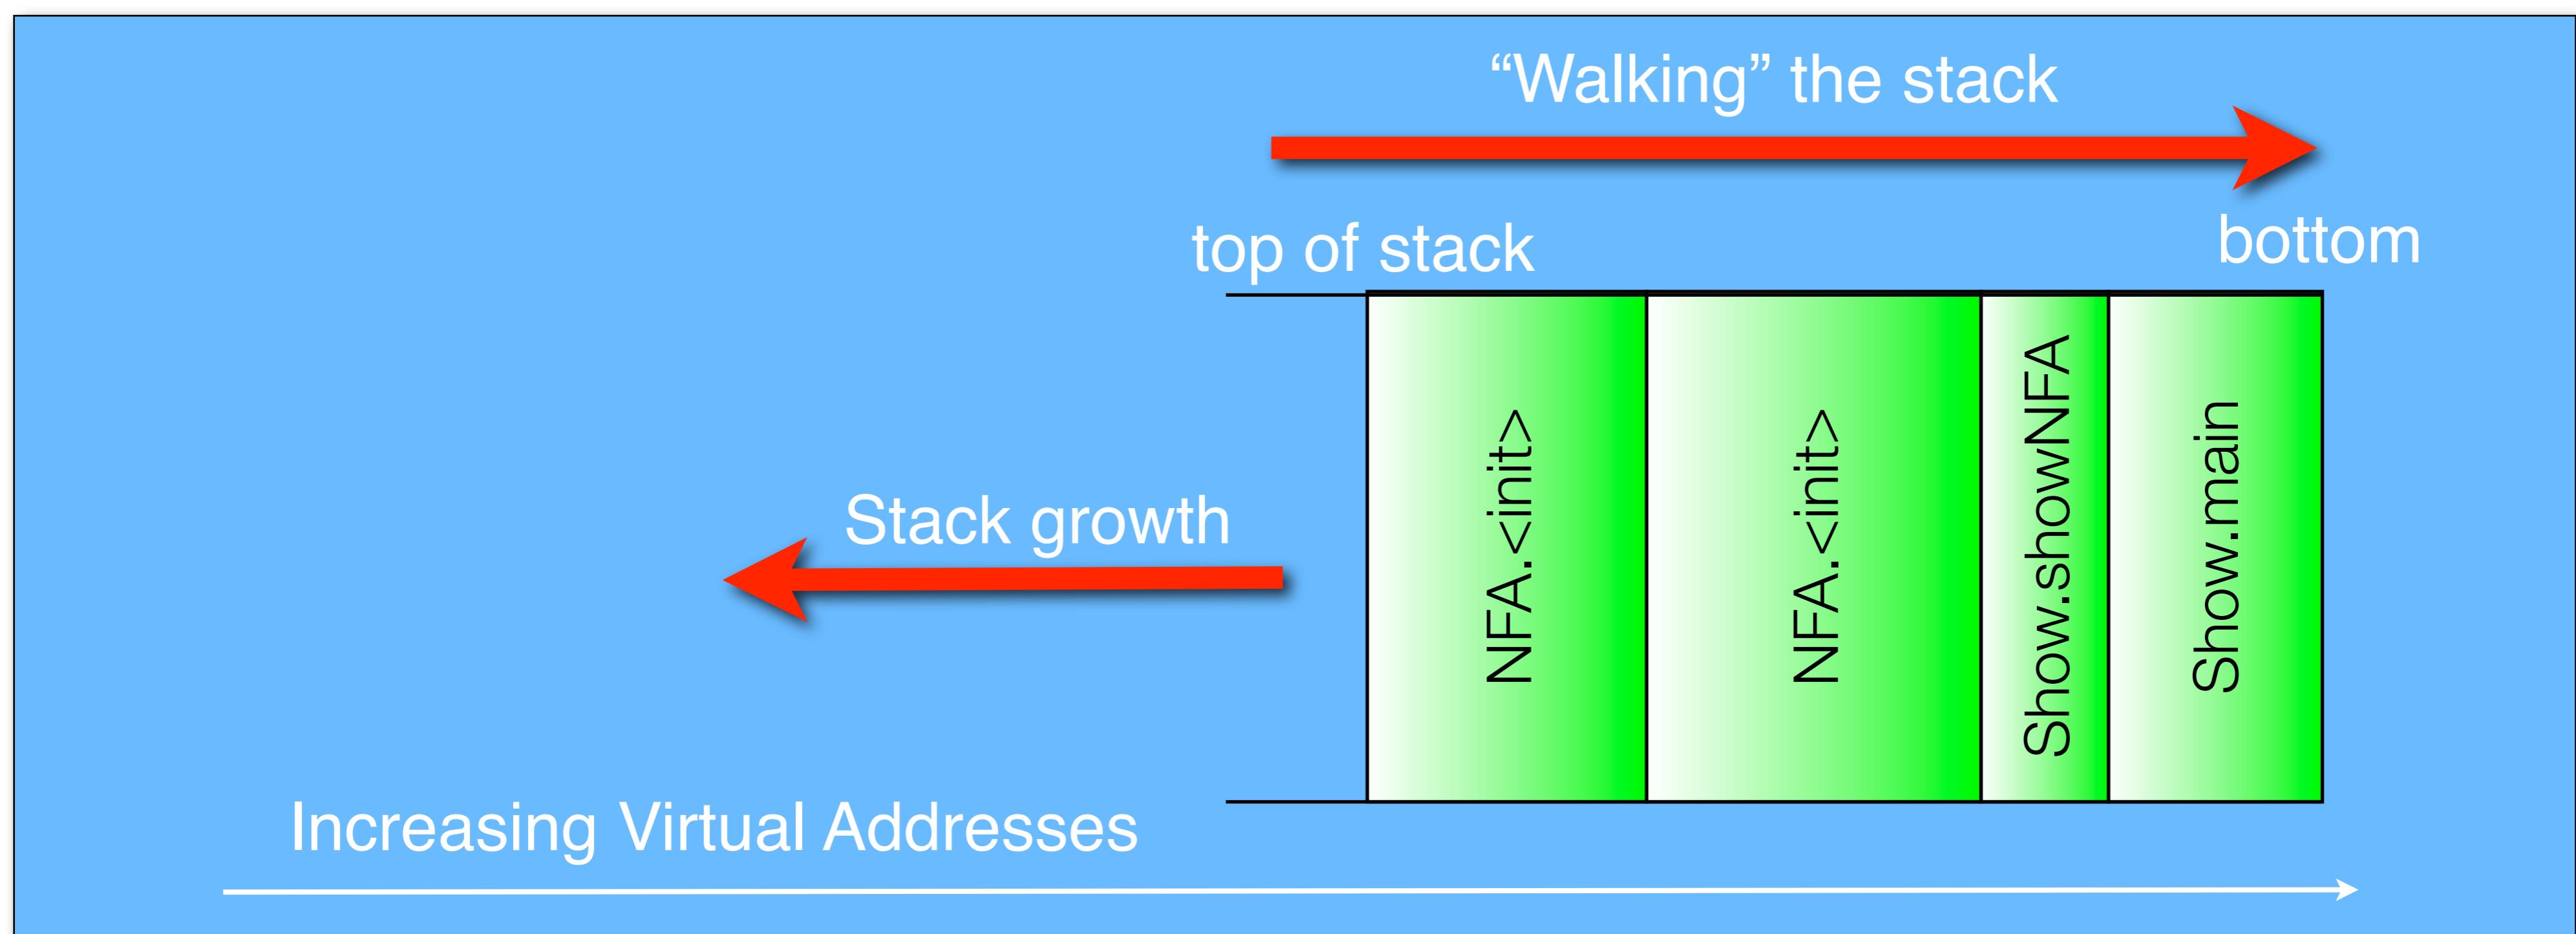
0000000100000ea0	pushq	%rbp
0000000100000ea1	movq	%rsp,%rbp
0000000100000ea4	movl	%edi,0xec(%rbp)
0000000100000ea7	movl	%esi,0xe8(%rbp)
0000000100000eaa	movl	0xe8(%rbp),%eax
0000000100000ead	addl	0xec(%rbp),%eax
0000000100000eb0	movl	%eax,0xfc(%rbp)
0000000100000eb3	movl	0xfc(%rbp),%eax
0000000100000eb6	leave	
0000000100000eb7	ret	

_main:

0000000100000eb8	pushq	%rbp
0000000100000eb9	movq	%rsp,%rbp
0000000100000ebc	subq	\$0x20,%rsp
0000000100000ec0	movl	%edi,0xec(%rbp)
0000000100000ec3	movq	%rsi,0xe0(%rbp)
0000000100000ec7	movl	\$0x00000014,%esi
0000000100000ecc	movl	\$0x0000000a,%edi
0000000100000ed1	calla	0x00000ea0
0000000100000ed6	movl	%eax,0xfc(%rbp)
0000000100000ed9	movl	0xfc(%rbp),%esi
0000000100000edc	leaq	0x00000041(%rip),%rdi
0000000100000ee3	movl	\$0x00000000,%eax
0000000100000ee8	callq	0x00000efa
0000000100000eed	movl	\$0x00000000,%eax
0000000100000ef2	leave	
0000000100000ef3	ret	

Stack Trace

```
Exception in thread "main" NotYetImplementedException:  
at NFA.<init>(NFA.java:25)  
at NFA.<init>(NFA.java:16)  
at Show.showNFA(Show.java:68)  
at Show.main(Show.java:23)
```



Advantages & Disadvantages

Advantages.

- **Negligible** (in most cases) **runtime overhead**.
- **Efficient** use of space.
- **Recursion** possible.
- **Offset** of local variable **within frame** usually **constant**.

Limitations.

- Stack space is a **limited resource**.
- Stack frame **size fixed** (in many languages).
- Some **offset computations** required at runtime.
- **Object lifetime limited** to one subroutine invocation.

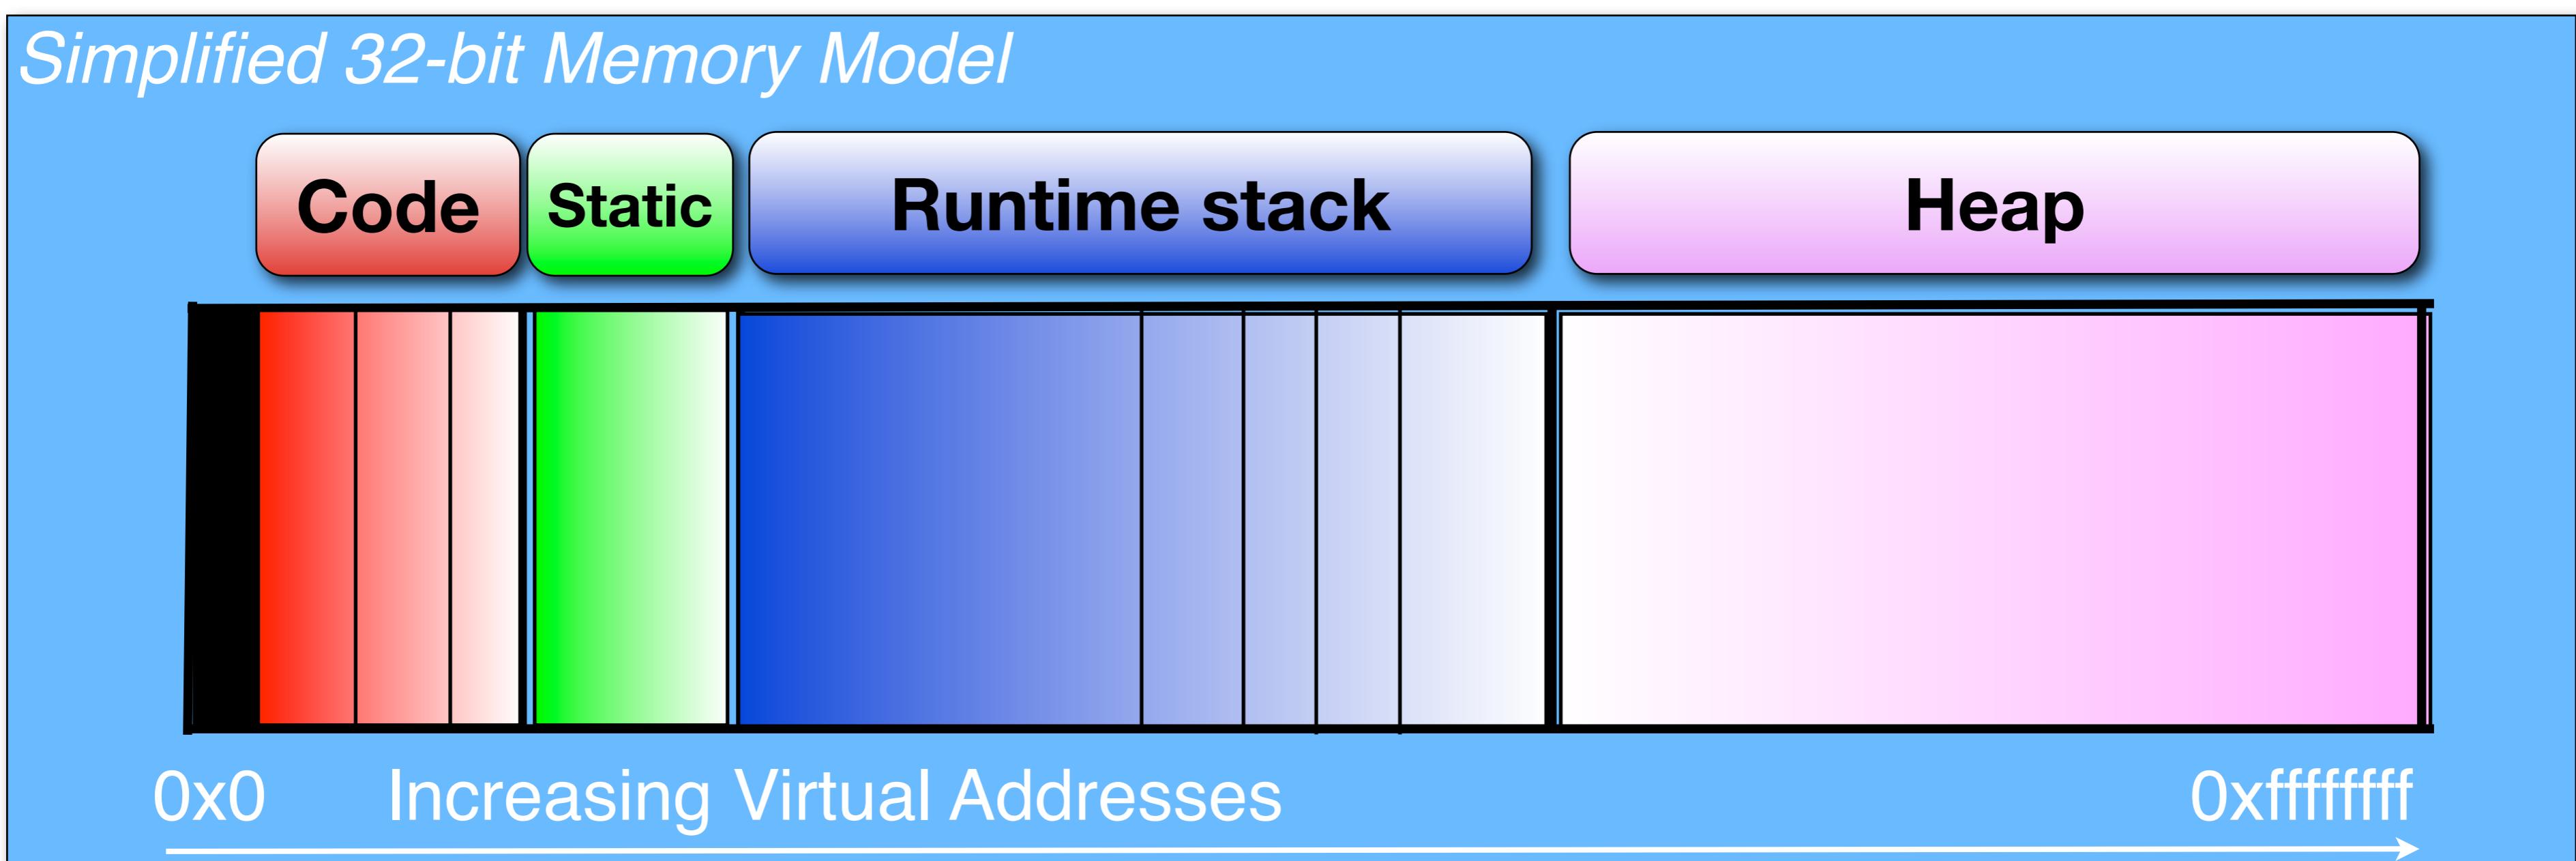
*Advice: use stack allocation when possible.
(except for large buffers)*

The Heap

(no relation to the data structure of the same name)

Arbitrary object lifetimes.

- Allocation and deallocation **at any point in time**.
- Can persists after subroutine completion.
- Very **flexible**; required for **dynamic allocations**.
- **Most expensive to manage**.



Memory Management

Allocation.

- Often explicit.
 - C++: **new**
- Compiler can generate implicit calls to **allocator**.
 - E.g., Prolog.

Deallocation.

- Often explicit.
 - C++: **delete**
- Sometimes done automatically.



Increasing Virtual Addresses

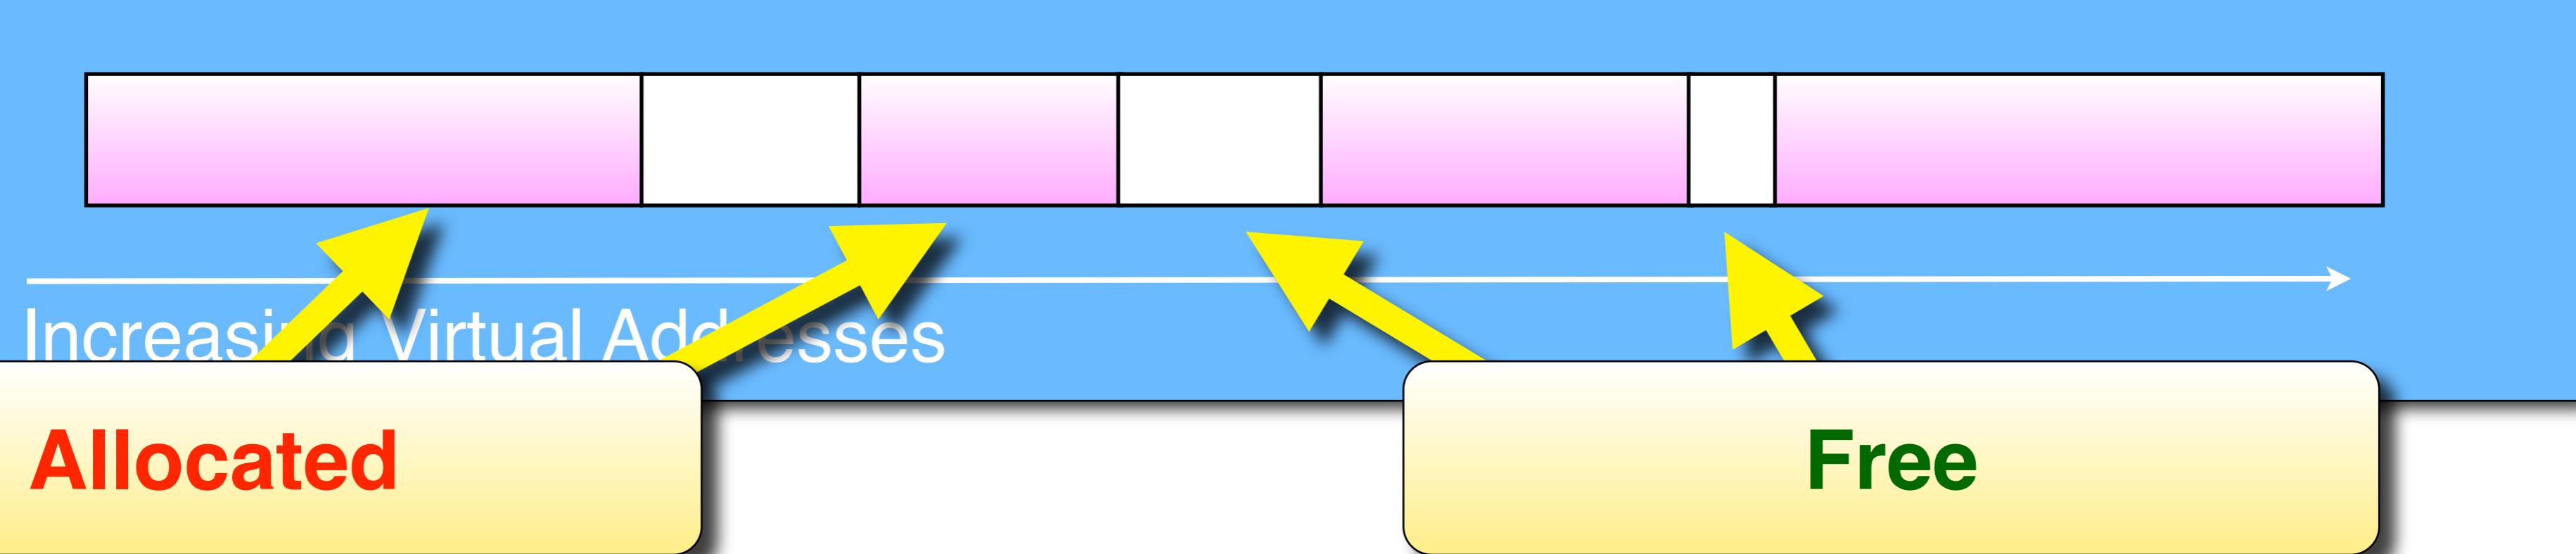
Memory Management

Allocation.

- Often explicit.
 - C++: **new**
- Compiler can generate implicit calls to **allocator**.
 - E.g., Prolog.

Deallocation.

- Often explicit.
 - C++: **delete**
- Sometimes done automatically.



Memory Management

Allocation.

- Often explicit.
 - C++: **new**
- Compiler can generate implicit calls to **allocator**.
 - E.g., Prolog

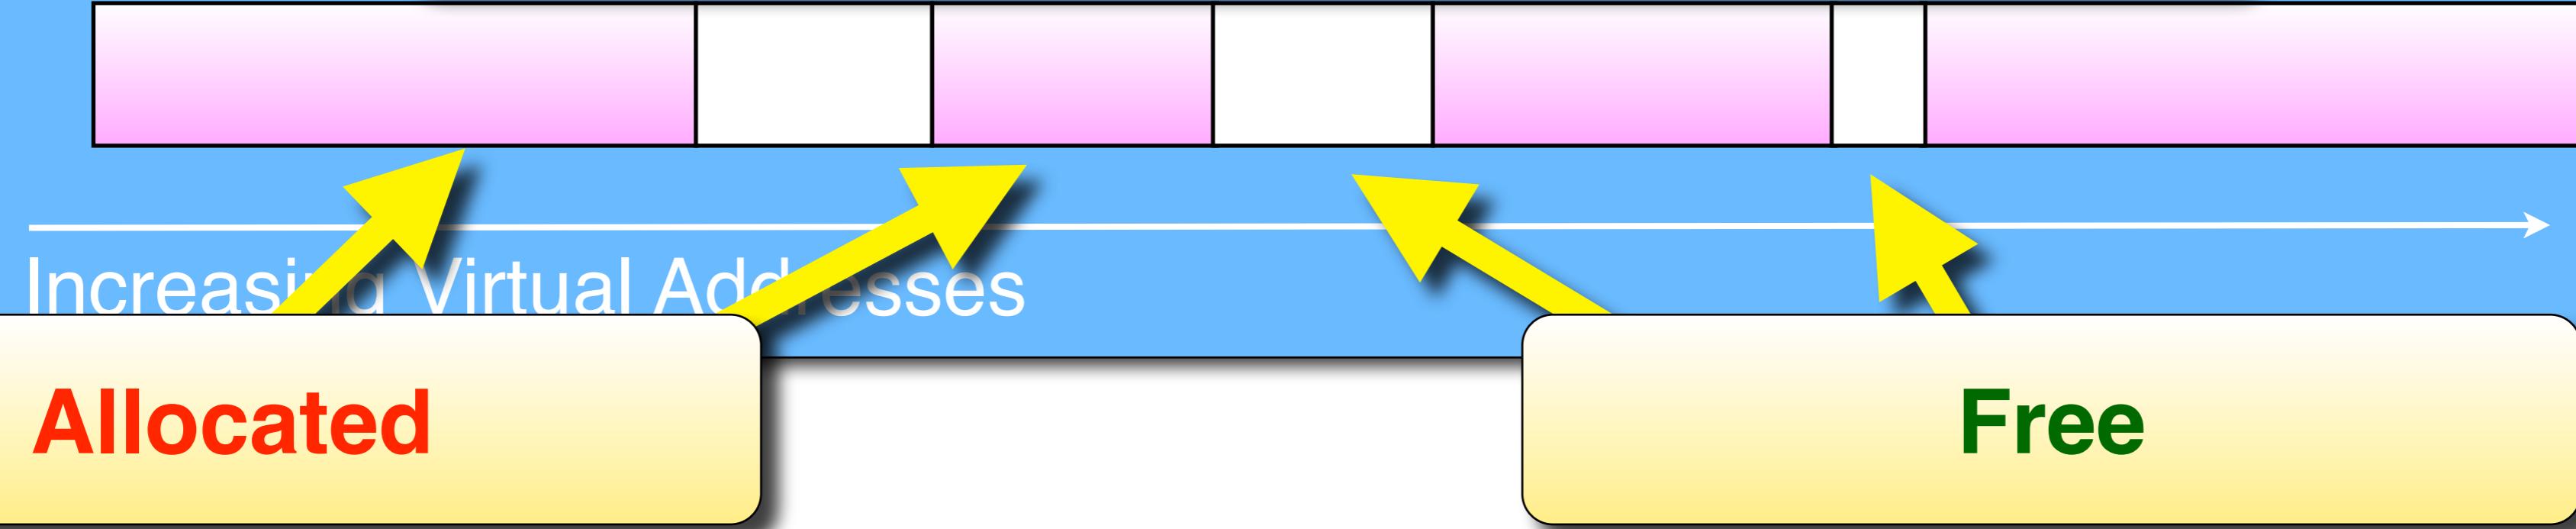
Deallocation.

- Often explicit.
 - C++: **delete**
- Sometimes done automatically.

Allocation Problem:

Given a size **S**, find a **contiguous** region of unallocated memory of length at least **S**.

(and be very, very **quick** about it)



Common Techniques

Allocator implementation.

→ Variable size.

- Heuristics: First-fit, best-fit, last-fit, worst-fit.
- List traversals, slow coalescing when deallocated.

→ Fixed-size blocks.

- 2^n or Fibonacci sequence.
- “Buddy allocator,” “slab allocator”
- Memory blocks are split until desired block size is reached.
- Quick coalescing: on free block is merged with its “buddy.”

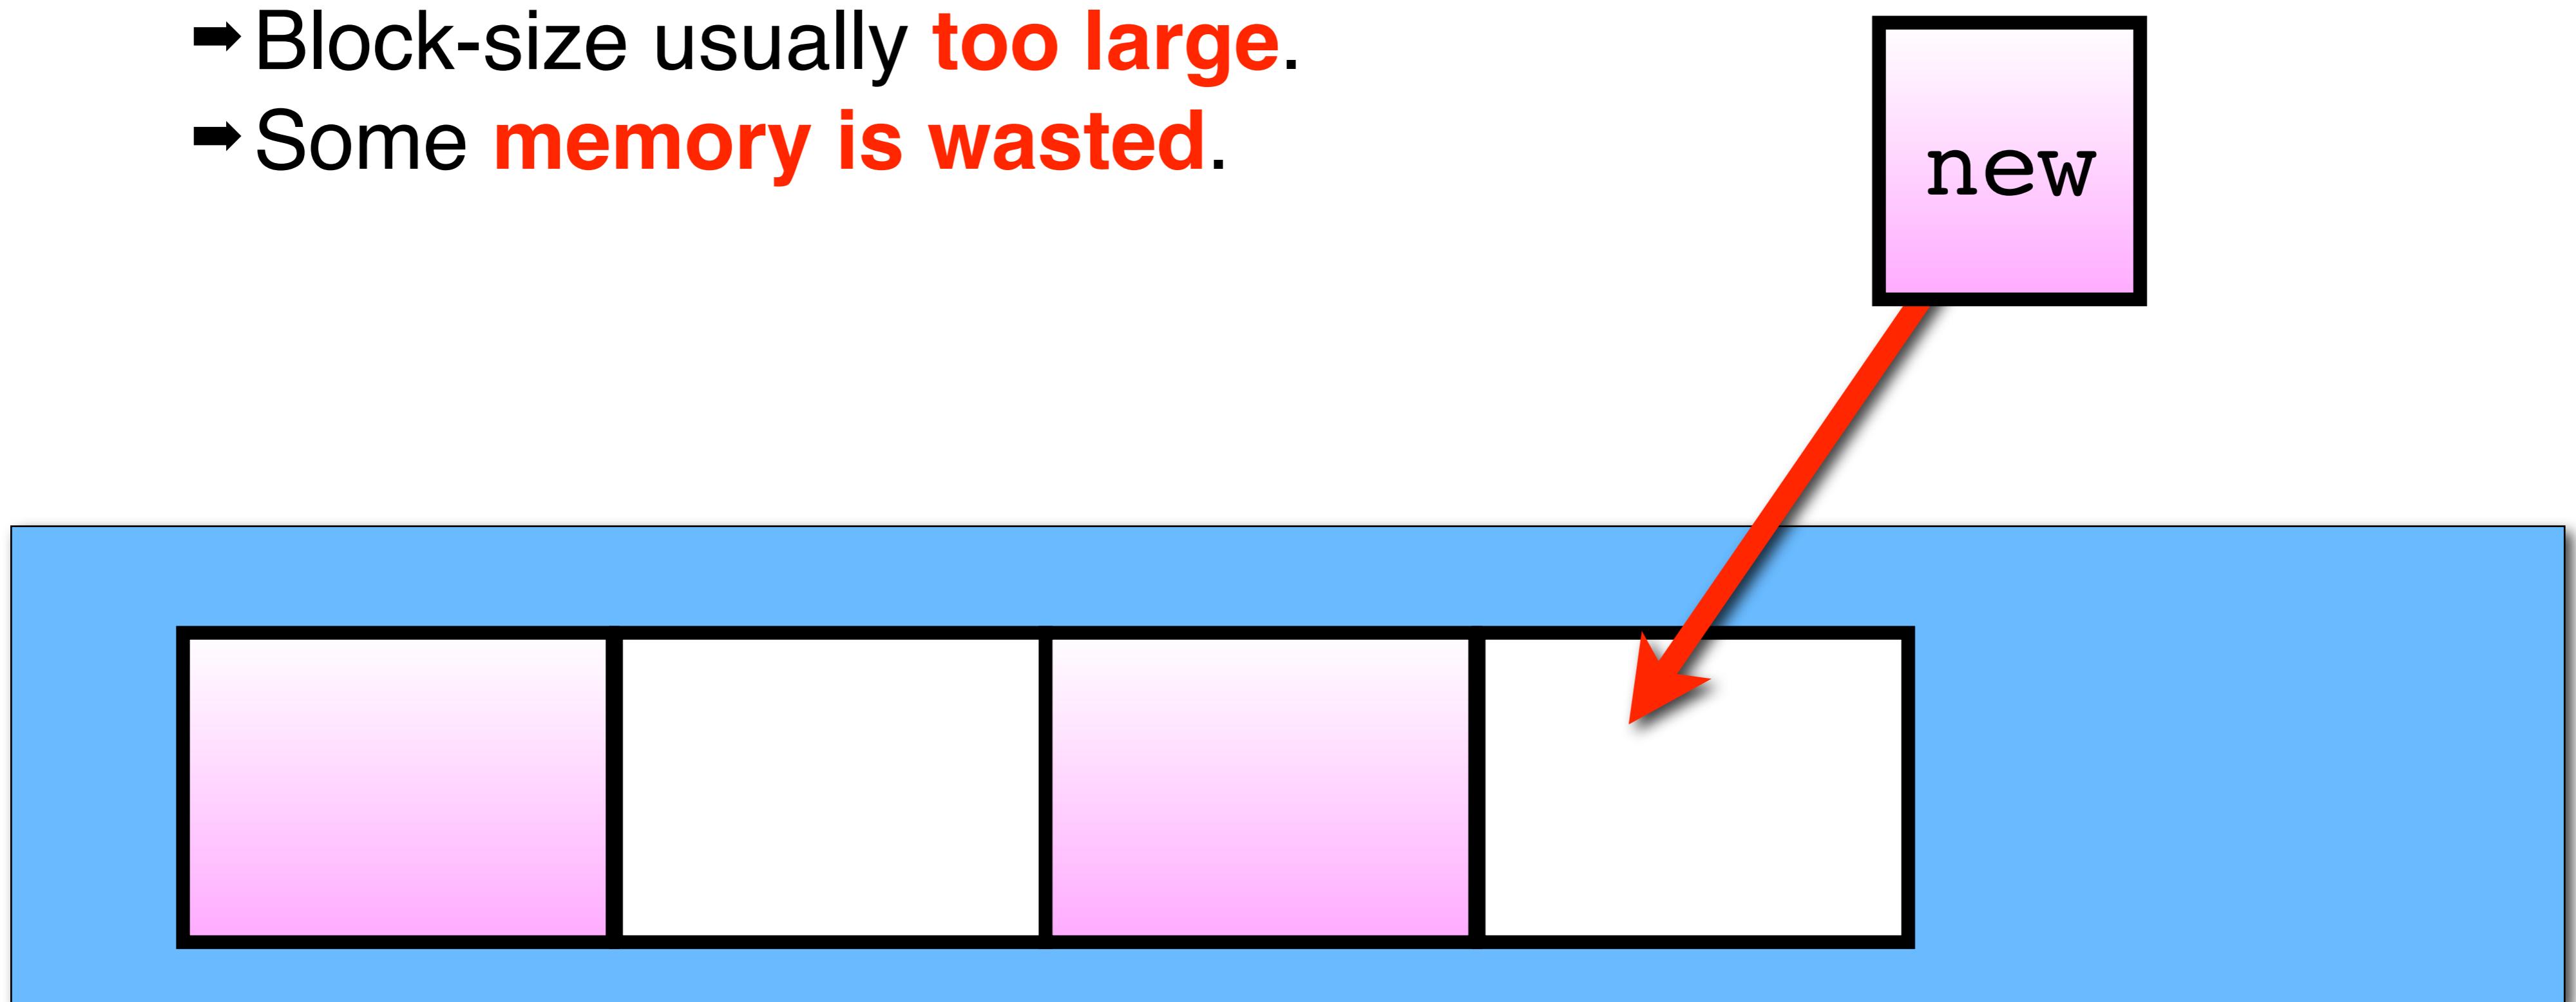
In practice.

- Most modern OSs use fixed-size blocks.
- Allocator performance crucial to many workloads.
- Allocators for multicore systems are still being researched.

Internal Fragmentation

Negative impact of **fixed-size blocks.**

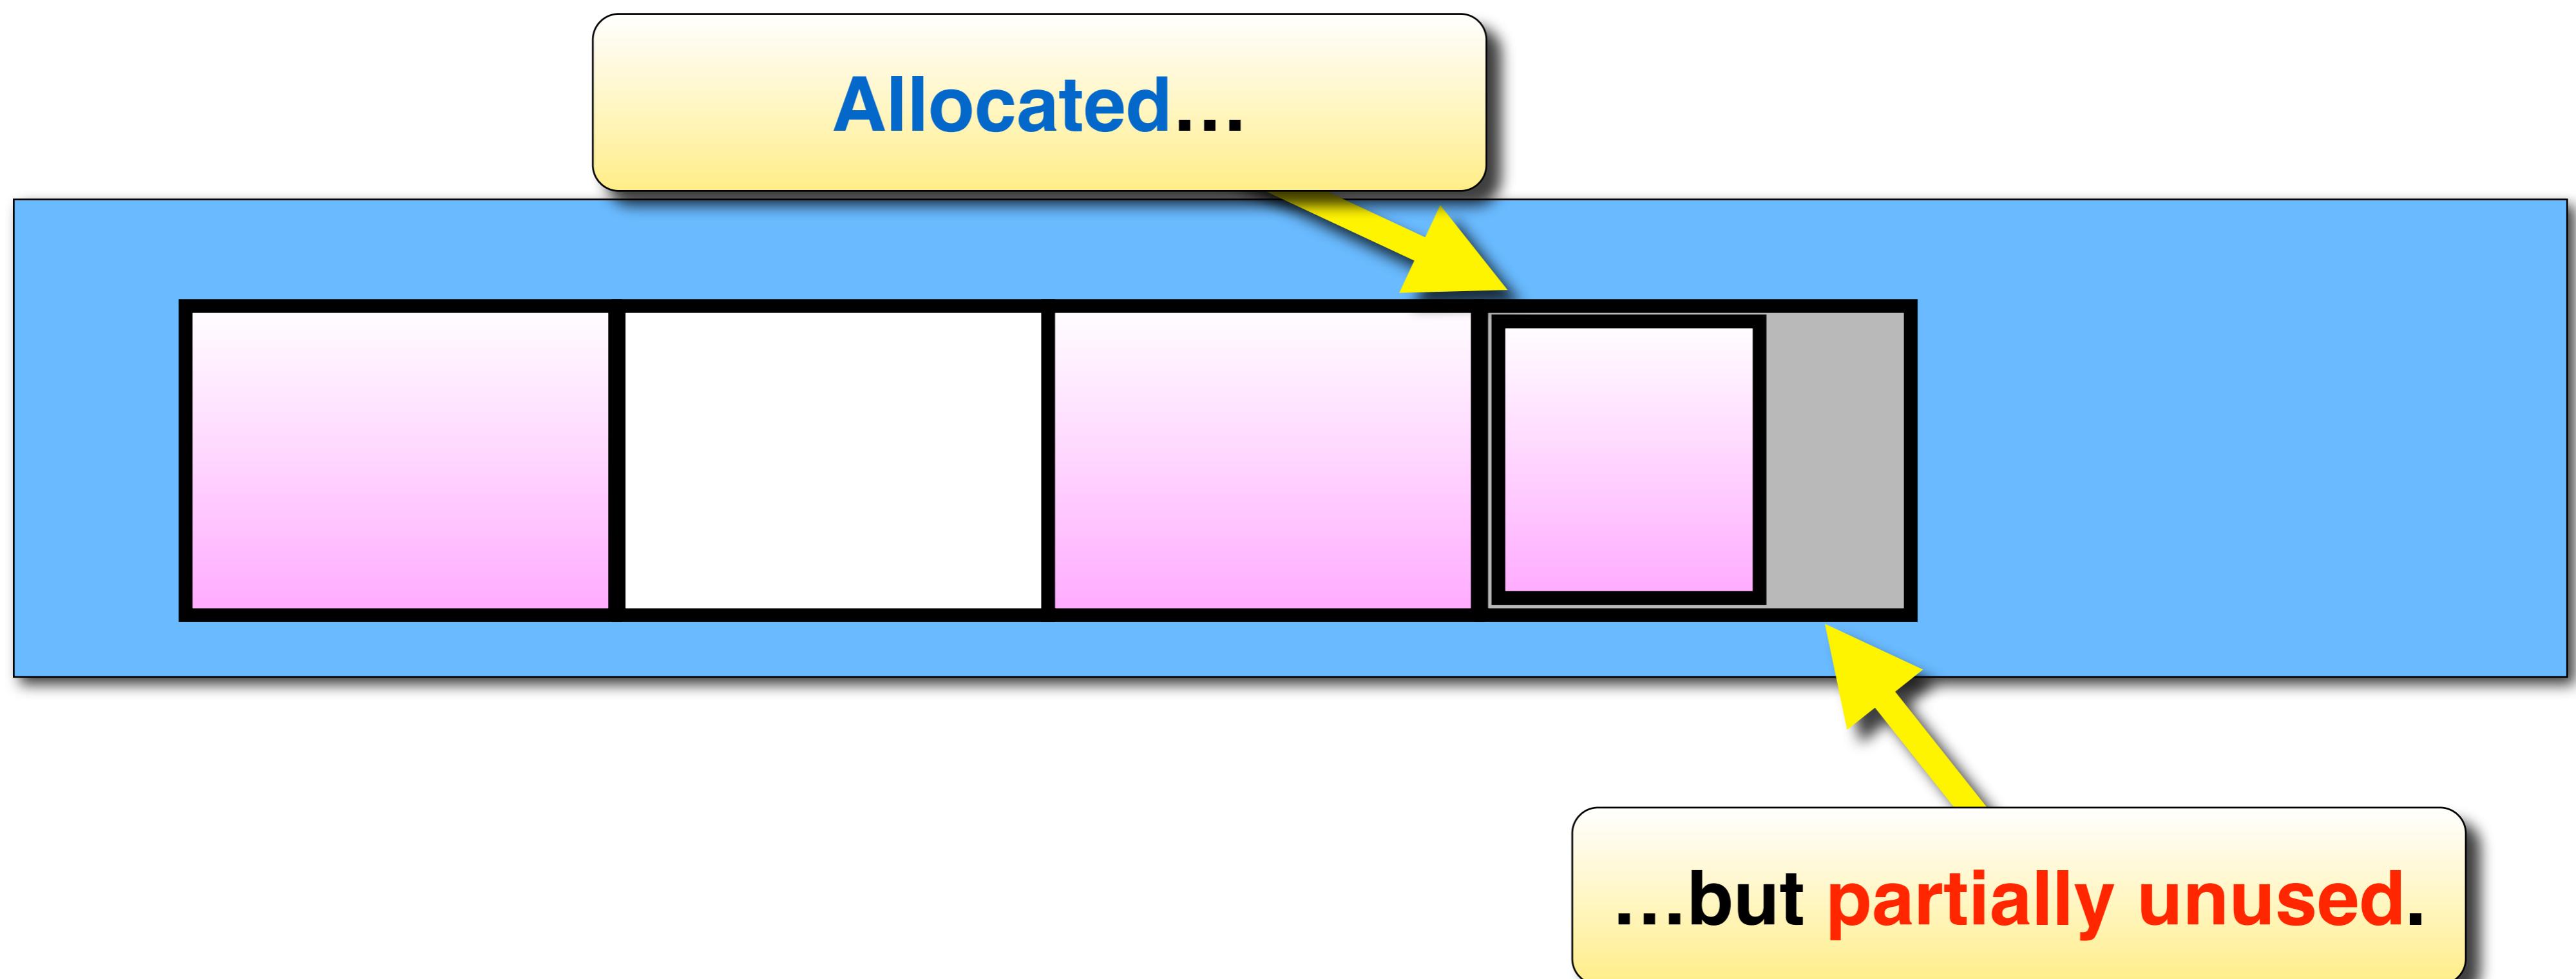
- Block-size usually **too large**.
- Some **memory is wasted**.



Internal Fragmentation

Negative impact of fixed-size blocks.

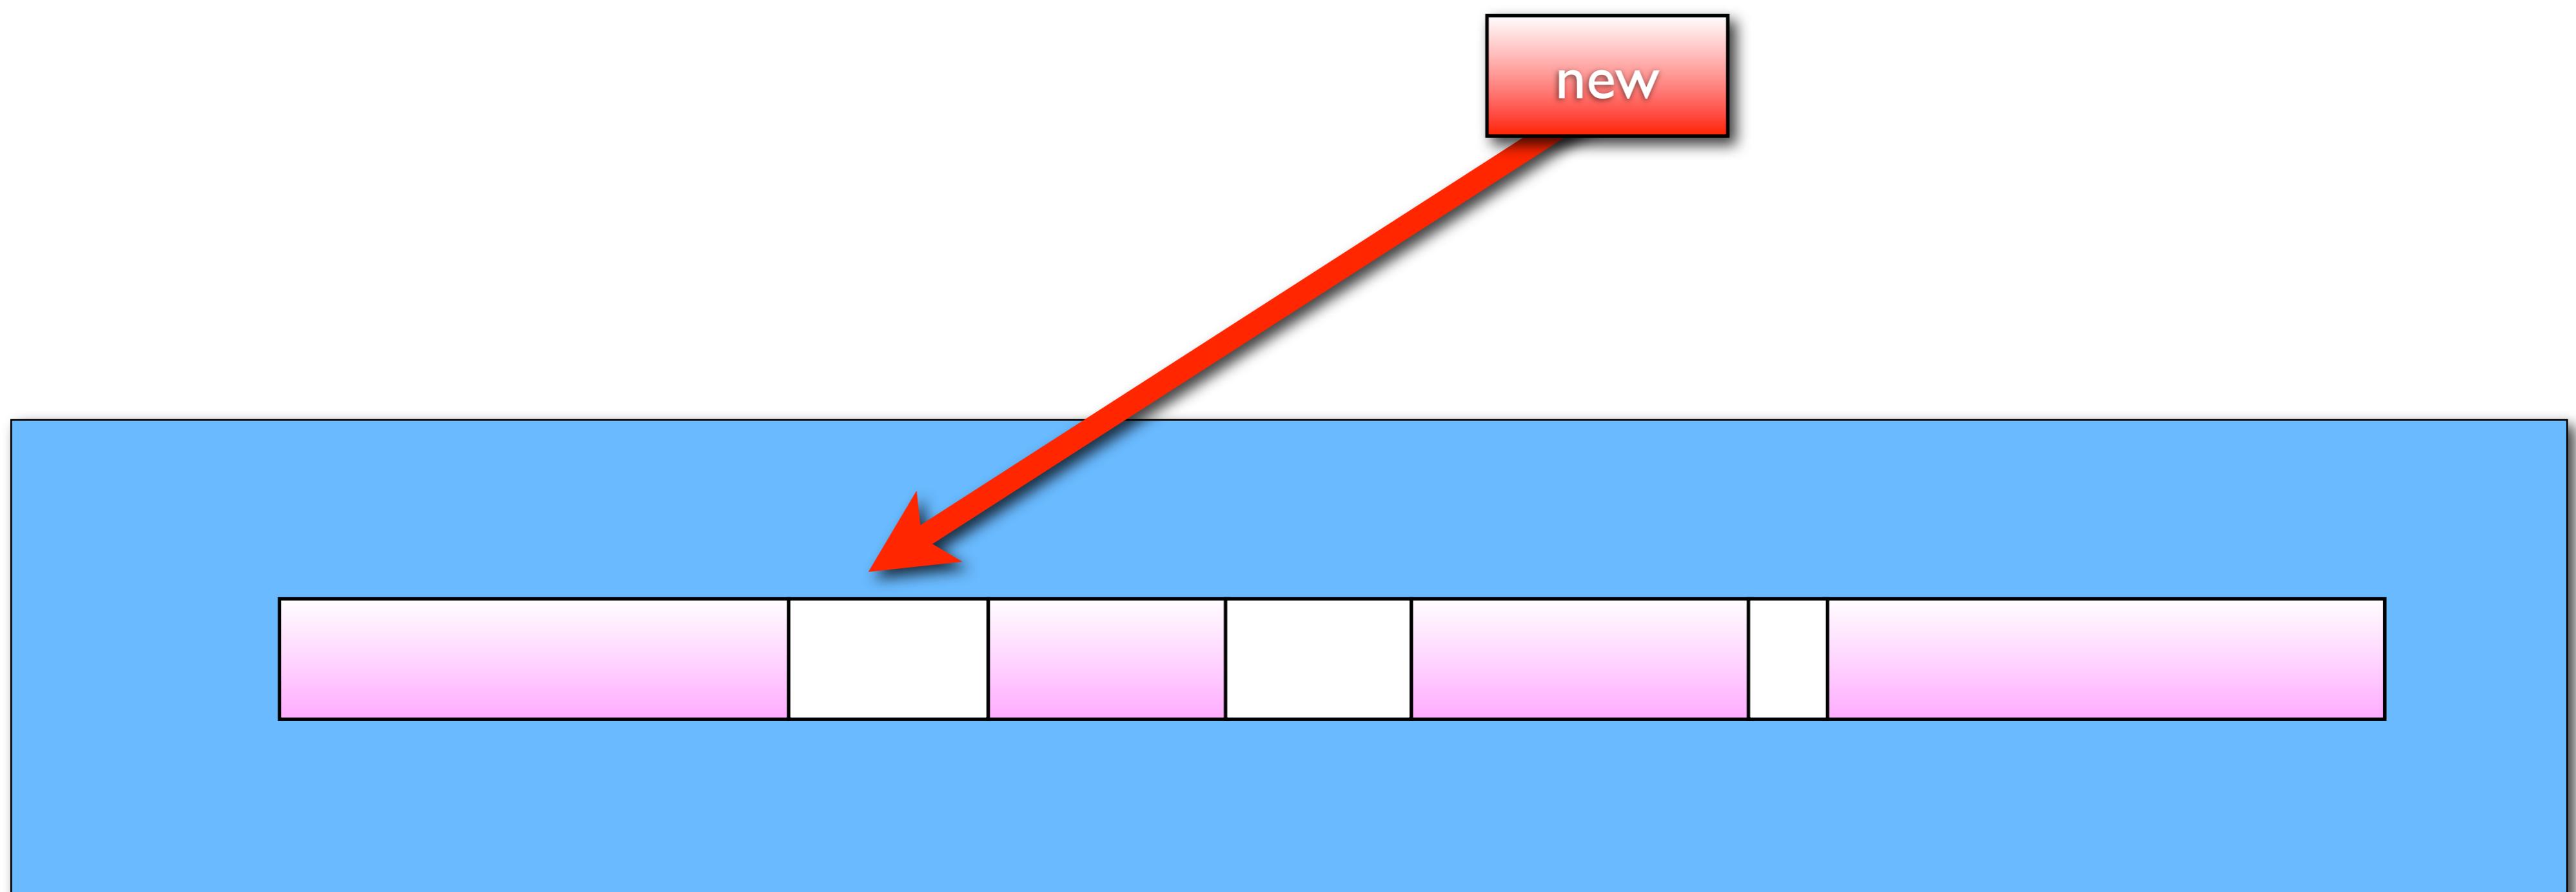
- Block-size usually **too large**.
- Some **memory is wasted**.



External Fragmentation

Non-contiguous free memory.

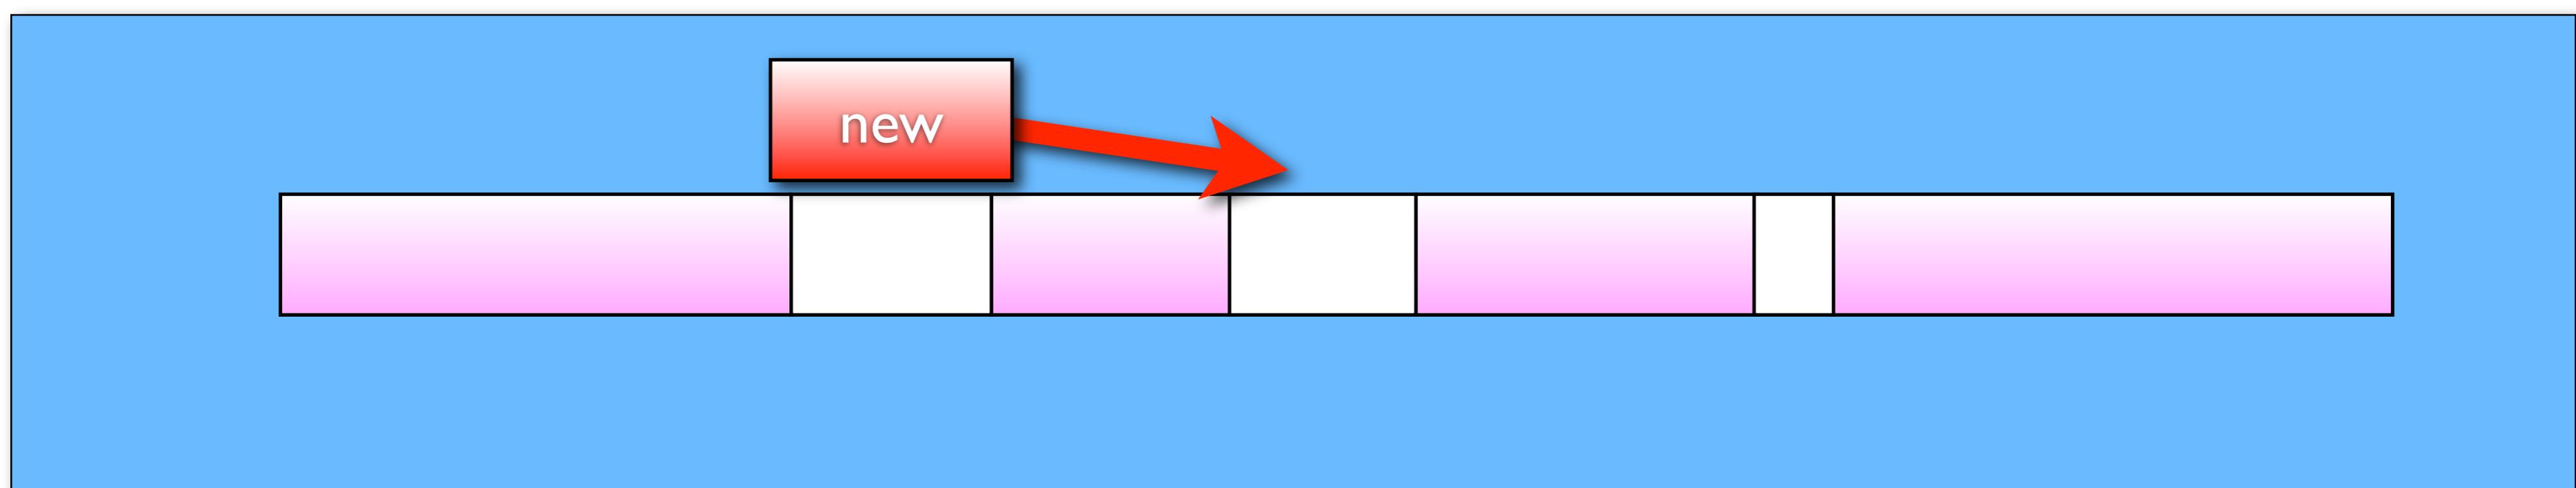
- In total, there is **sufficient available space...**
- ...but there is none of the free blocks is large enough by itself.



External Fragmentation

Non-contiguous free memory.

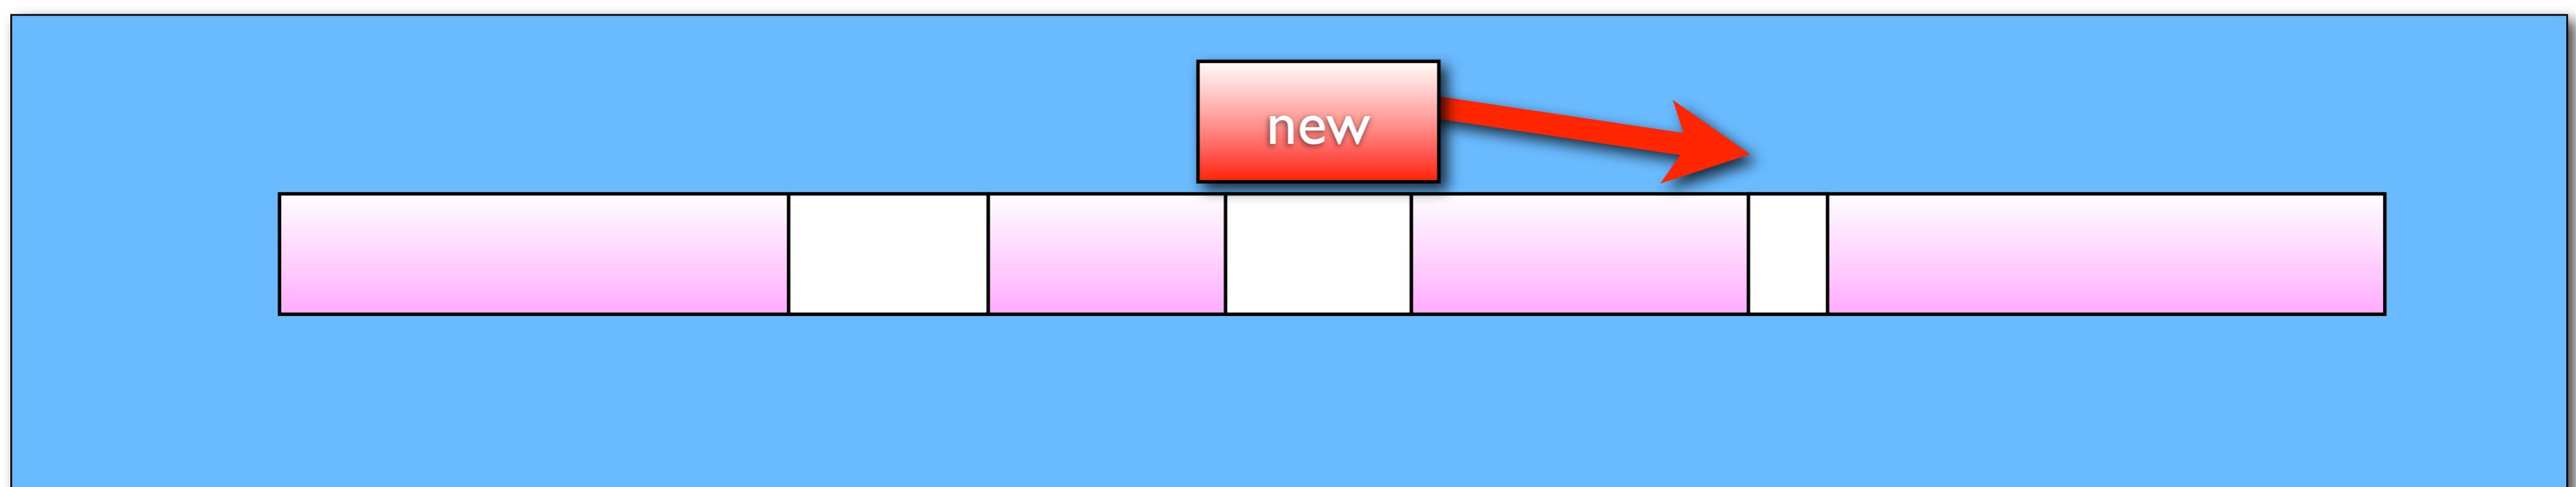
- In total, there is **sufficient available space...**
- ...but there is none of the free blocks is large enough by itself.



External Fragmentation

Non-contiguous free memory.

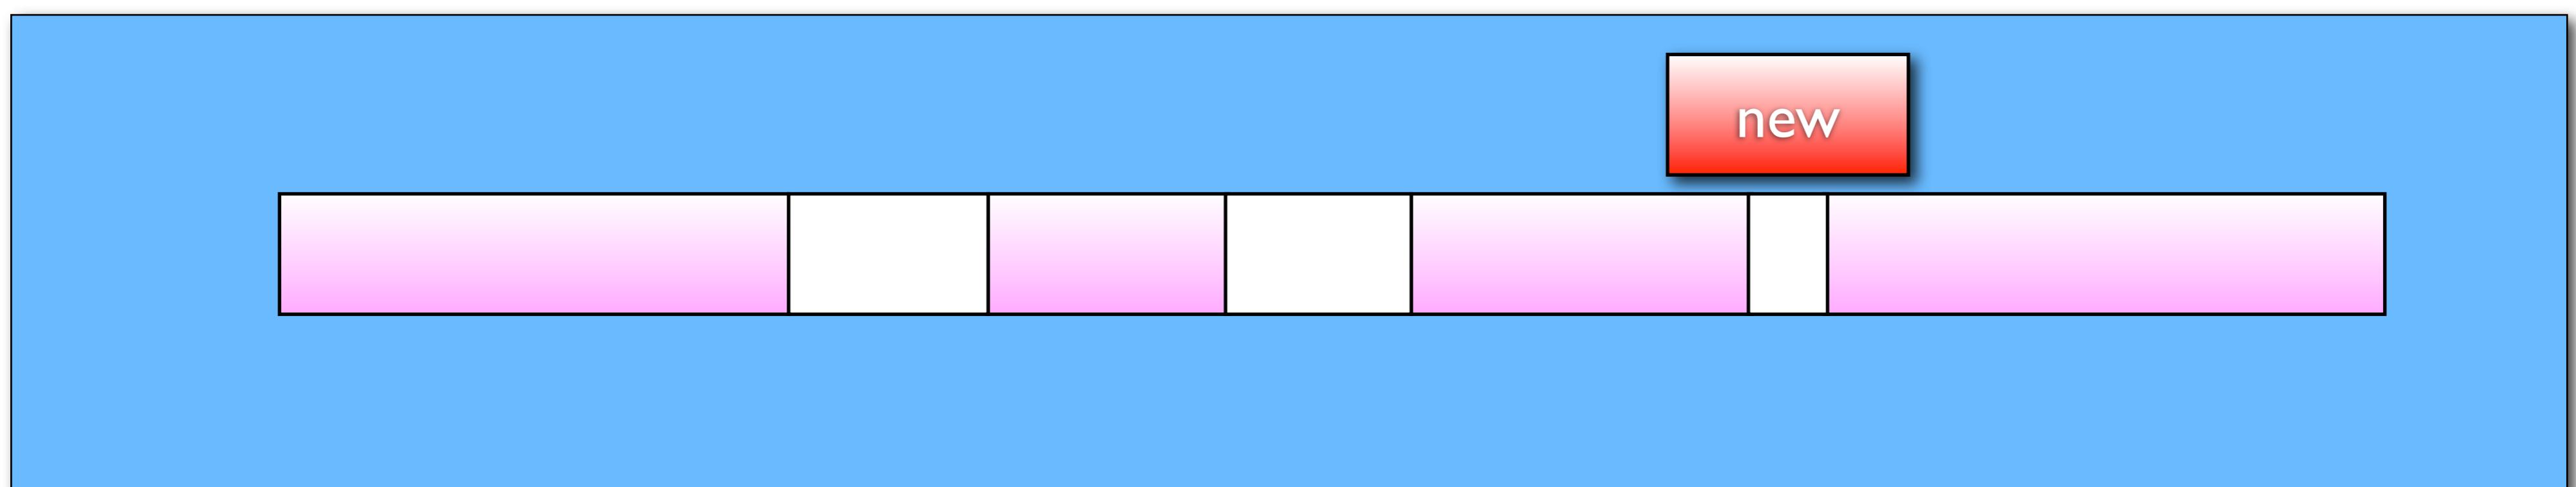
- In total, there is **sufficient available space...**
- ...but there is none of the free blocks is large enough by itself.



External Fragmentation

Non-contiguous free memory.

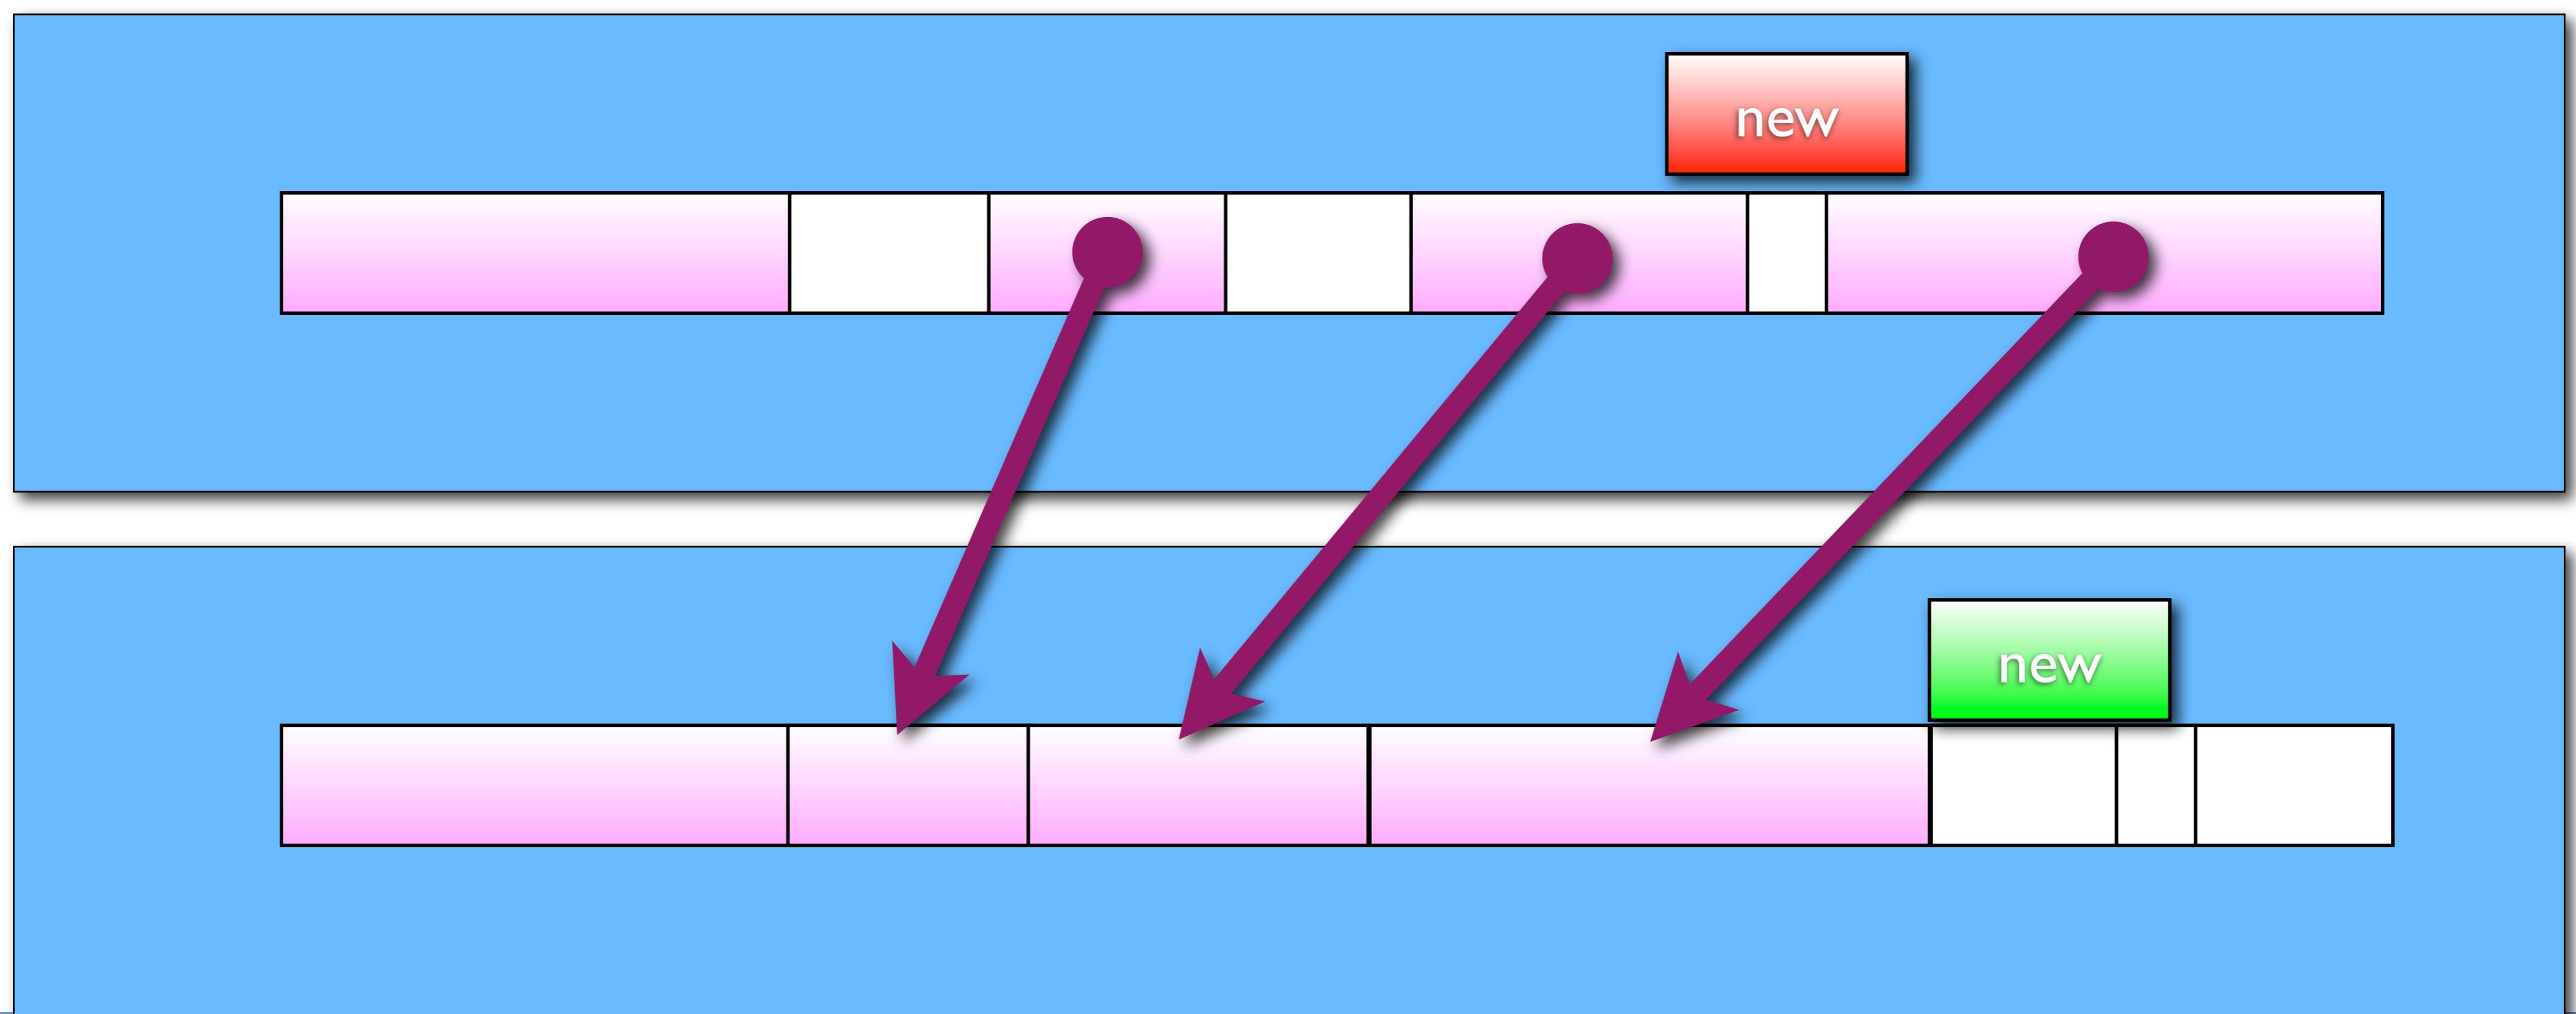
- In total, there is **sufficient available space...**
- ...but there is none of the free blocks is large enough by itself.



Compacting the Heap

Merge free space.

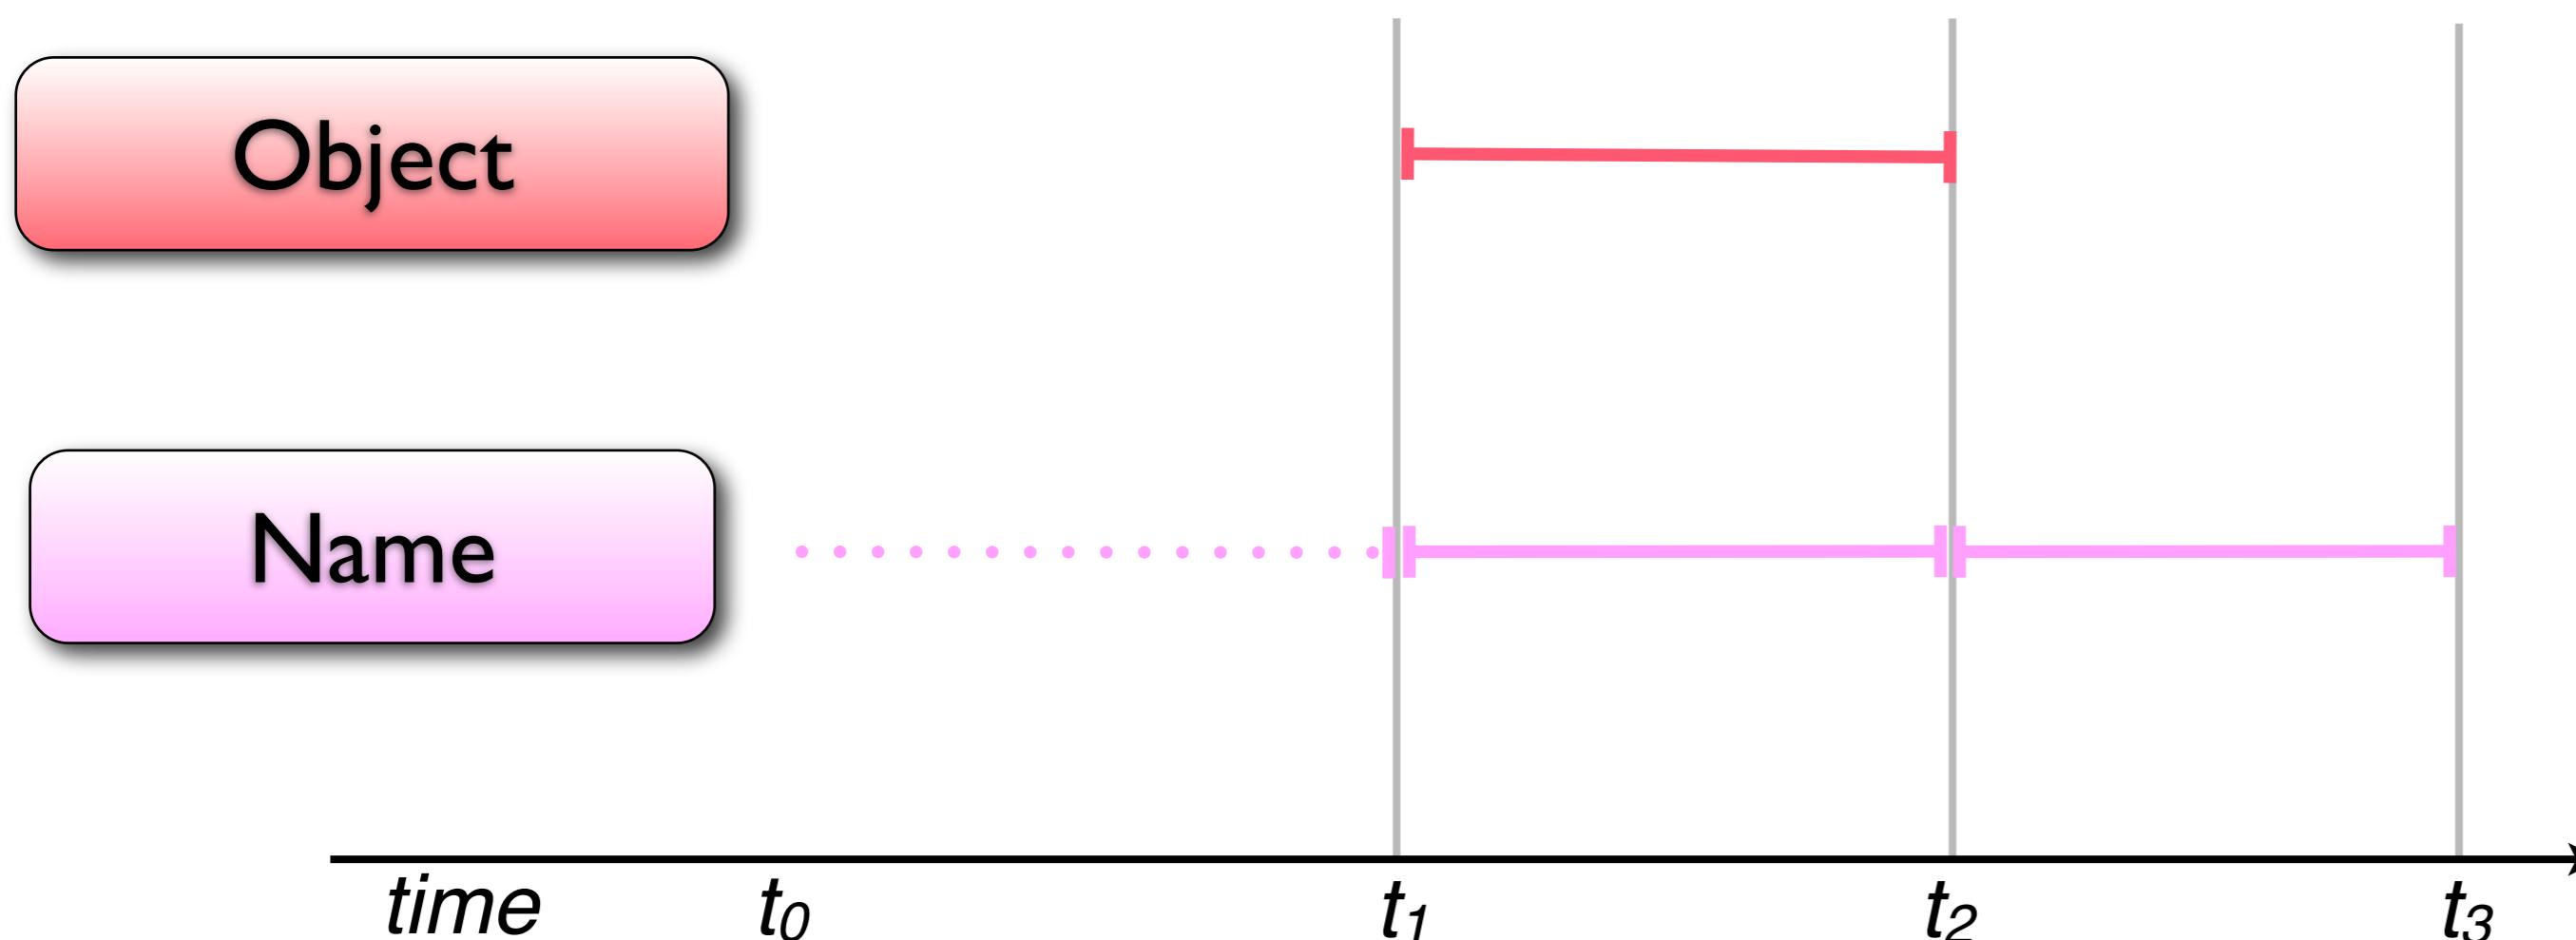
- **Copy** existing allocations & **update all references**.
- Very difficult to implement...



“Dangling” References

Binding / object lifetime mismatch.

- **Binding exists longer than object.**
- Object de-allocated too early; **access now illegal.**
- “Use-after-free bug”
(free is the C deallocation routine)
- “Dangling” pointer or reference.

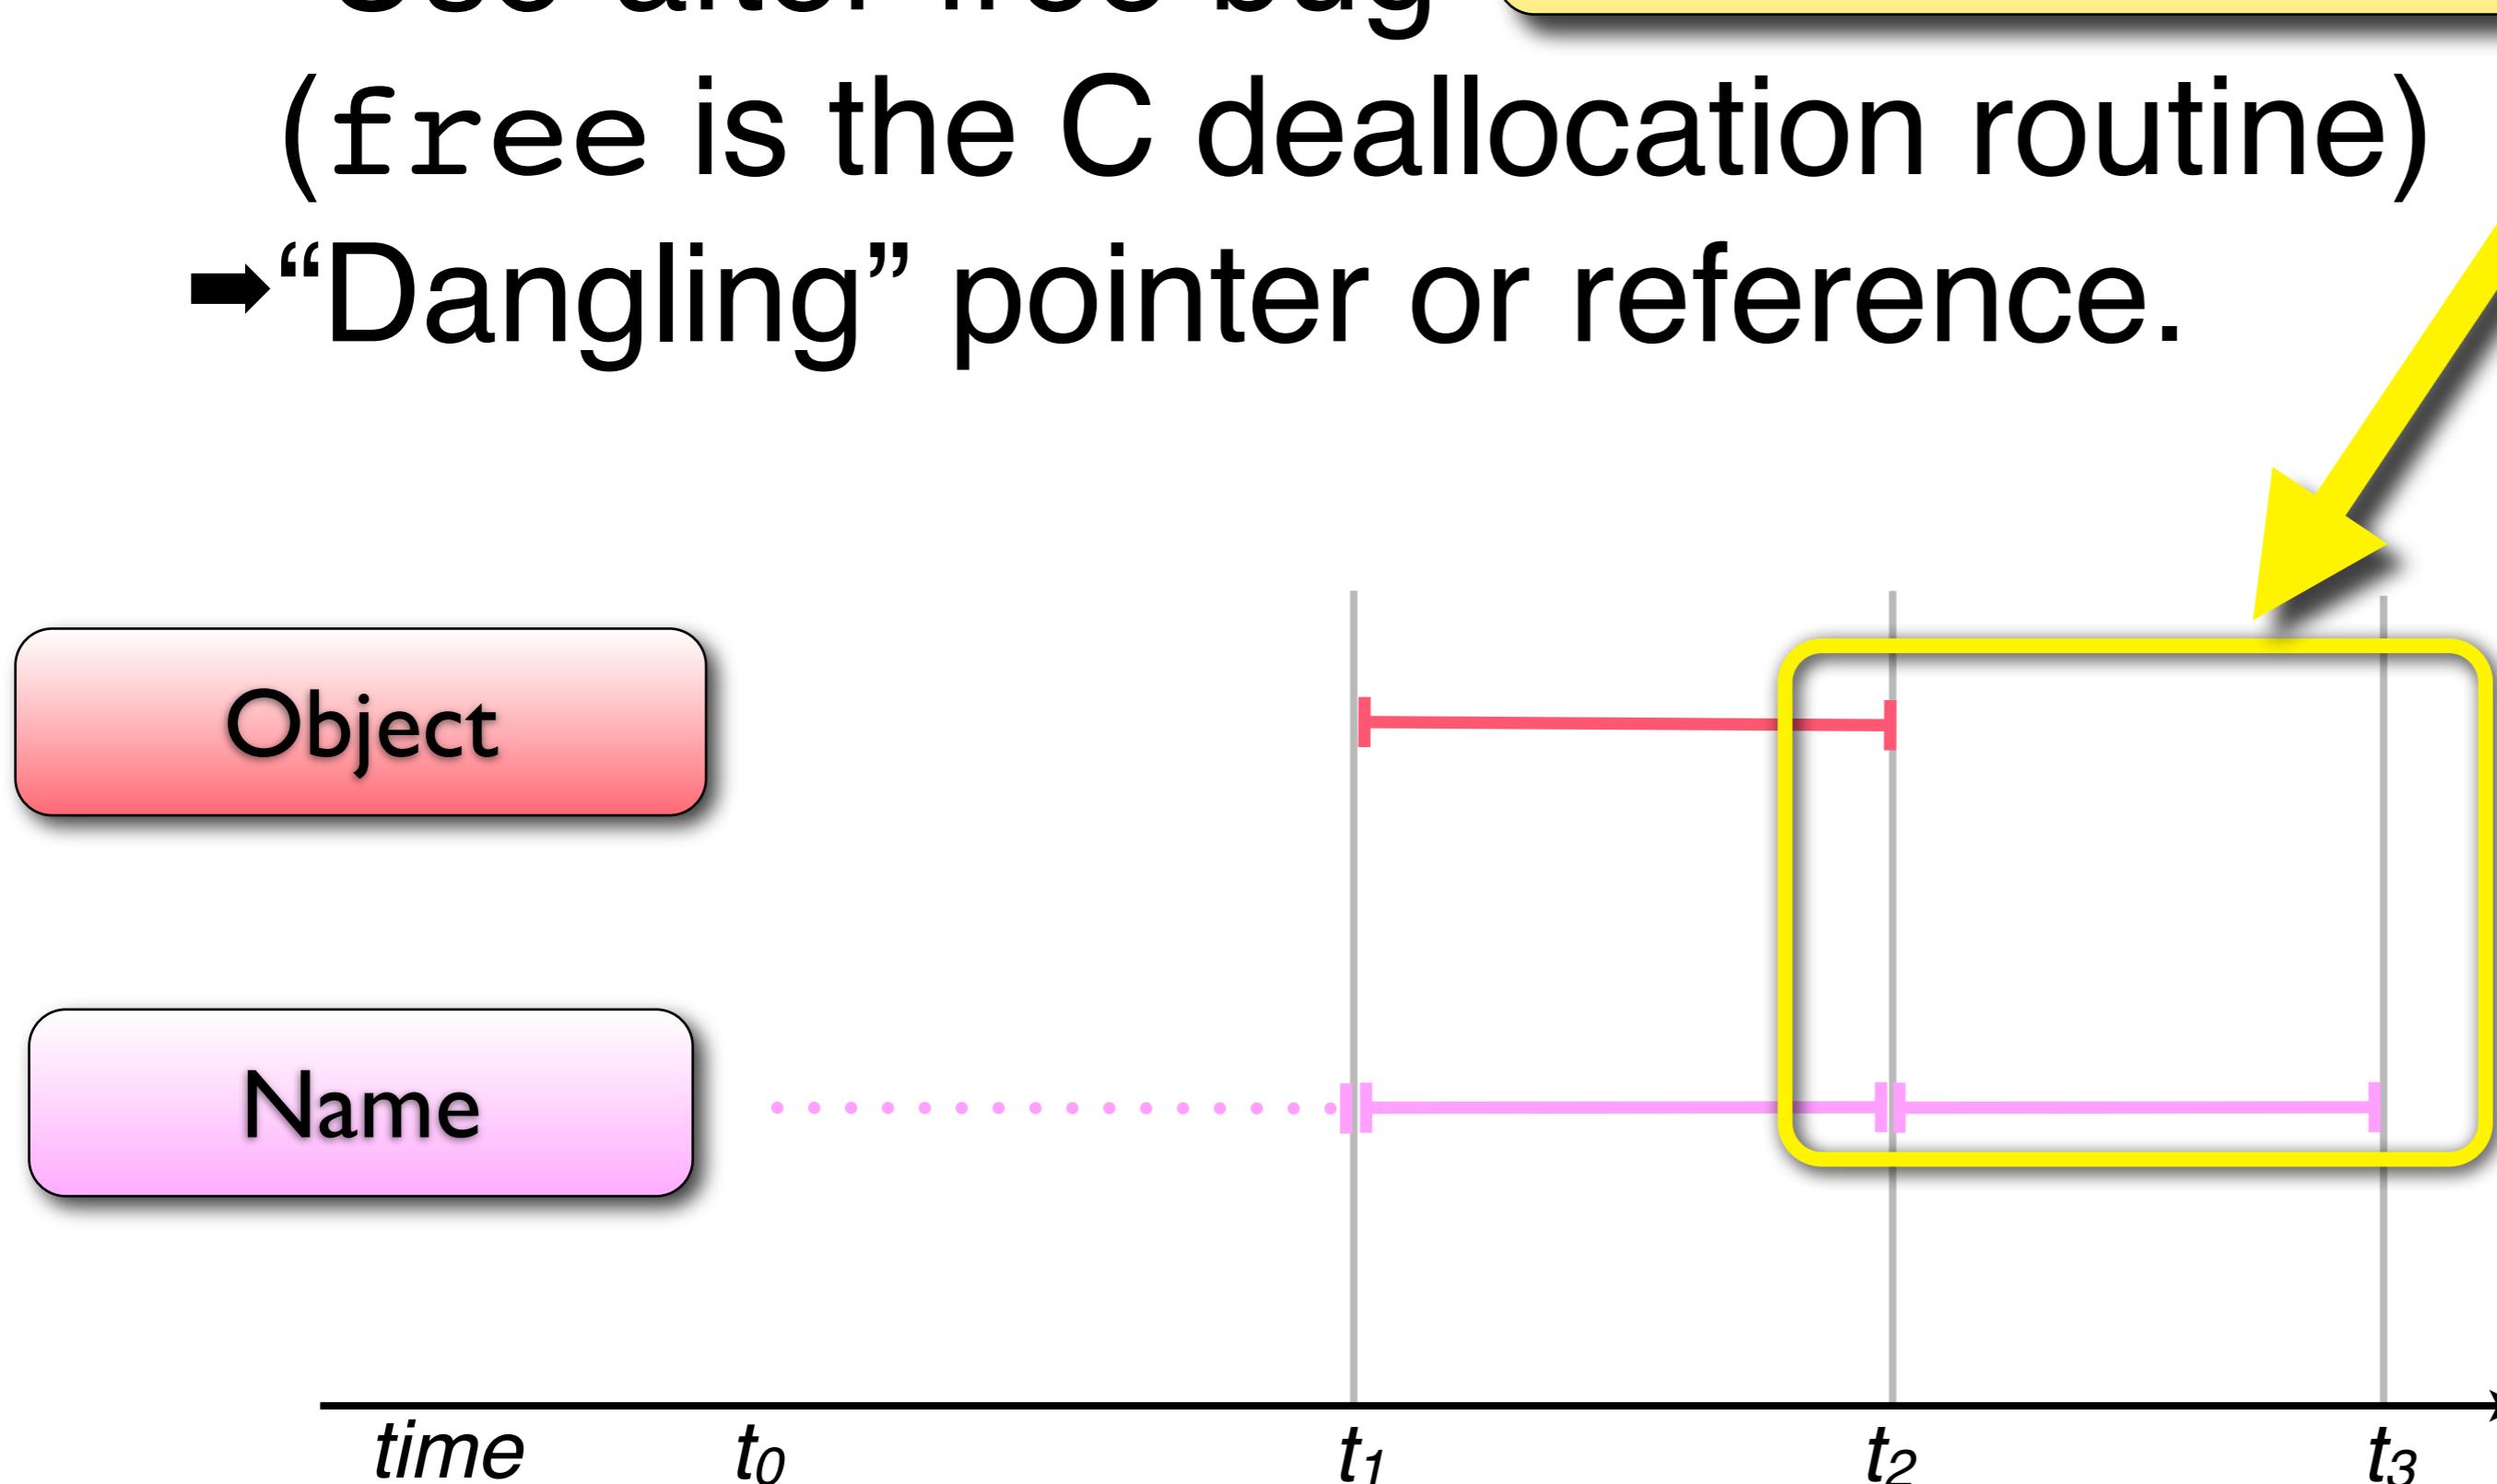


“Dangling” References

Binding / object lifetime mismatch.

- **Binding exists longer than object**
- Object de-allocated
- “Use-after-free bug”
(free is the C deallocation routine)
- “Dangling” pointer or reference.

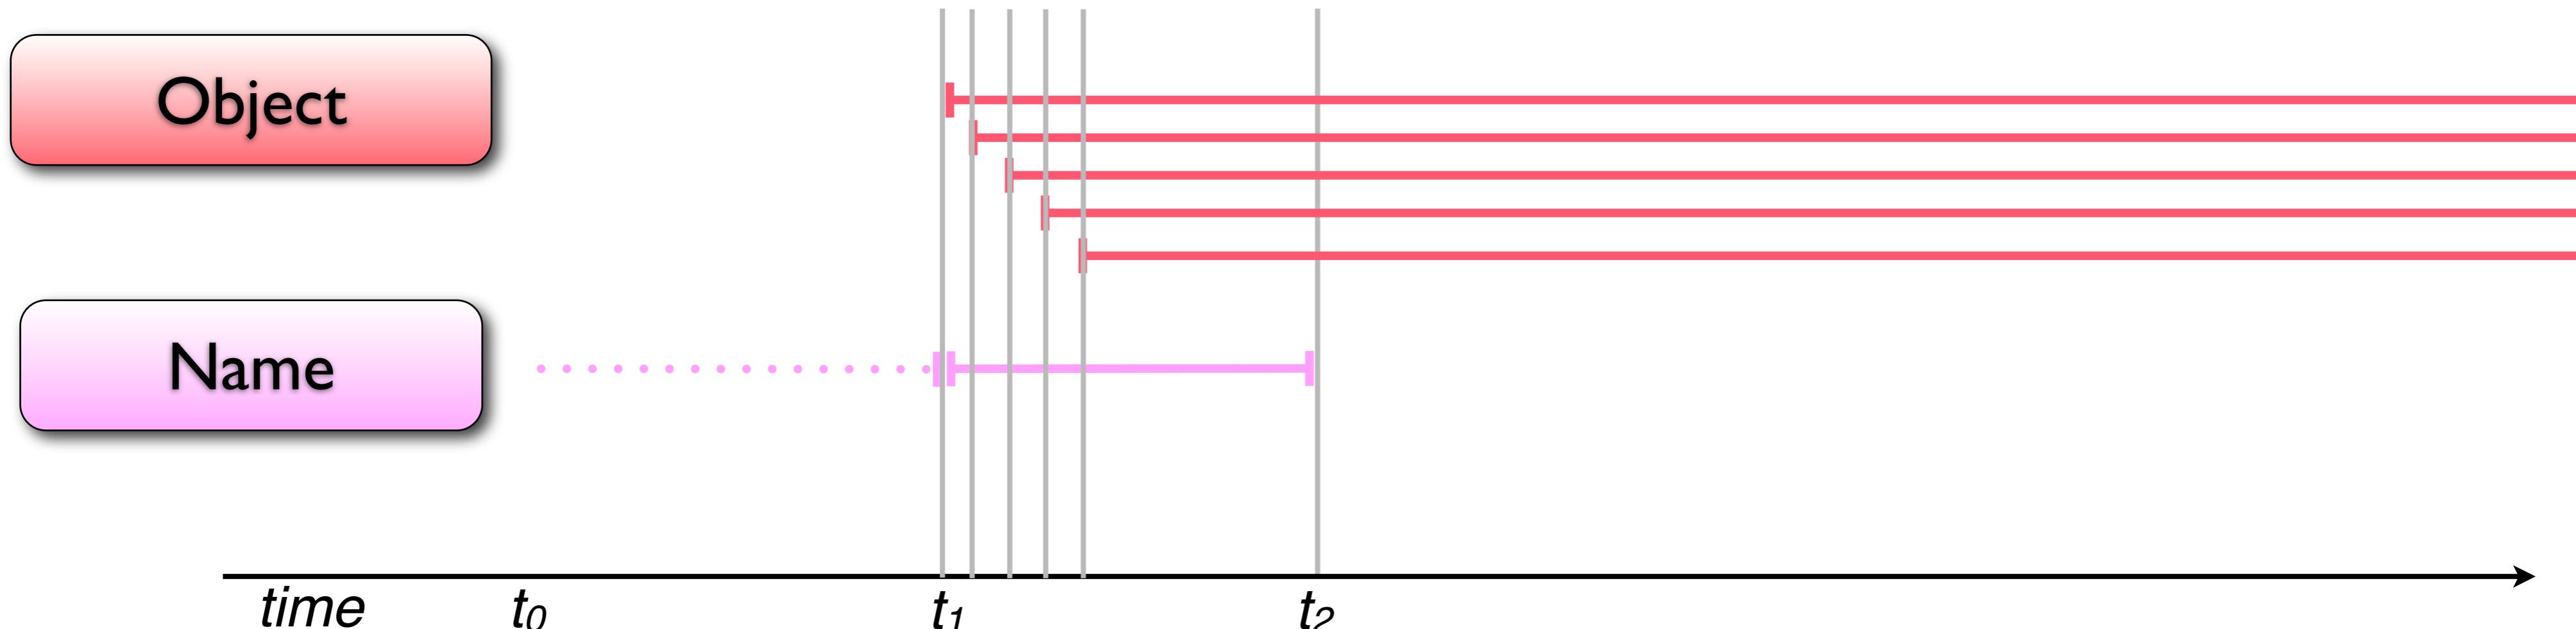
Name bound, but object no longer exists.
Reference is “stale” and “dangles.”



Memory “Leaks”

Omitted deallocation.

- Objects that “**live forever**.”
- Even **if no longer required**.
 - Possibly no longer referenced.
- **Waste memory**; can bring system down.
- A problem in virtually every non-trivial project.

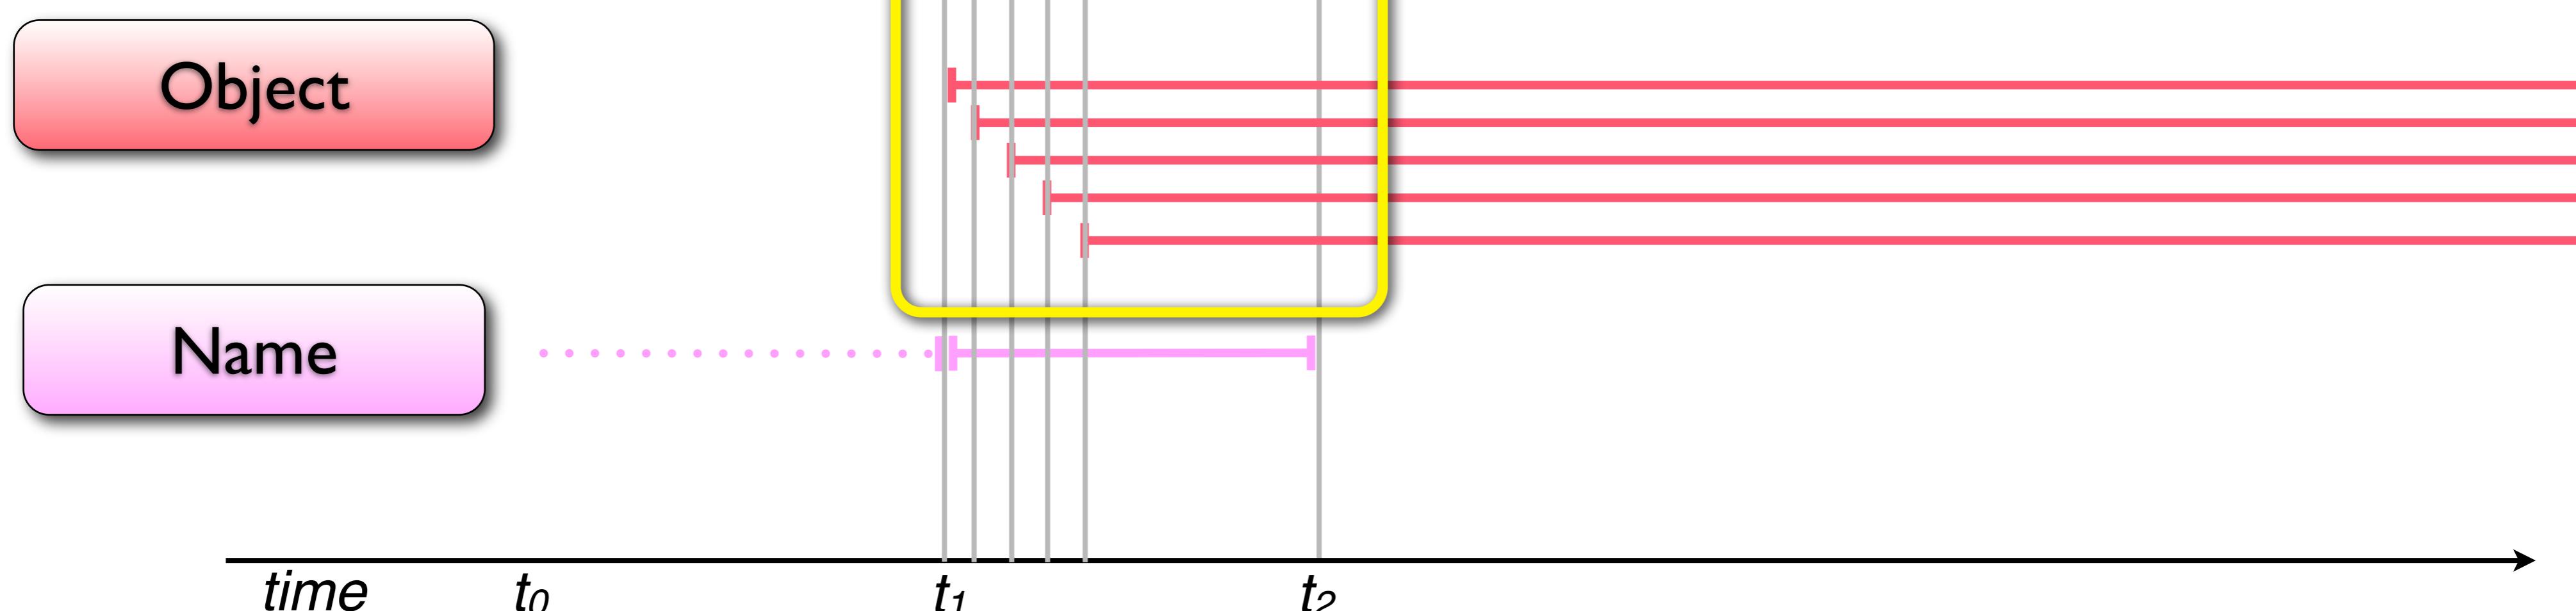


Memory “Leaks”

Omitted deallocation

- Objects that are no longer referenced
- Even if no longer referenced
 - Possibly no longer referenced.
- Waste memory; can bring system down.
- A problem in virtually every non-trivial project.

Objects “forgotten” about, but “stick around”
and waste space until program termination.



Garbage Collection

Manual deallocation is error-prone.

- “Dangling references.”
- “**Use after free.**”
- Possibly unnecessarily **conservative**.
- “Memory leaks.”

Garbage collection.

- **Automatically** deallocates objects **when it is safe to do so.**
- Automated heap management; programmer can **focus on solving real problem.**
- We will focus on garbage collection techniques later in the semester.

Summary & Advise

Static Allocation

Not dynamically sizable; lifetime spans virtually whole program execution; use only sparingly.

Stack Allocation

Lifetime restricted to subroutine invocation;
allocation and deallocation is cheap;
use whenever possible.

Heap Allocation

Arbitrary lifetimes;
use garbage collection whenever possible;
use for large buffers and long-living objects.